

Understanding Garbage Collectors in object-oriented programming

Lecture 4 - Tuning the Garbage Collector

Nahuel Palumbo

✉ nahuel.palumbo@inria.fr



Automatic Memory Management

Garbage Collectors

```
data = malloc(size)  
...  
... use data ...  
...  
free(data)
```

Manually
Memory Management



```
data = Data new  
...  
... use data ...  
...  
????????????
```

Free the space when data
is not used anymore



Compute the size
Allocate in the memory

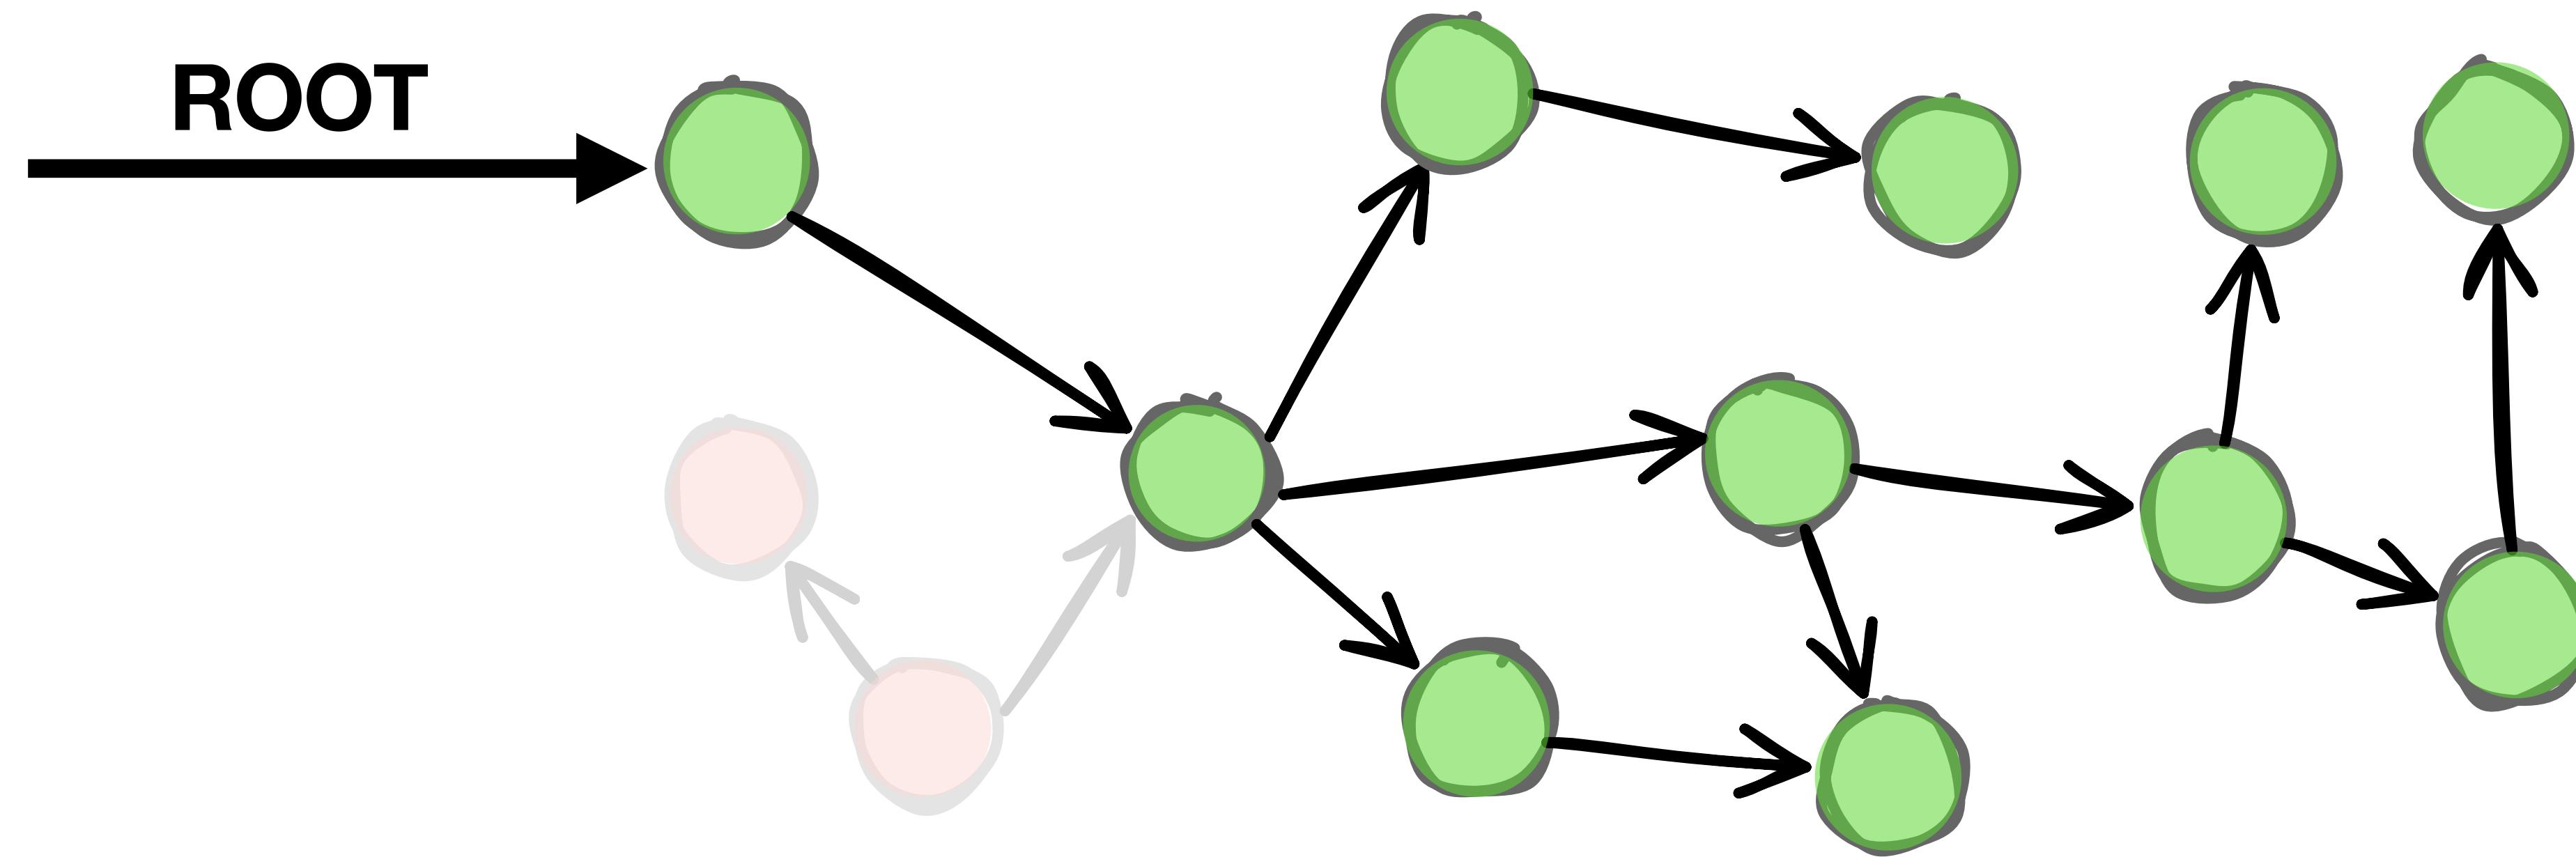


Maybe move the data for
better use of the memory



Application's Allocation Patterns

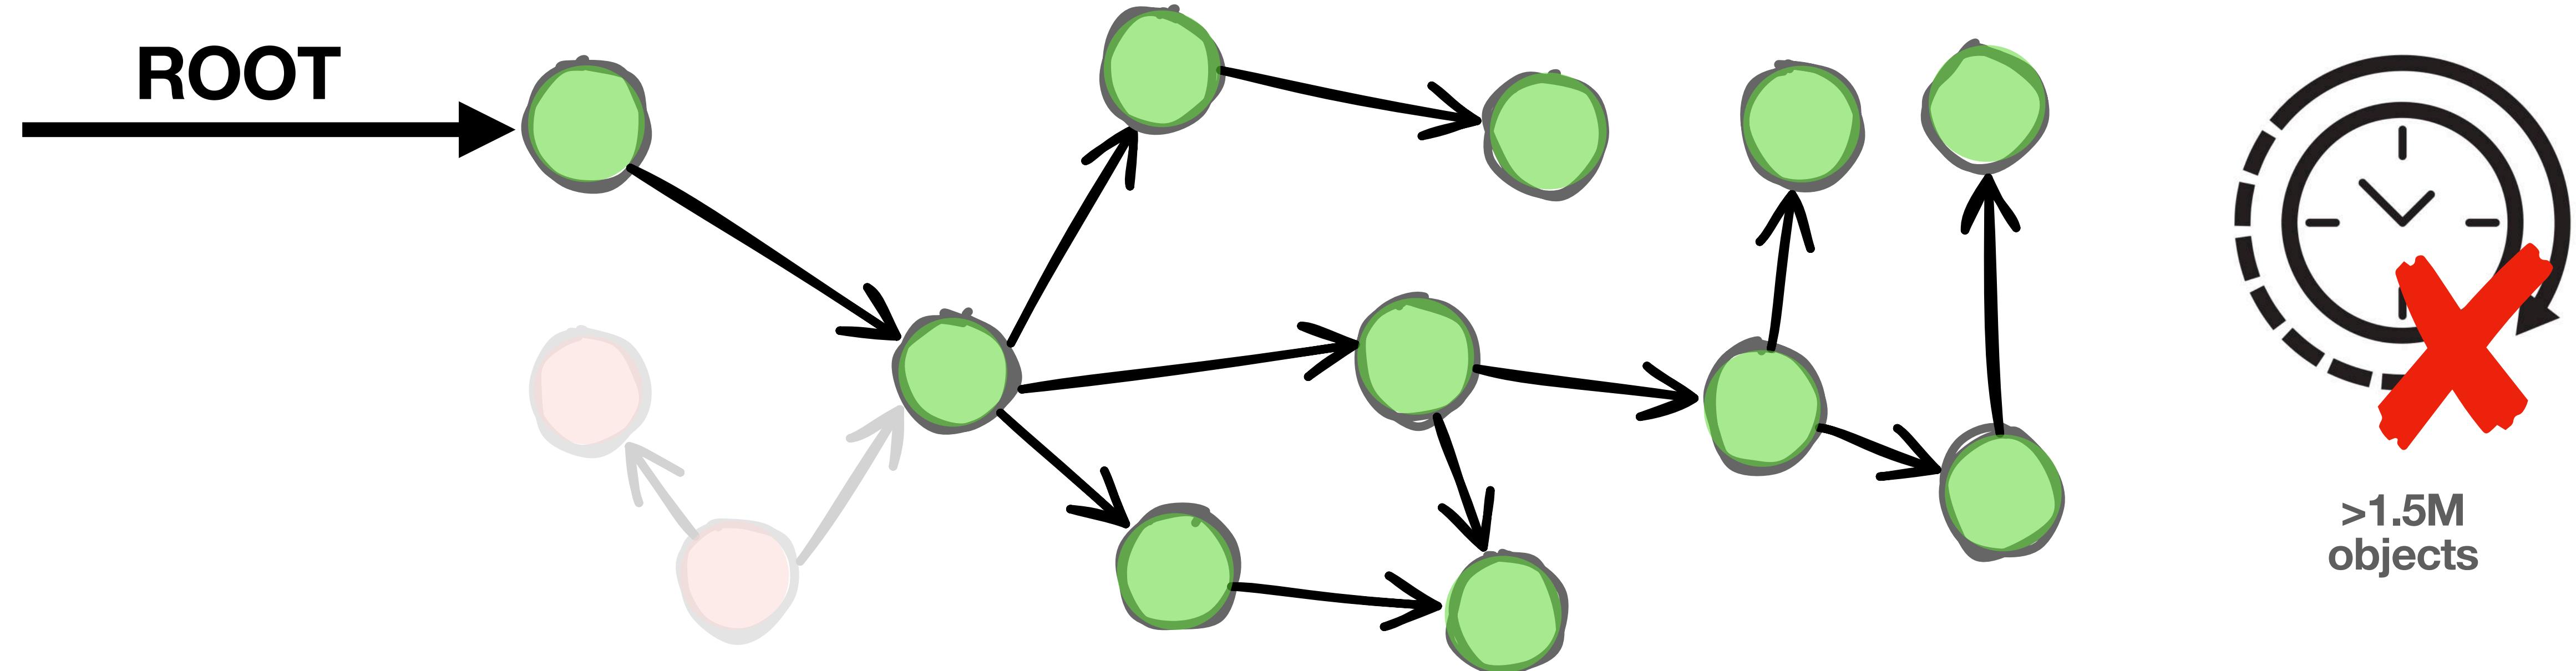
When an object dies?



“Must be accessible from the roots”

Application's Allocation Patterns

When an object dies?

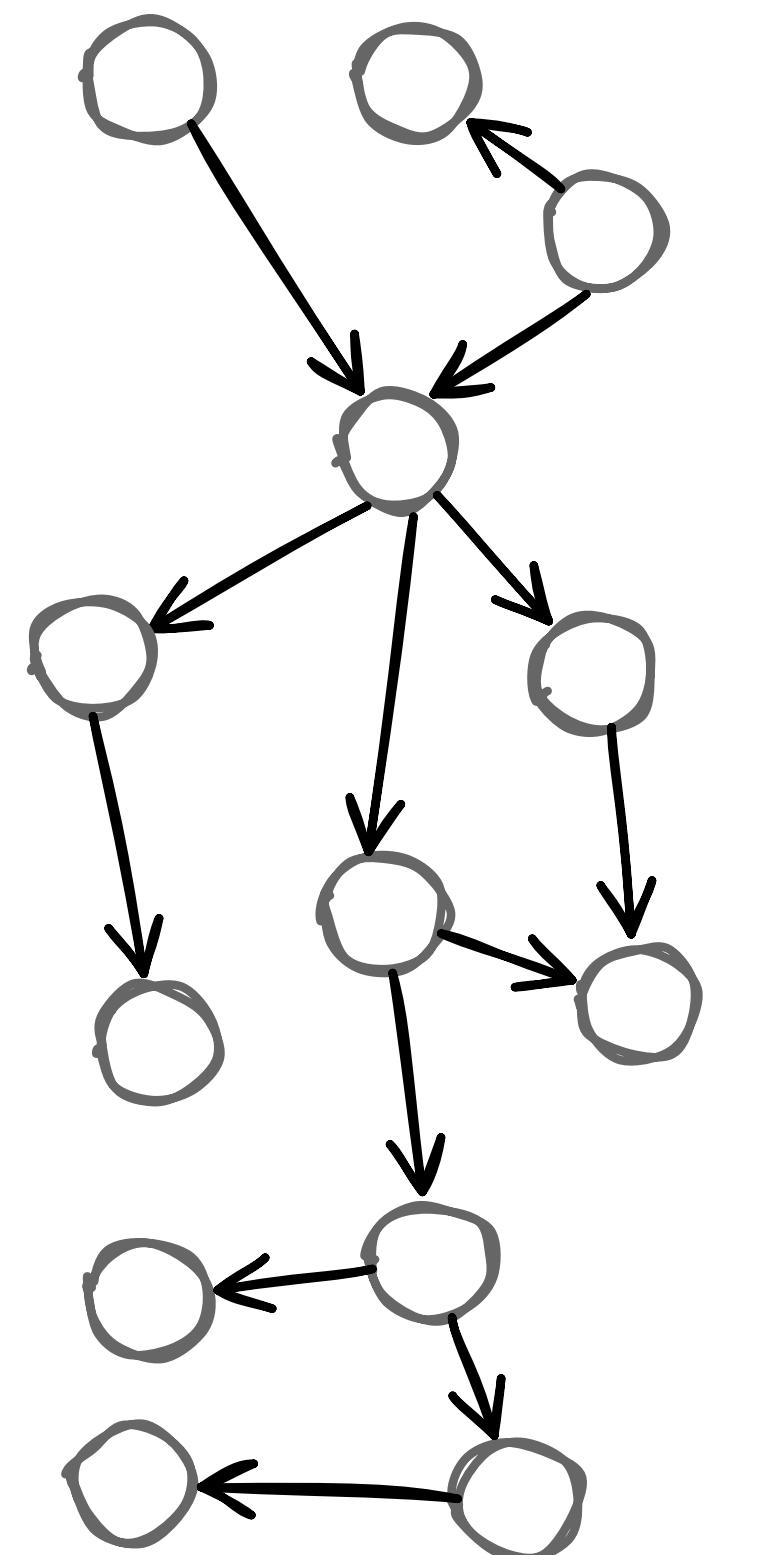


“Must be accessible from the roots”



Application's Allocation Patterns

How do the applications use the memory?

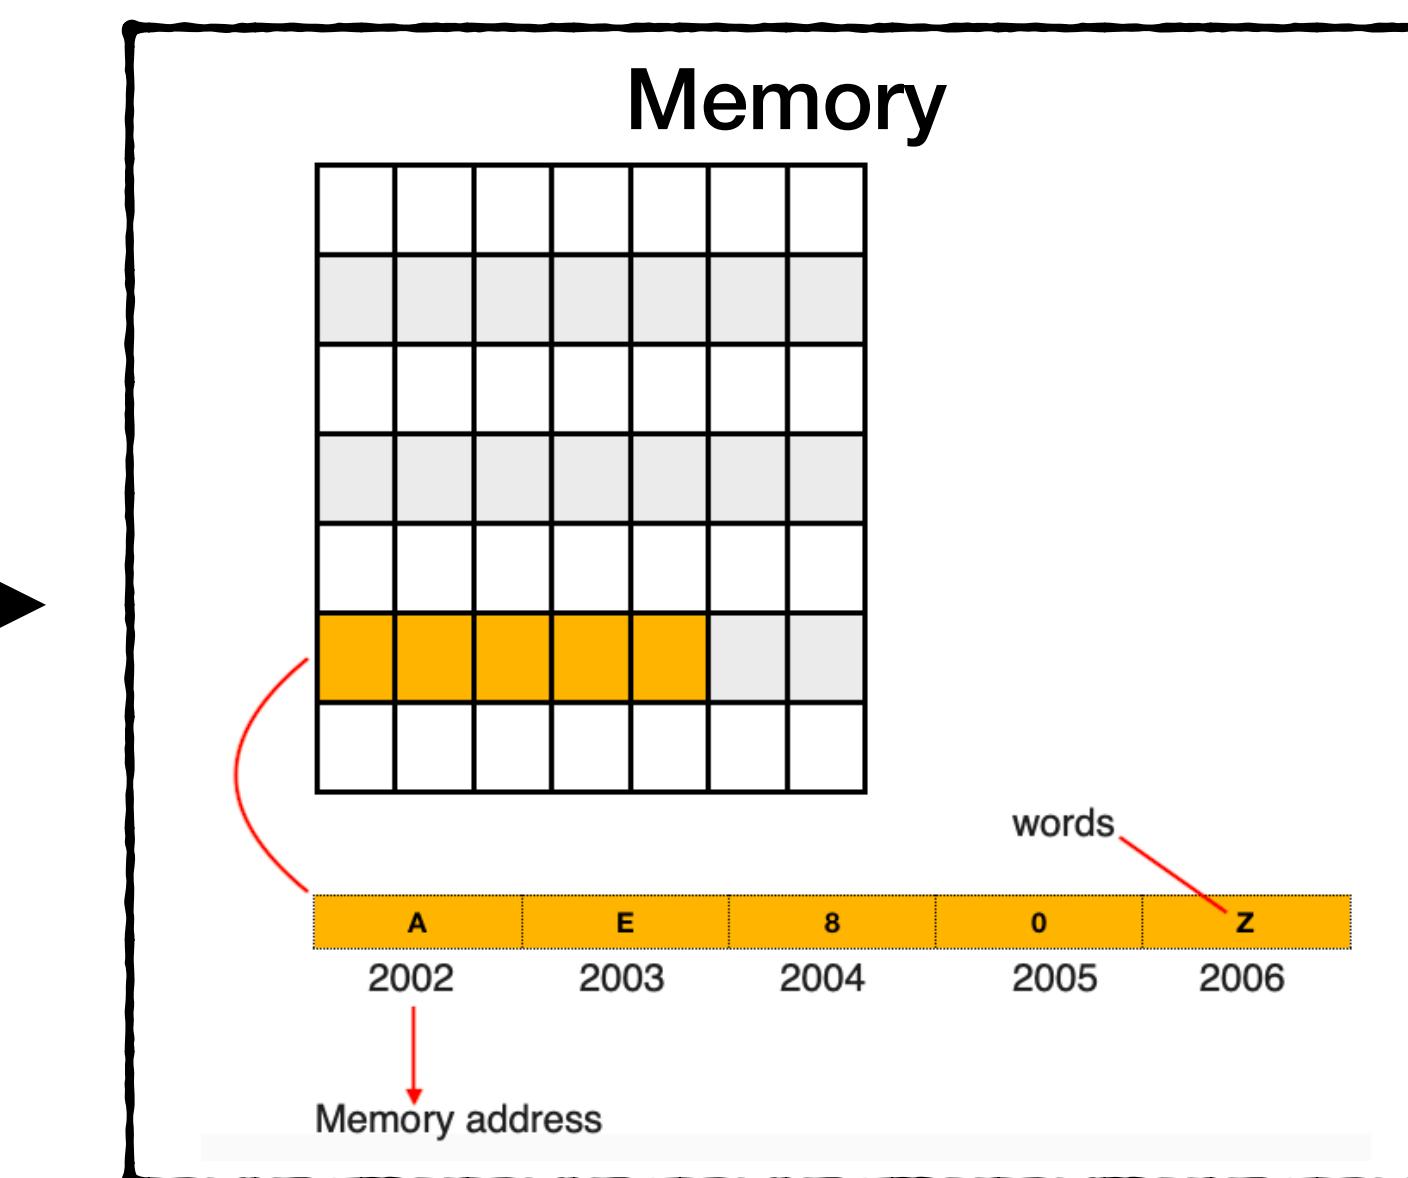


Allocations are particular for each application

→ Hard to predict

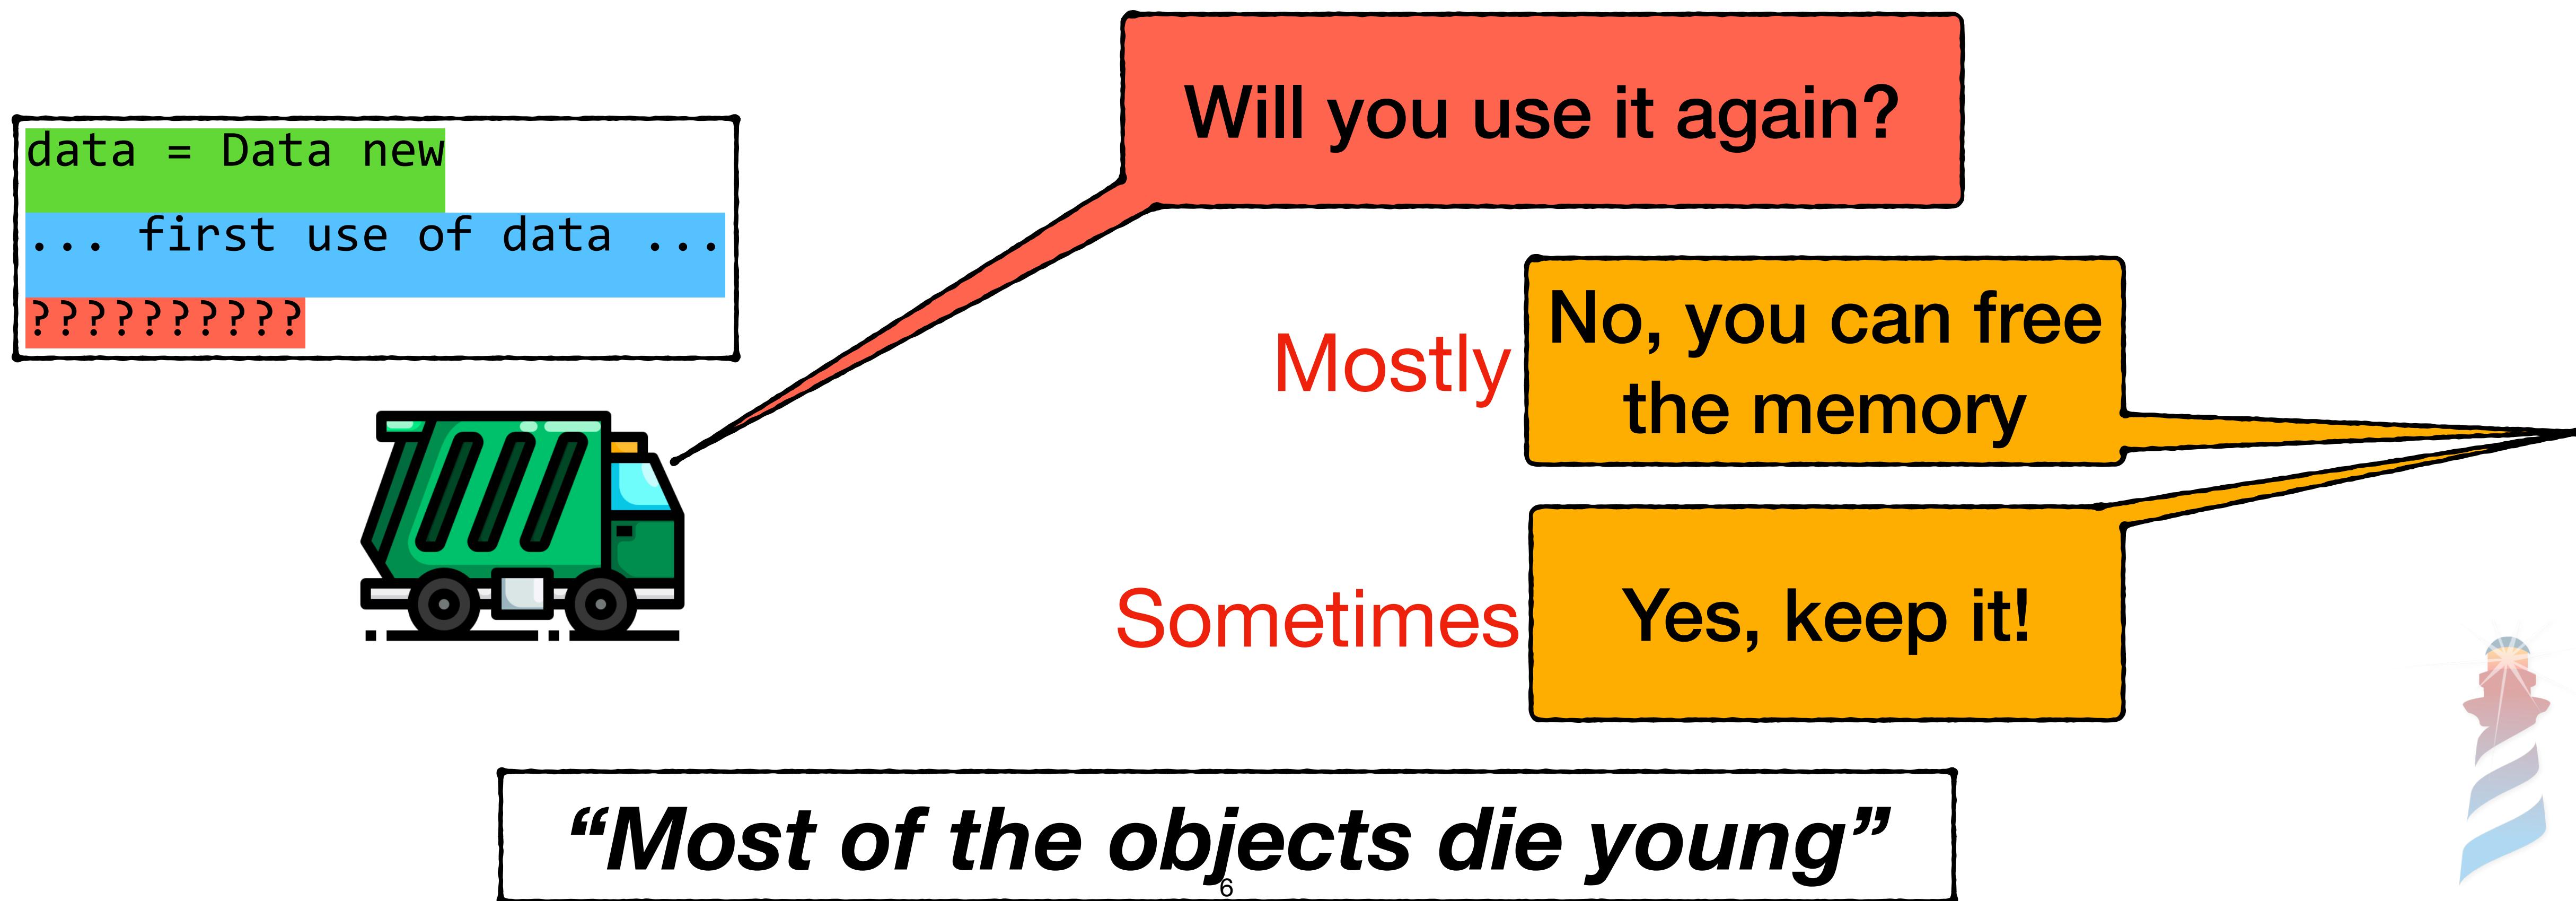
There are some general heuristics

→ Easy to predict



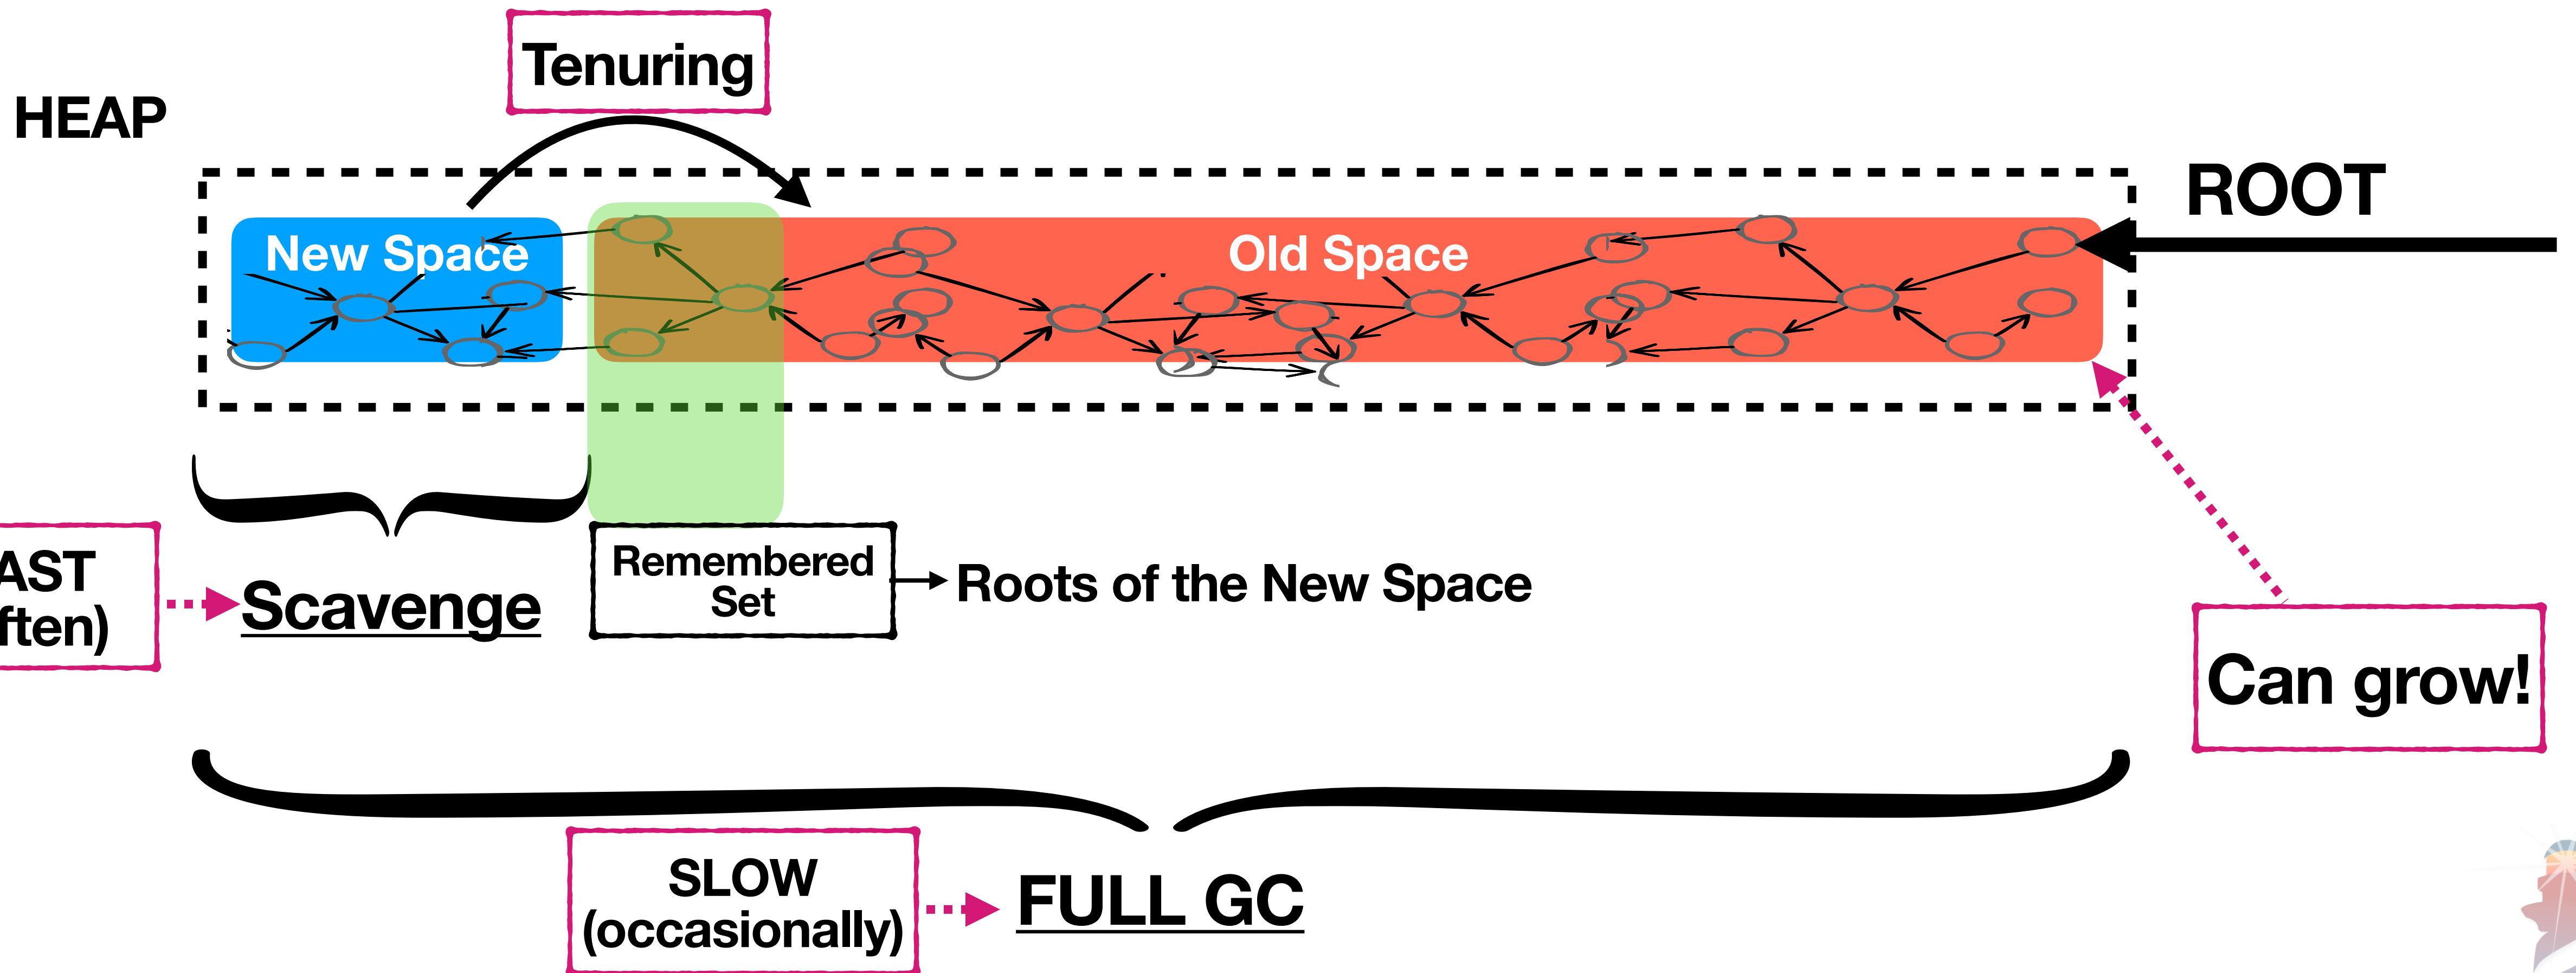
Application's Allocation Patterns

Weak generational hypothesis



Generational Garbage Collector

High-Performance Automatic Memory Management for OOP



Pathological Allocation Pattern

Generational Garbage Collectors' problem

Weak generational
Stable memory use



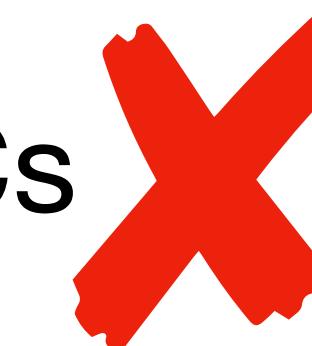
Few Full GCs
Fast Scavengers



Long lifetimes
Memory-starved



Many Full GCs
Scavenger overhead



DataFrame

“Pharo is slow”



DataFrame

Playground

Do it Publish Bindings Versions Pages

```
1 DataFrame readFromCsv: pathToFile asFileReference.
```

Line: 1:1 +L



DataFrame

“Pharo The GC is slow”

Data size	Total time (sec)	GC time (sec)	GC overhead
529 MB	43	7	16%
1.6 GB	150	38	25%
3.1 GB	5599	5158	92%



DataFrame

```
Playground
Do it Publish Bindings Versions Pages
1 DataFrame readFromCsv: pathToFile asFileReference.
Line: 1:1 +L
```



DataFrame

“Pharo The GC is slow”

Data size	Total time (sec)	GC time (sec)	GC overhead
529 MB	43	7	16%
1.6 GB	150	38	25%
3.1 GB	5599 >1h30m	5158	92% !!!!



DataFrame

The screenshot shows the Pharo playground interface. The title bar says "Playground". The toolbar includes "Do it", "Publish", "Bindings", "Versions", and "Pages". A message in the workspace area reads: "1 DataFrame readFromCsv: pathToFile asFileReference.". Below the workspace, it says "Line: 1:1". In the bottom right corner of the workspace, there is a blue button with the text "+L". To the right of the playground window, the text "~1h25m" is written in red.



Pathological Allocation Pattern

Garbage Collectors' problem

Weak generational
Stable memory use



Few Full GCs
Fast Scavengers



Long lifetimes
Memory-starved



Many Full GCs
Scavenger overhead



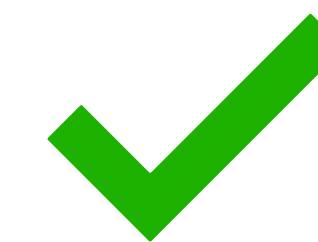
Pathological Allocation Pattern

Garbage Collectors' problem

Weak generational
Stable memory use



Few Full GCs
Fast Scavengers



Long lifetimes
Memory-starved

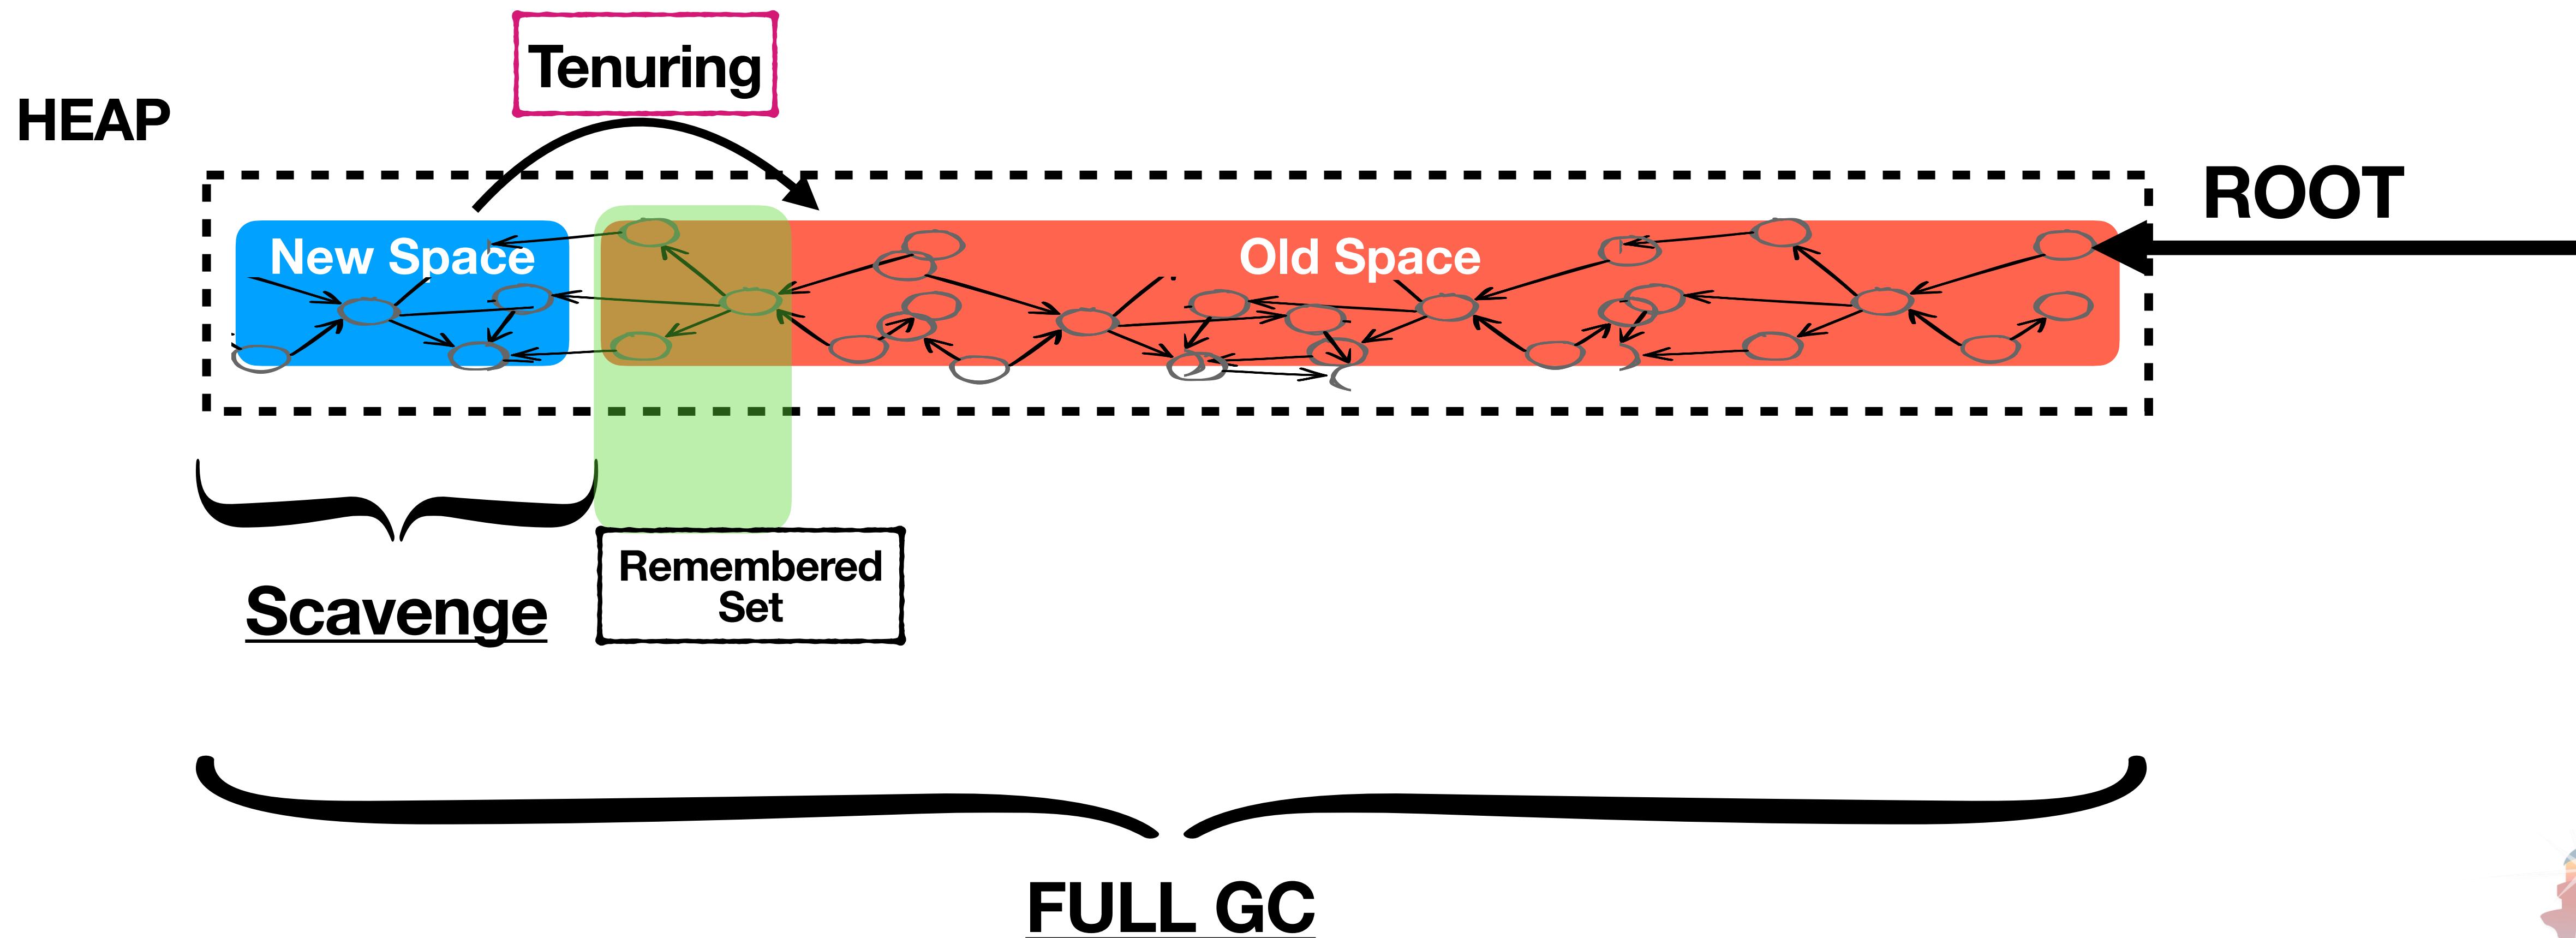


Many Full GCs
Scavenger overhead



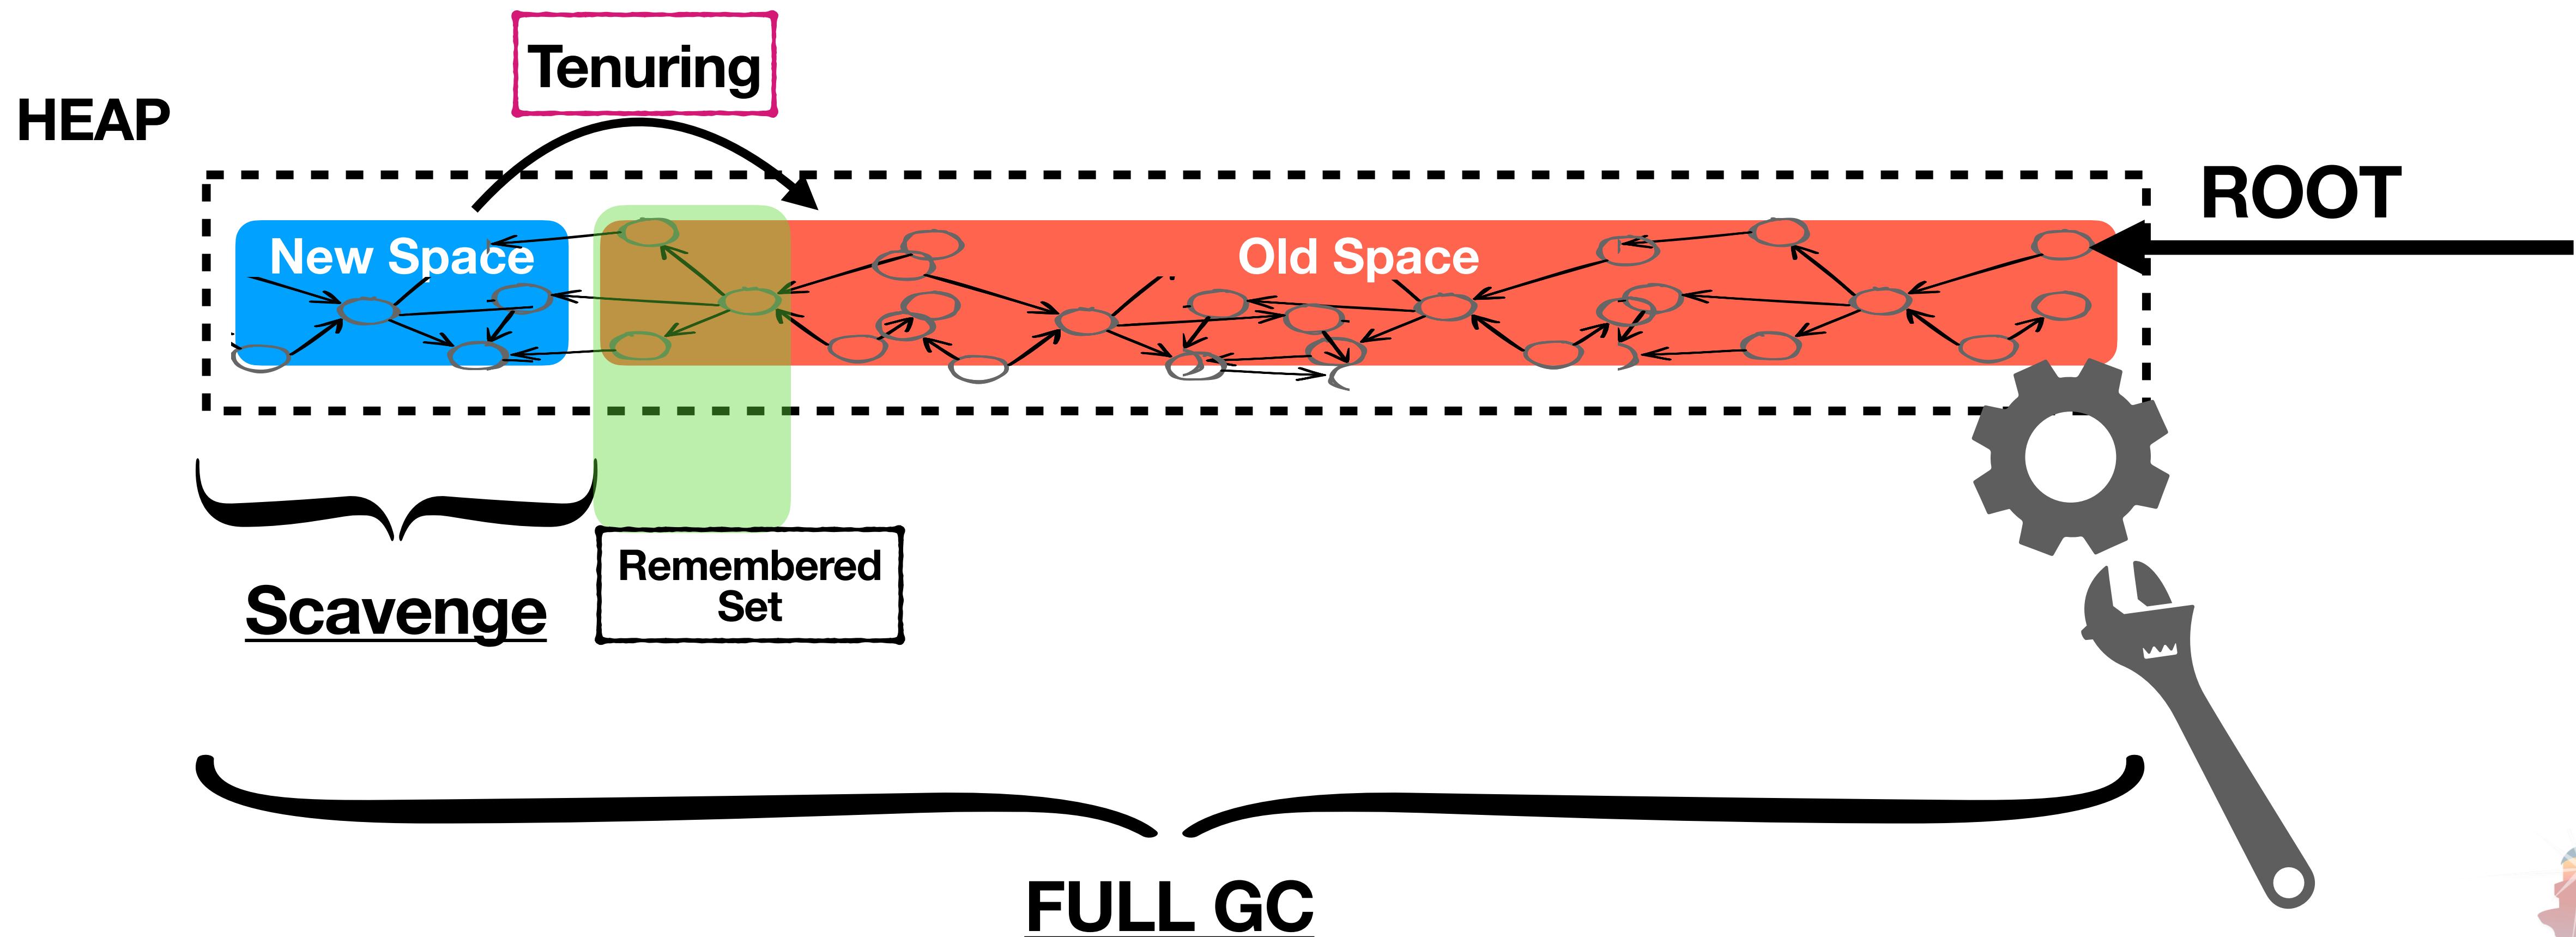
Pathological Allocation Pattern

Tuning the Garbage Collector



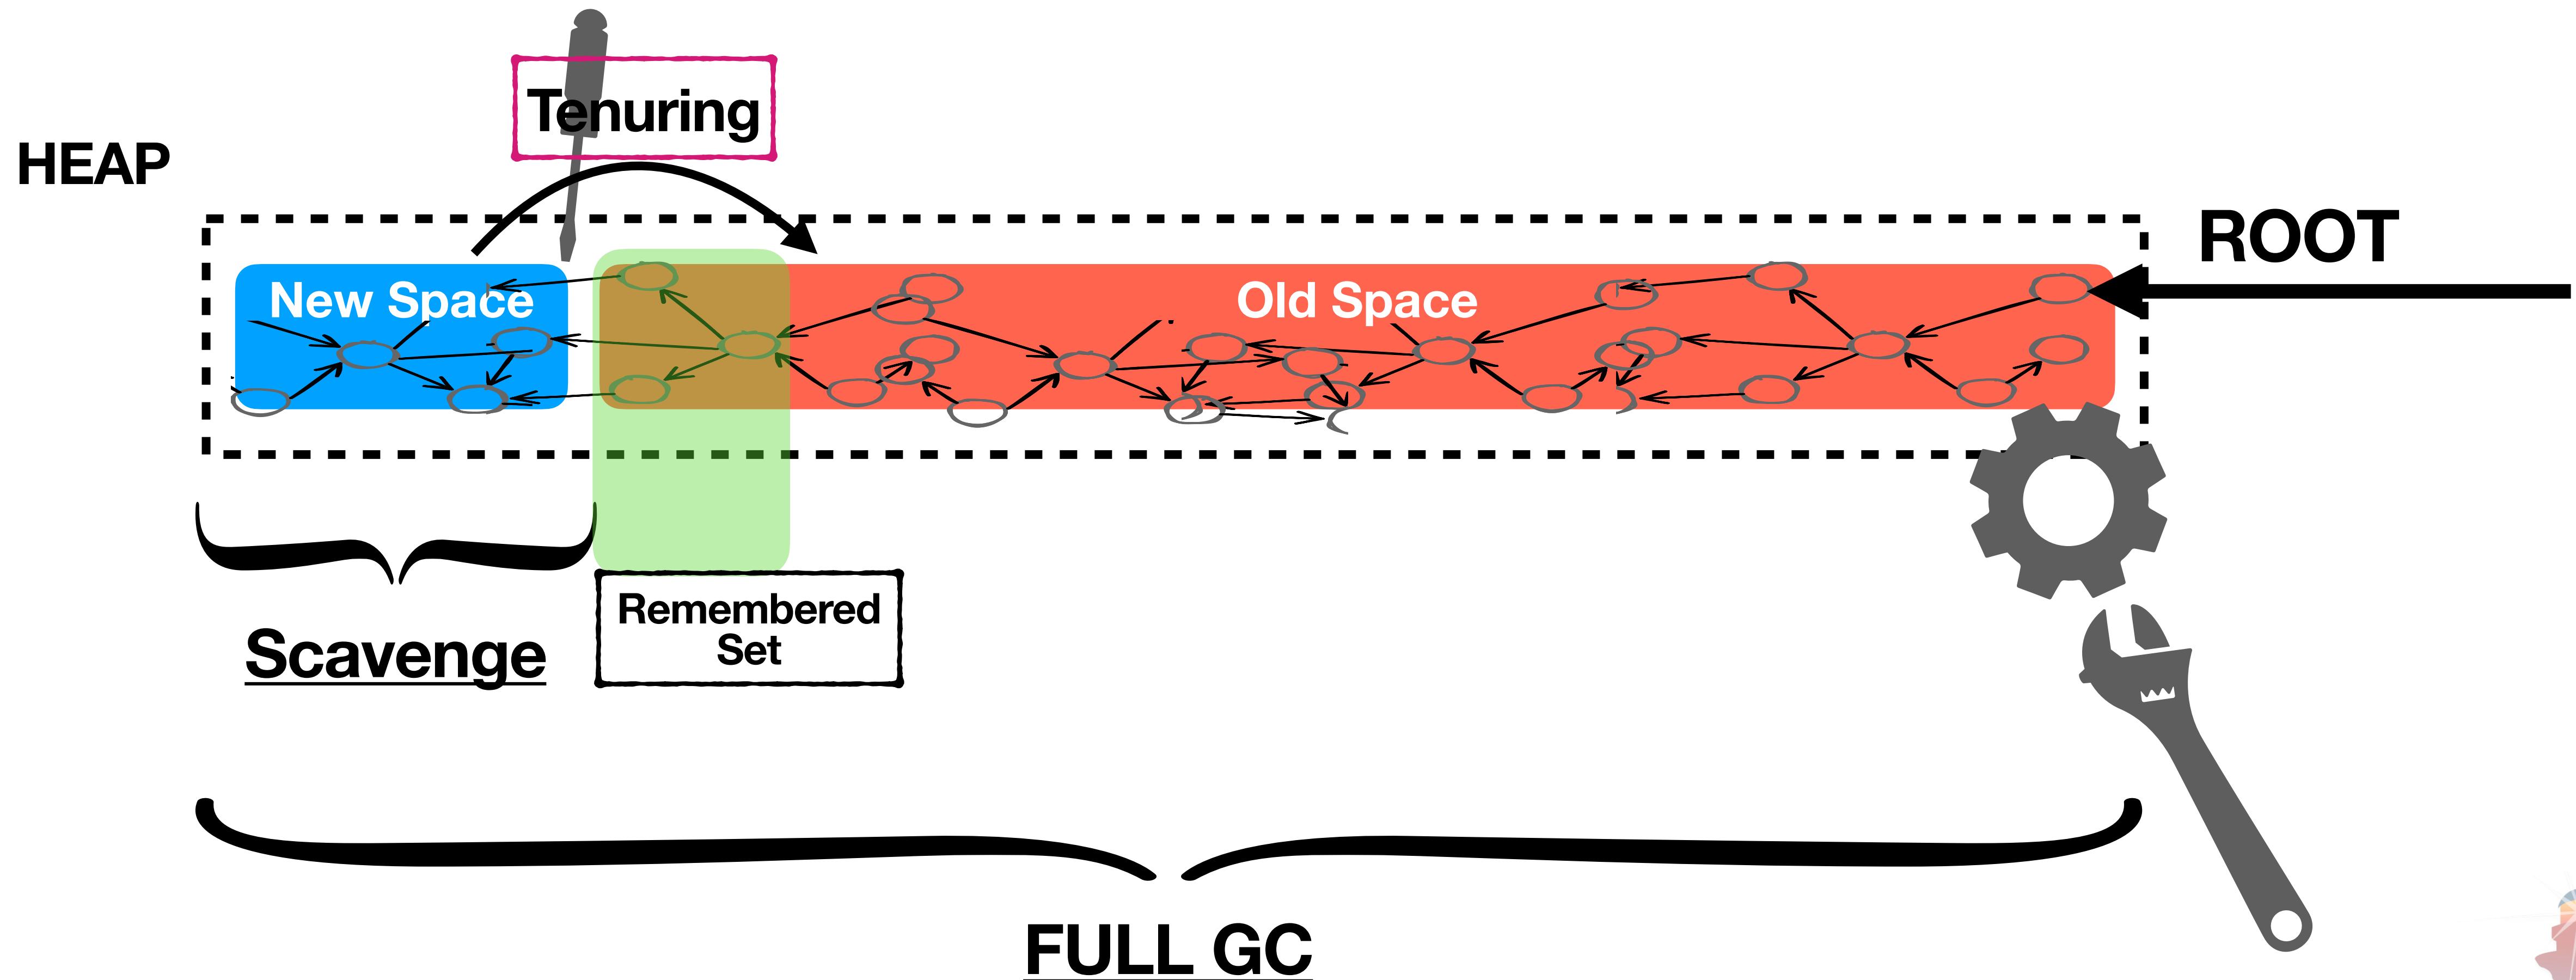
Pathological Allocation Pattern

Tuning the Garbage Collector



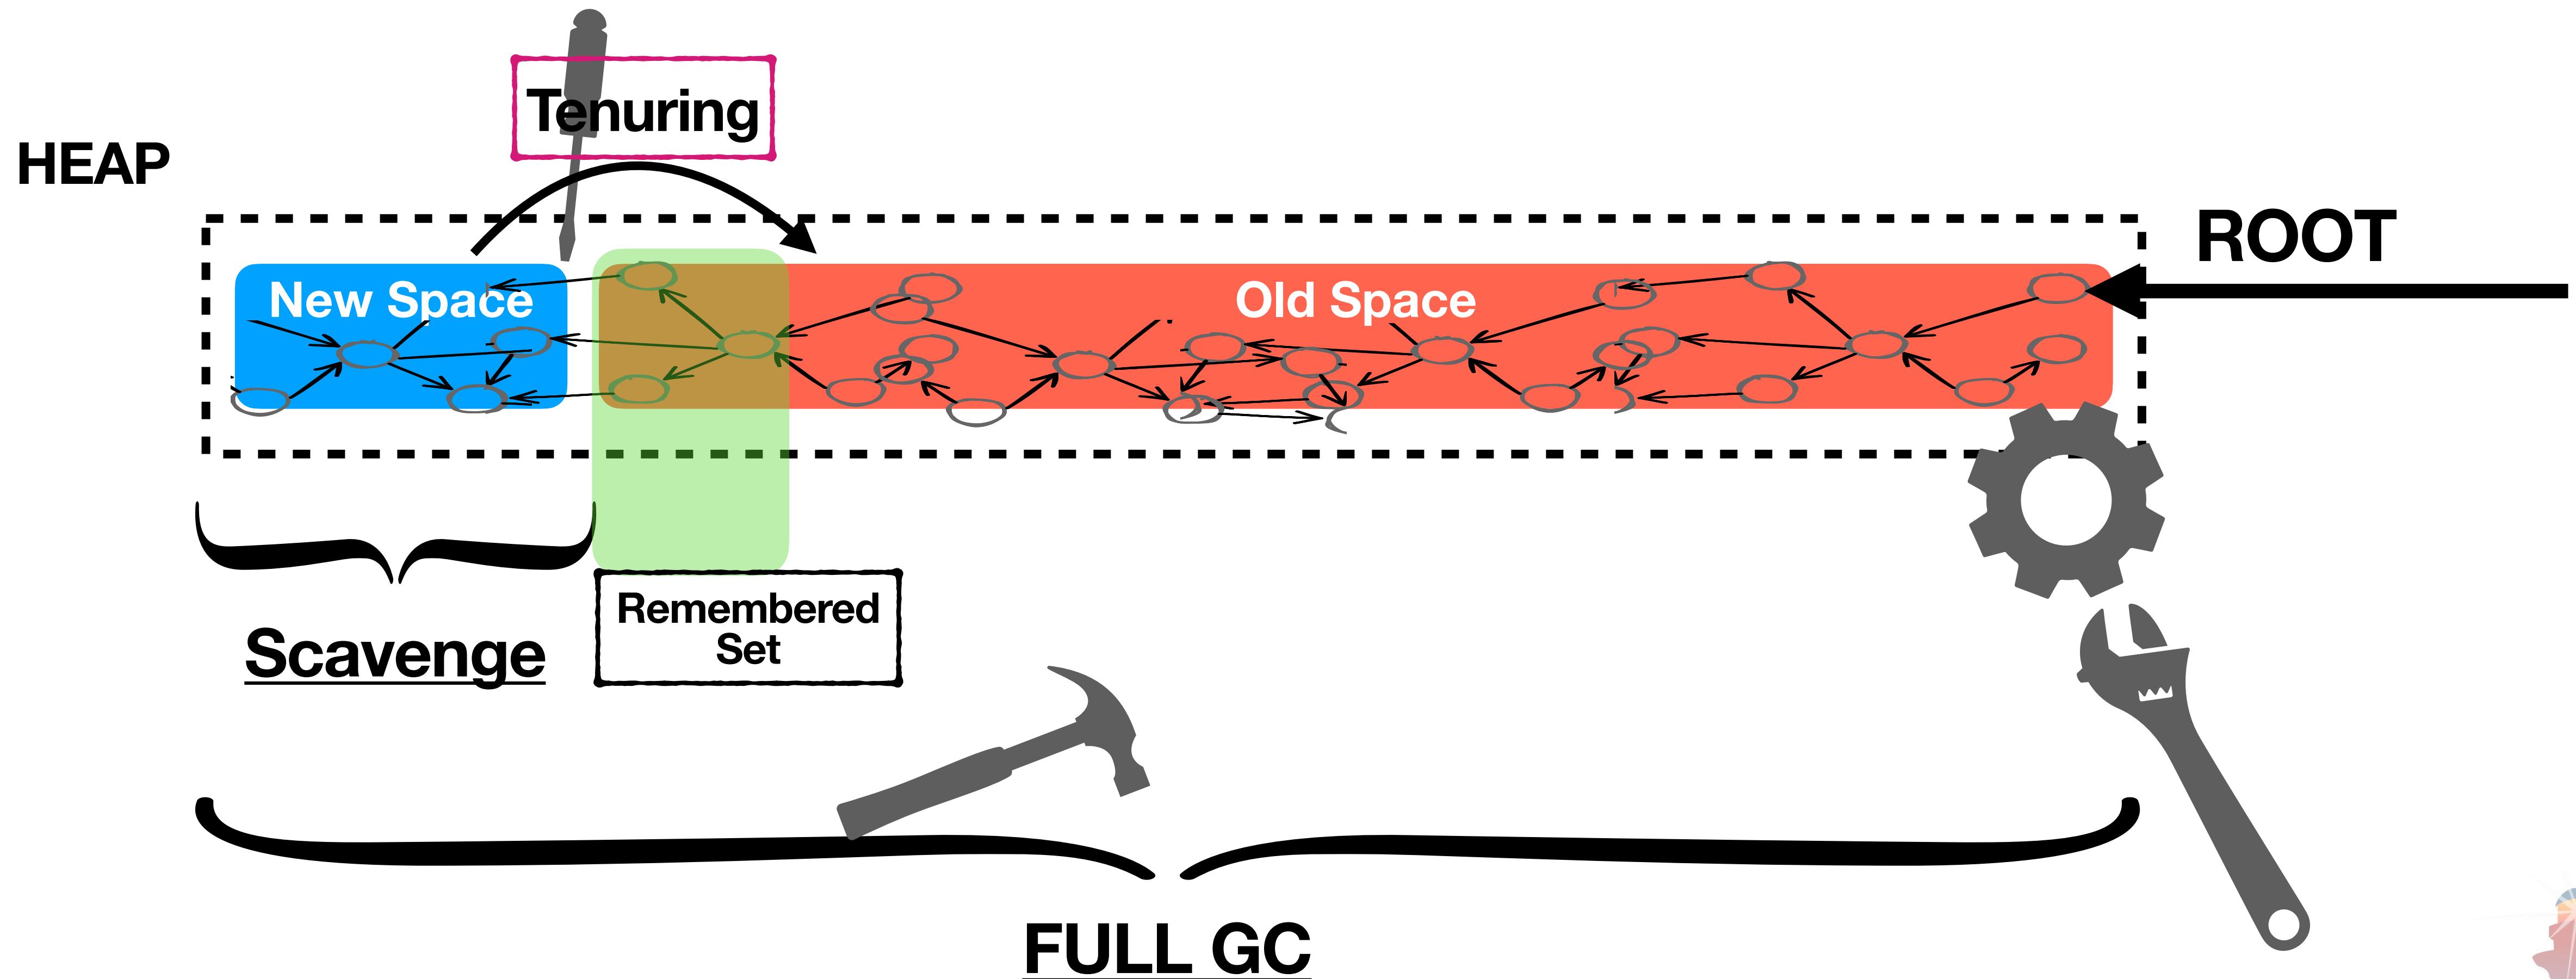
Pathological Allocation Pattern

Tuning the Garbage Collector



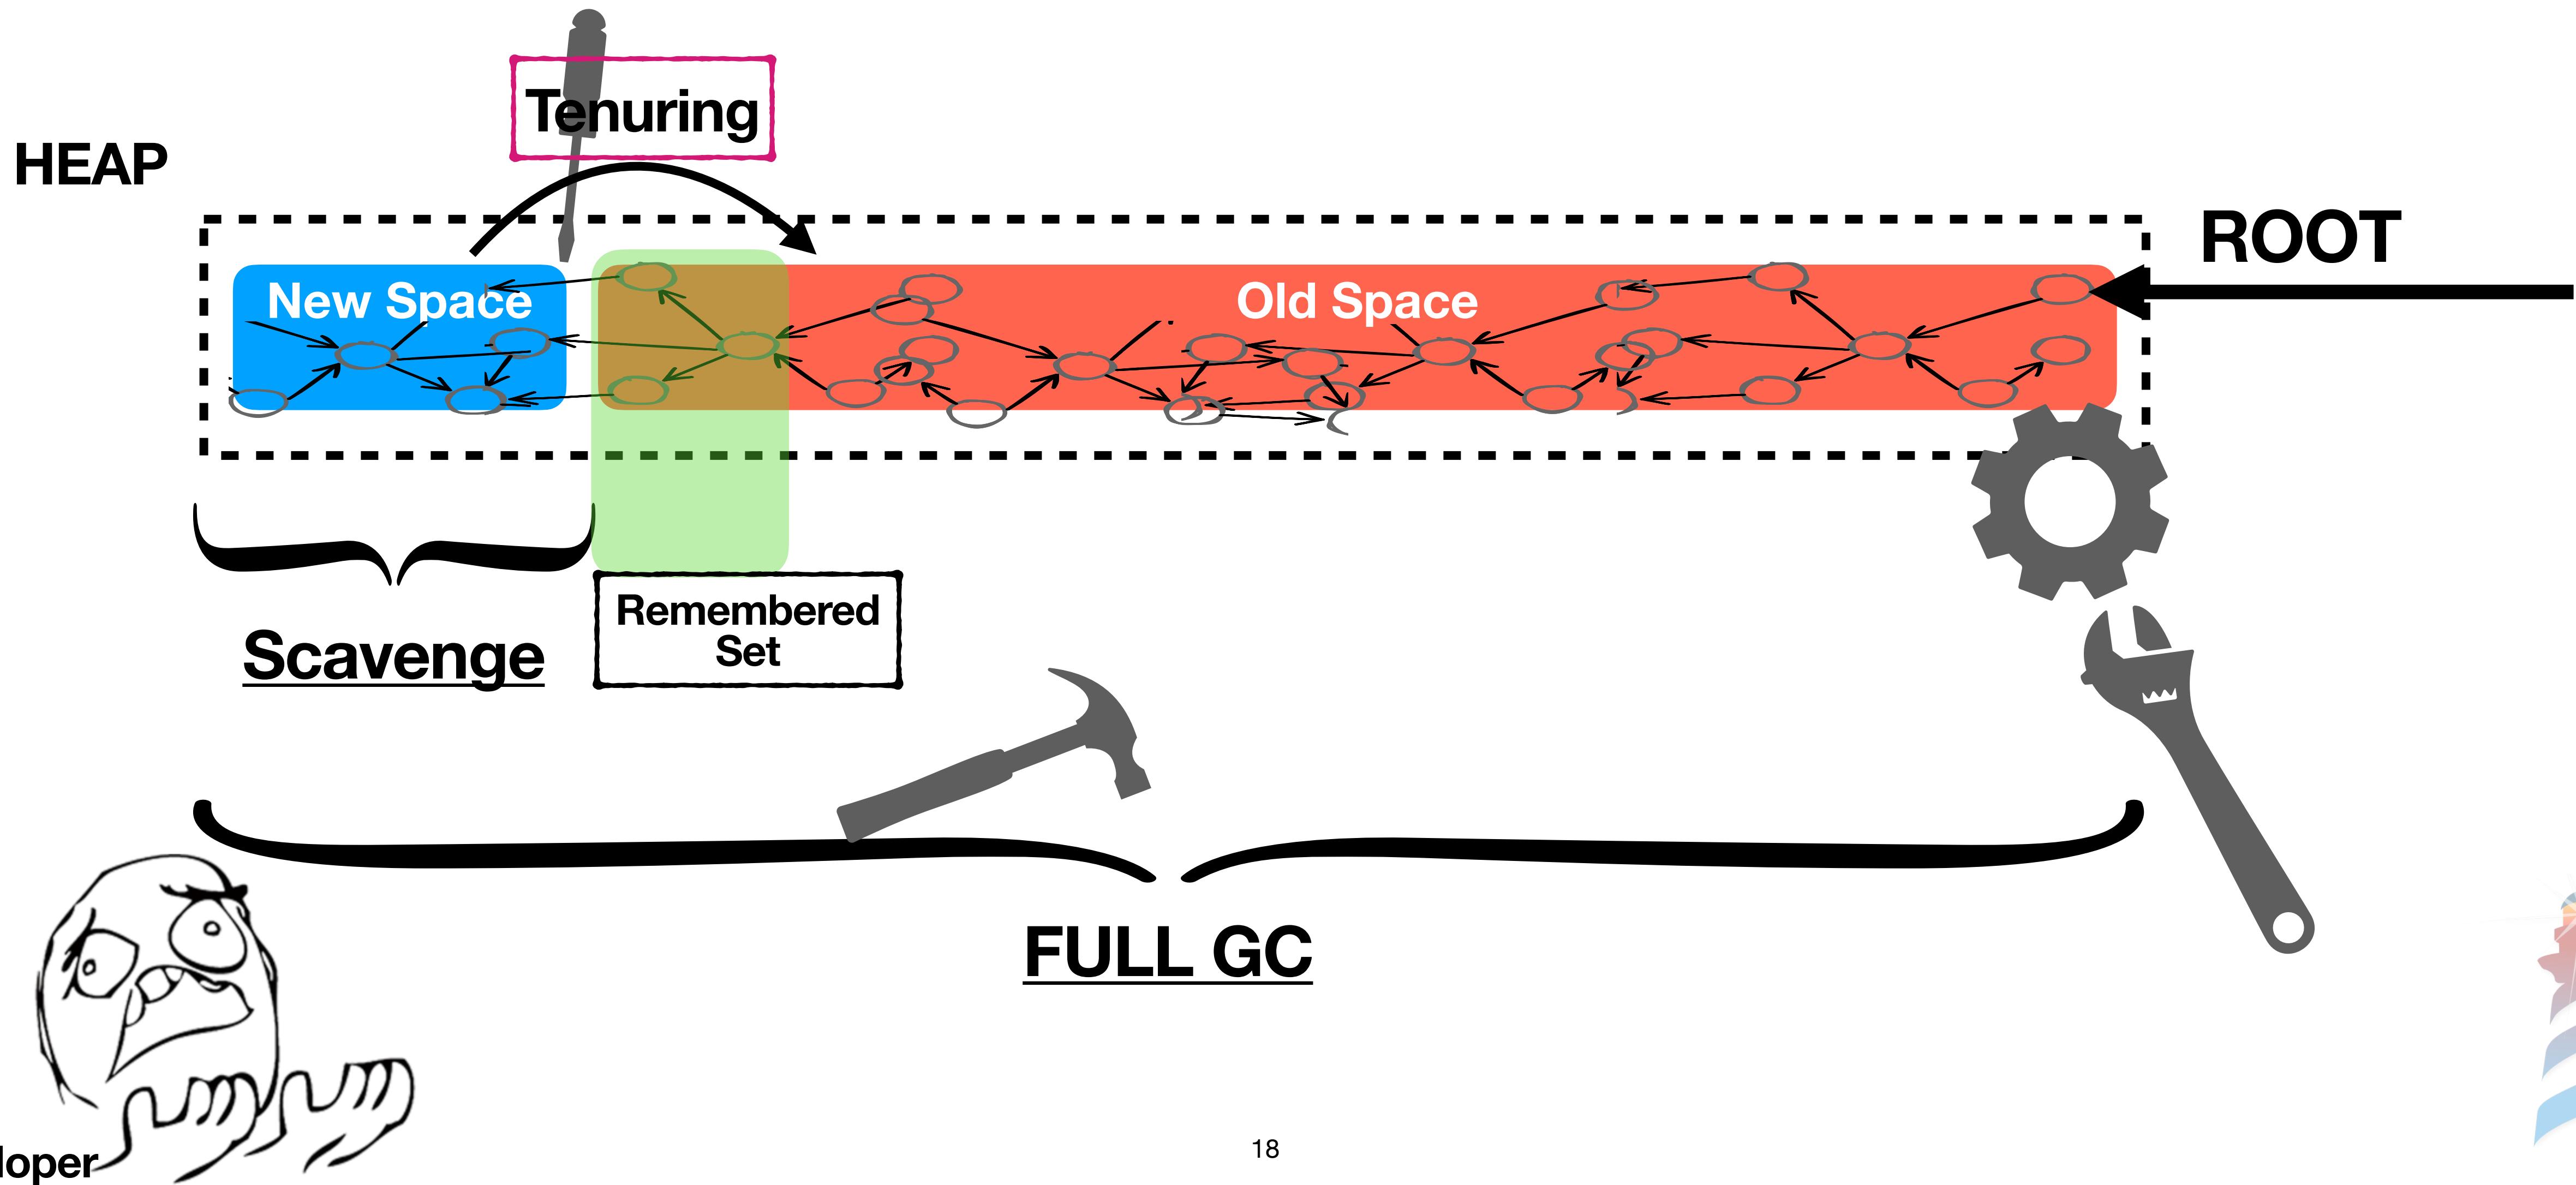
Pathological Allocation Pattern

Tuning the Garbage Collector

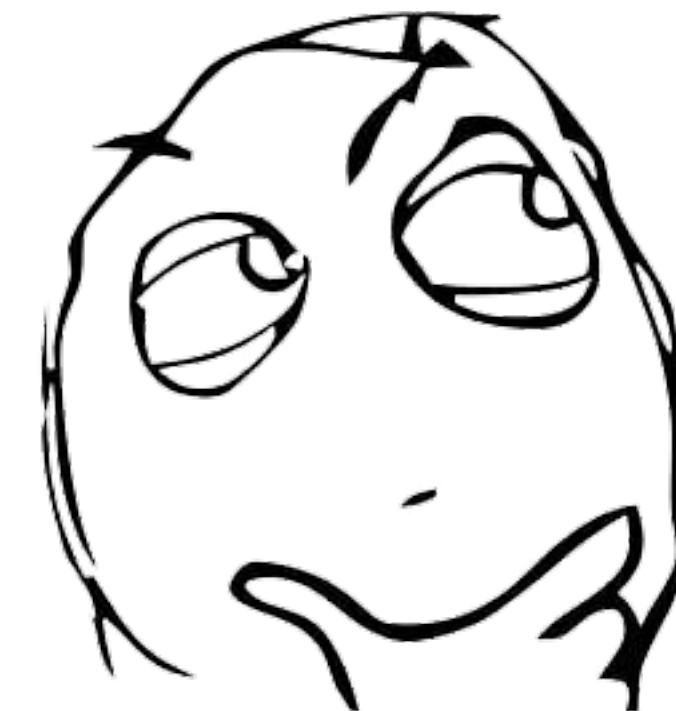


Pathological Allocation Pattern

Tuning the Garbage Collector



How should I tune the GC parameters for my application?



Garbage Collector Tuning in Pathological Allocation Pattern Applications

Nahuel Palumbo¹, Sebastian Jordan Montaño¹, Guillermo Polito¹, Pablo Tesone¹ and Stéphane Ducasse¹

¹*Univ. Lille, Inria, CNRS, Centrale Lille, UMR 9189 CRISTAL F-59000 Lille, France*

Abstract

Automatic memory management is often supported by Garbage Collectors (GC). GC usually impacts running application performance. For tuning properly, they expose some parameters to support the adaptation of their algorithms to specific applications' scenarios. In some cases, the developers should modify the GC parameter values to achieve high performance.

However, many application developers cannot be expected to perform expert analysis to determine which parameter values are the best for their application. There are techniques to find “good enough” parameter values. But, even if the overhead was reduced, it is still unknown the cause of the problem and how the GC tuning managed it.

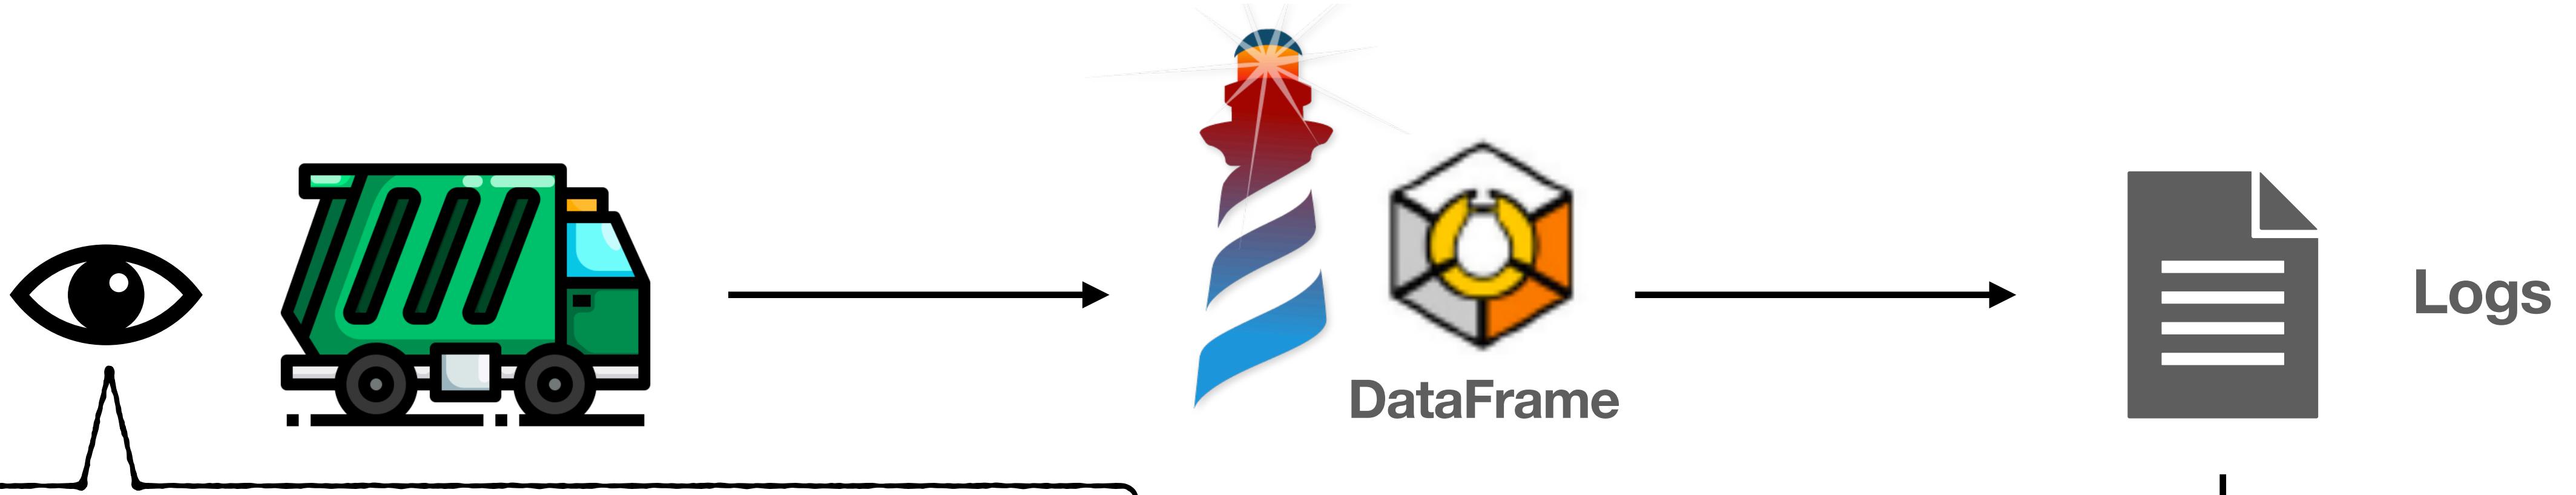
In this paper, we present a methodology to identify the causes of GC overhead in Pharo applications for tuning GC parameters. We describe the GC inside the PharoVM and its parameters, looking at how their variations change the allocation behaviour. We were able to analyse, identify and understand the GC performance issues present in one real application and suggest specific GC tuning actions. Using the suggested parameter values, we improved its performance by up to 12x and reduced GC overhead by up to 3.8x.

During the experiments, we also found: 1) a bug in the production PharoVM concerning the tenuring policy, 2) a misconception about one GC parameter even for the VM developers, and 3) some possible improvements for the current GC implementation.



Our methodology for GC tuning

Profile GC events



From Scavenges:

- Amount of memory used (before and after).
- Size of the Remembered Set (before and after).
- Tenuring info (amount of data - threshold).
- Executed time.

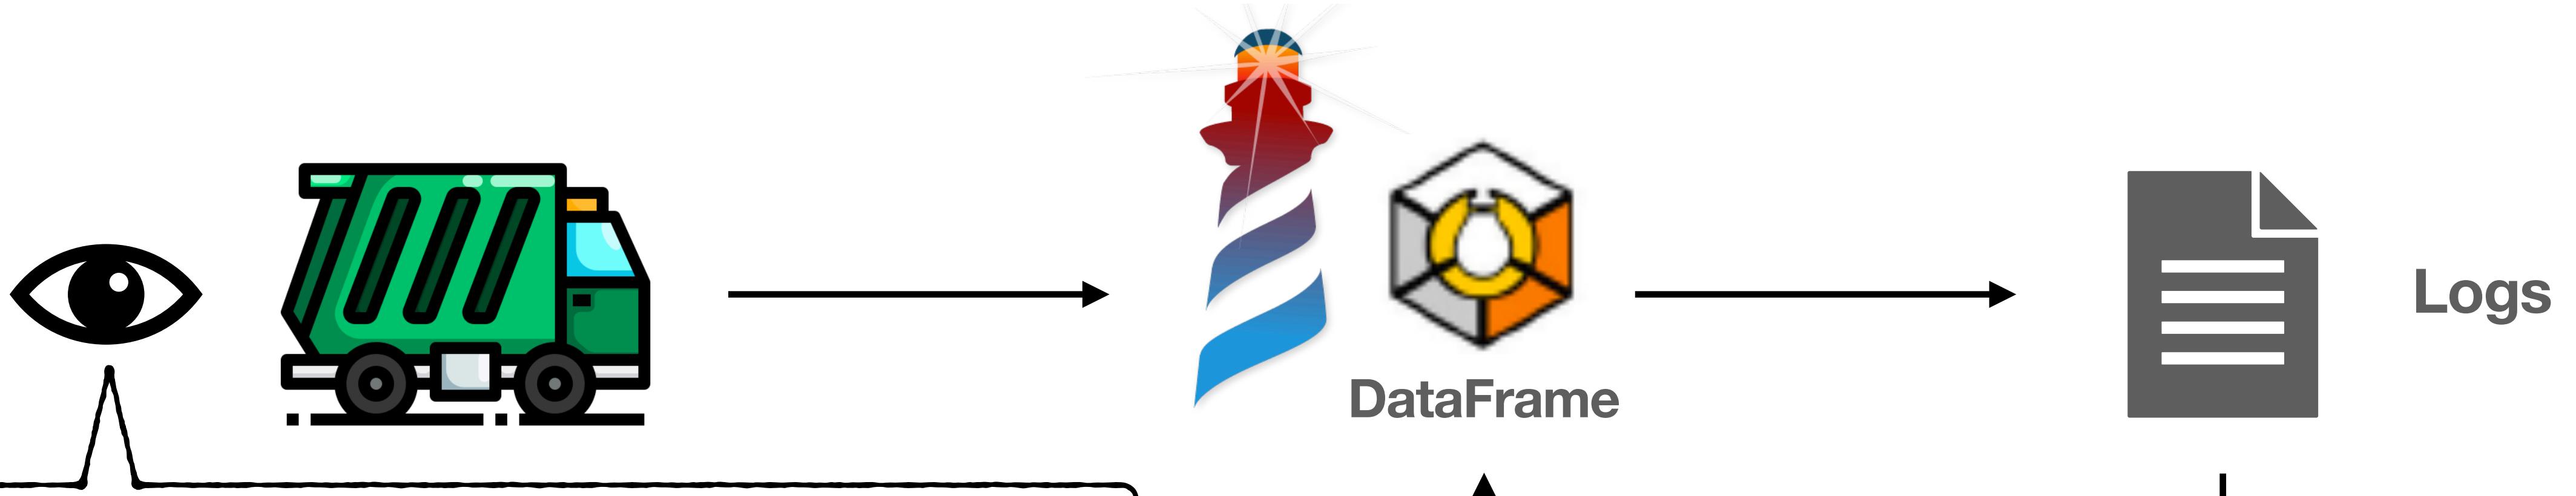
From FullGC:

- Time spent marking/sweeping/compacting.
- Executed time.



Our methodology for GC tuning

Profile GC events



From Scavenges:

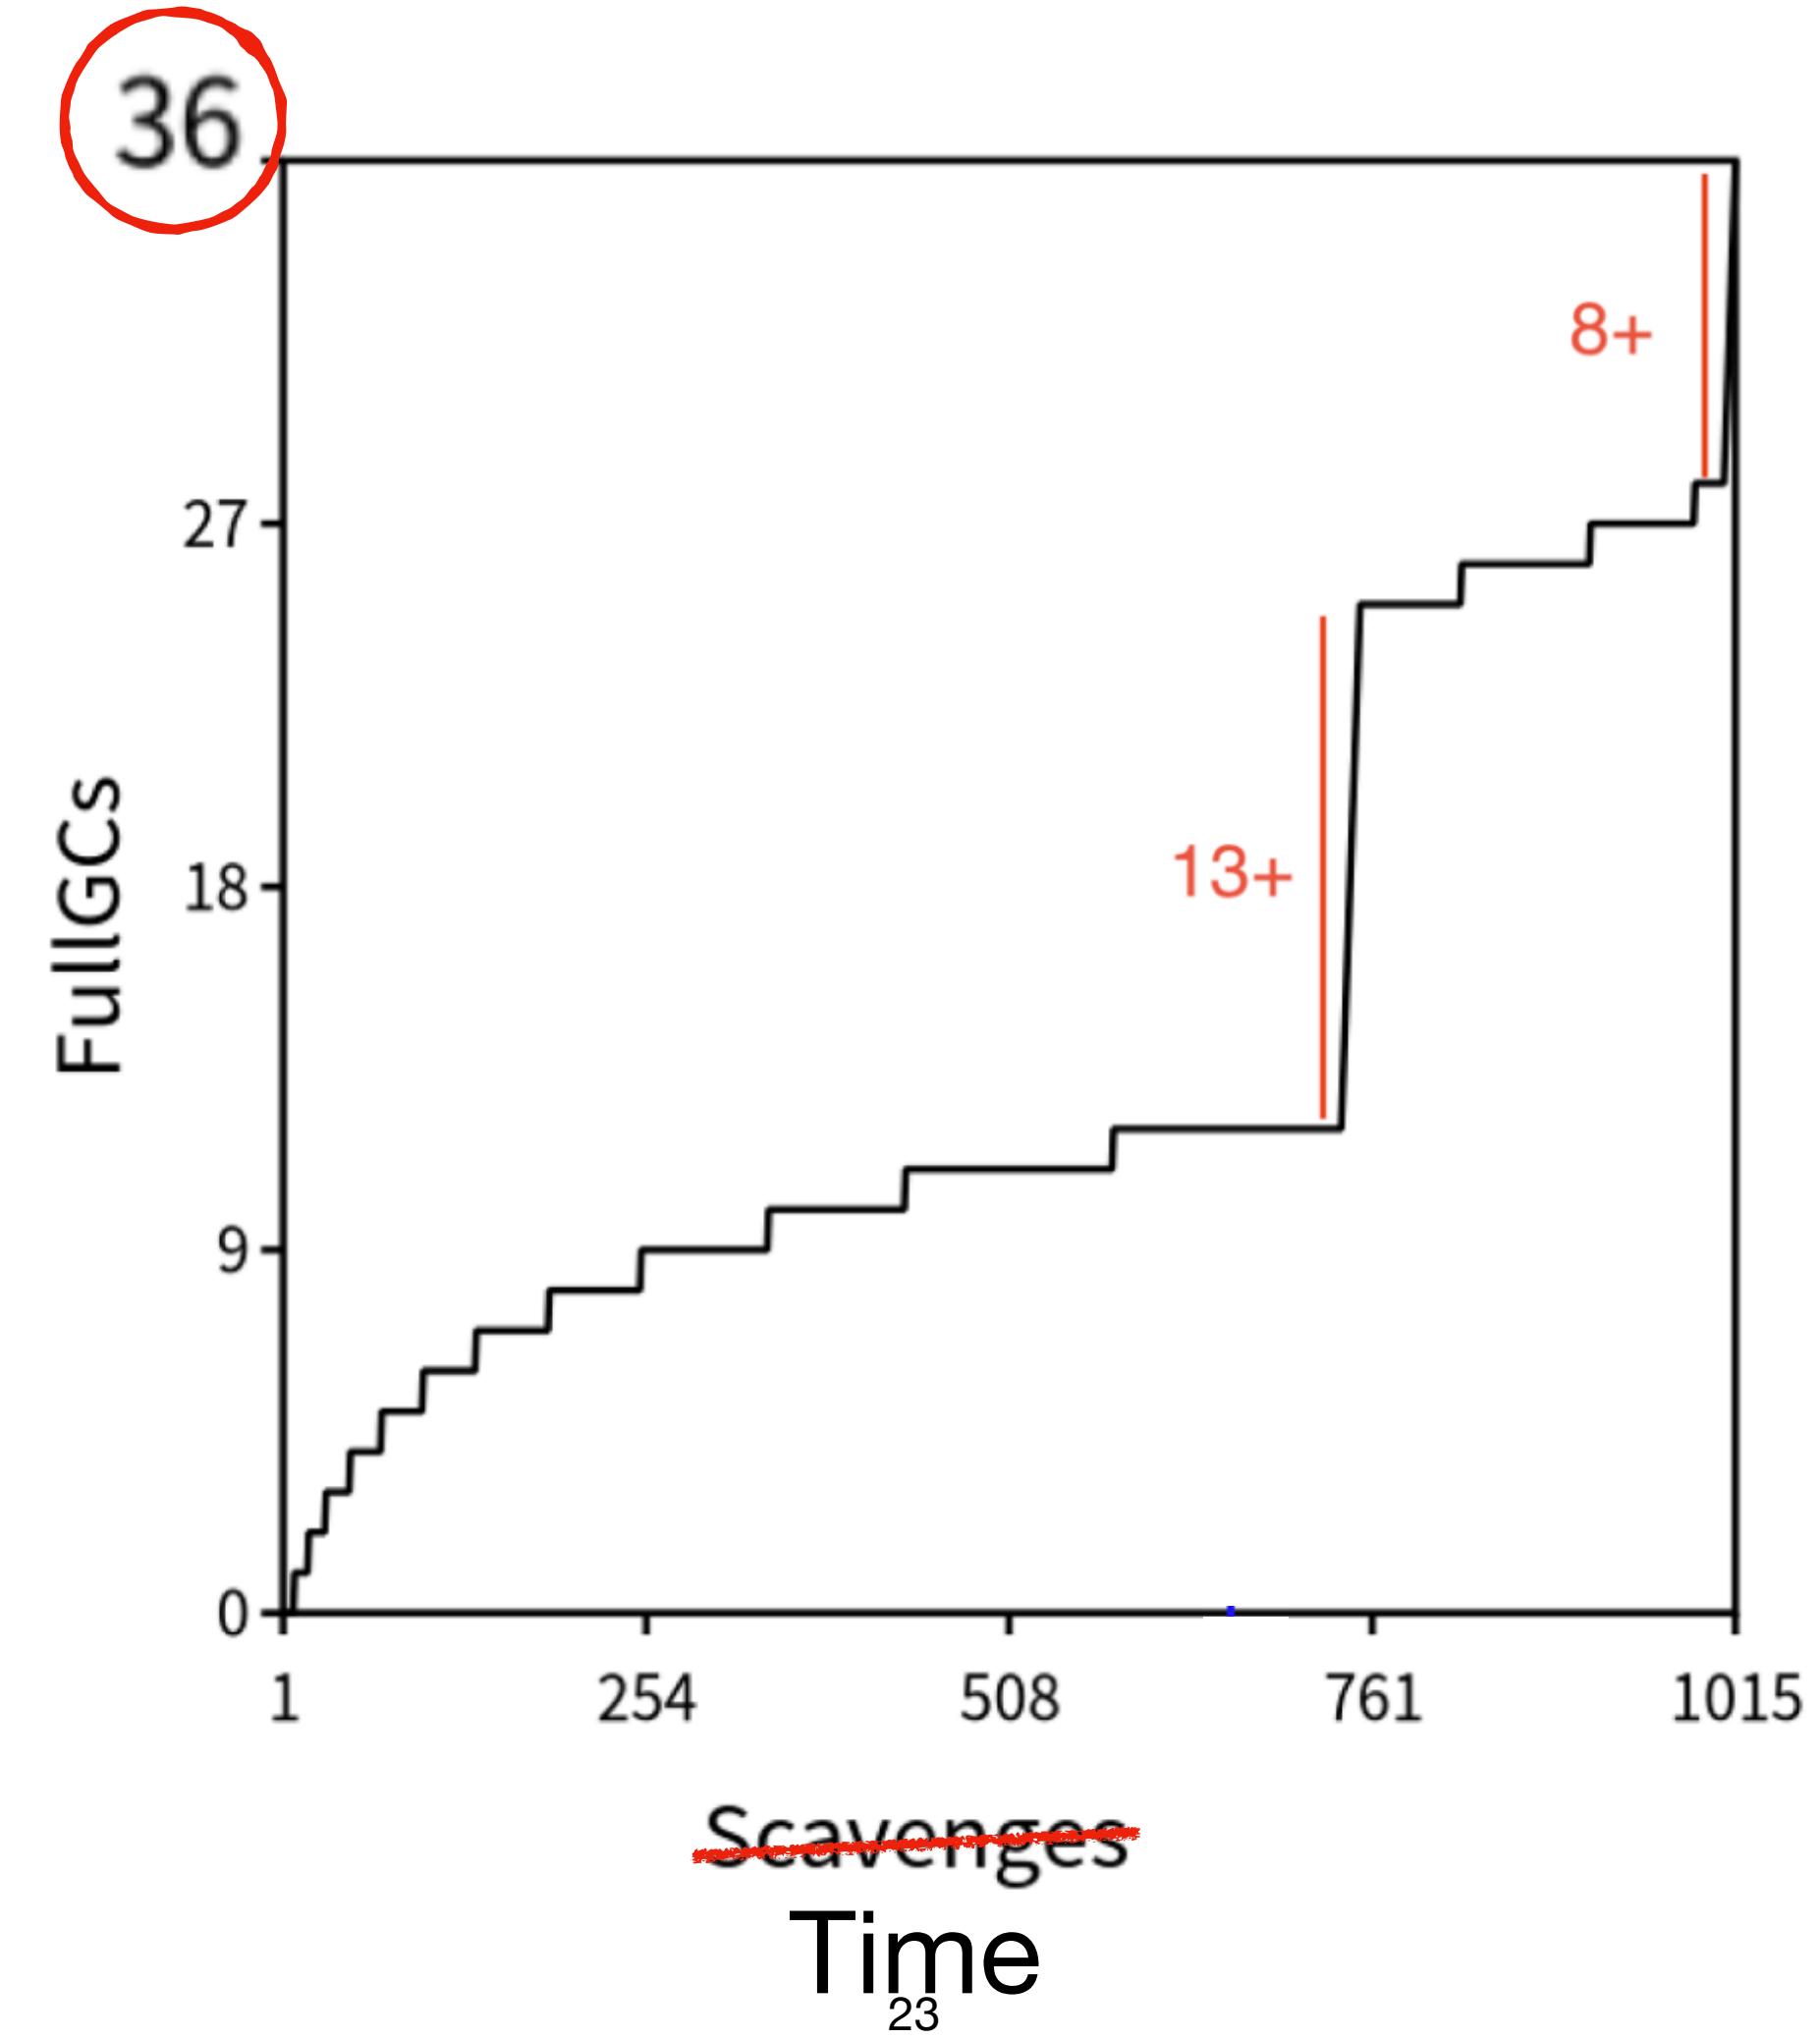
- Amount of memory used (before and after).
- Size of the Remembered Set (before and after).
- Tenuring info (amount of data - threshold).
- Executed time.

From FullGC:

- Time spent marking/sweeping/compacting.
- Executed time.

How Memory Grows

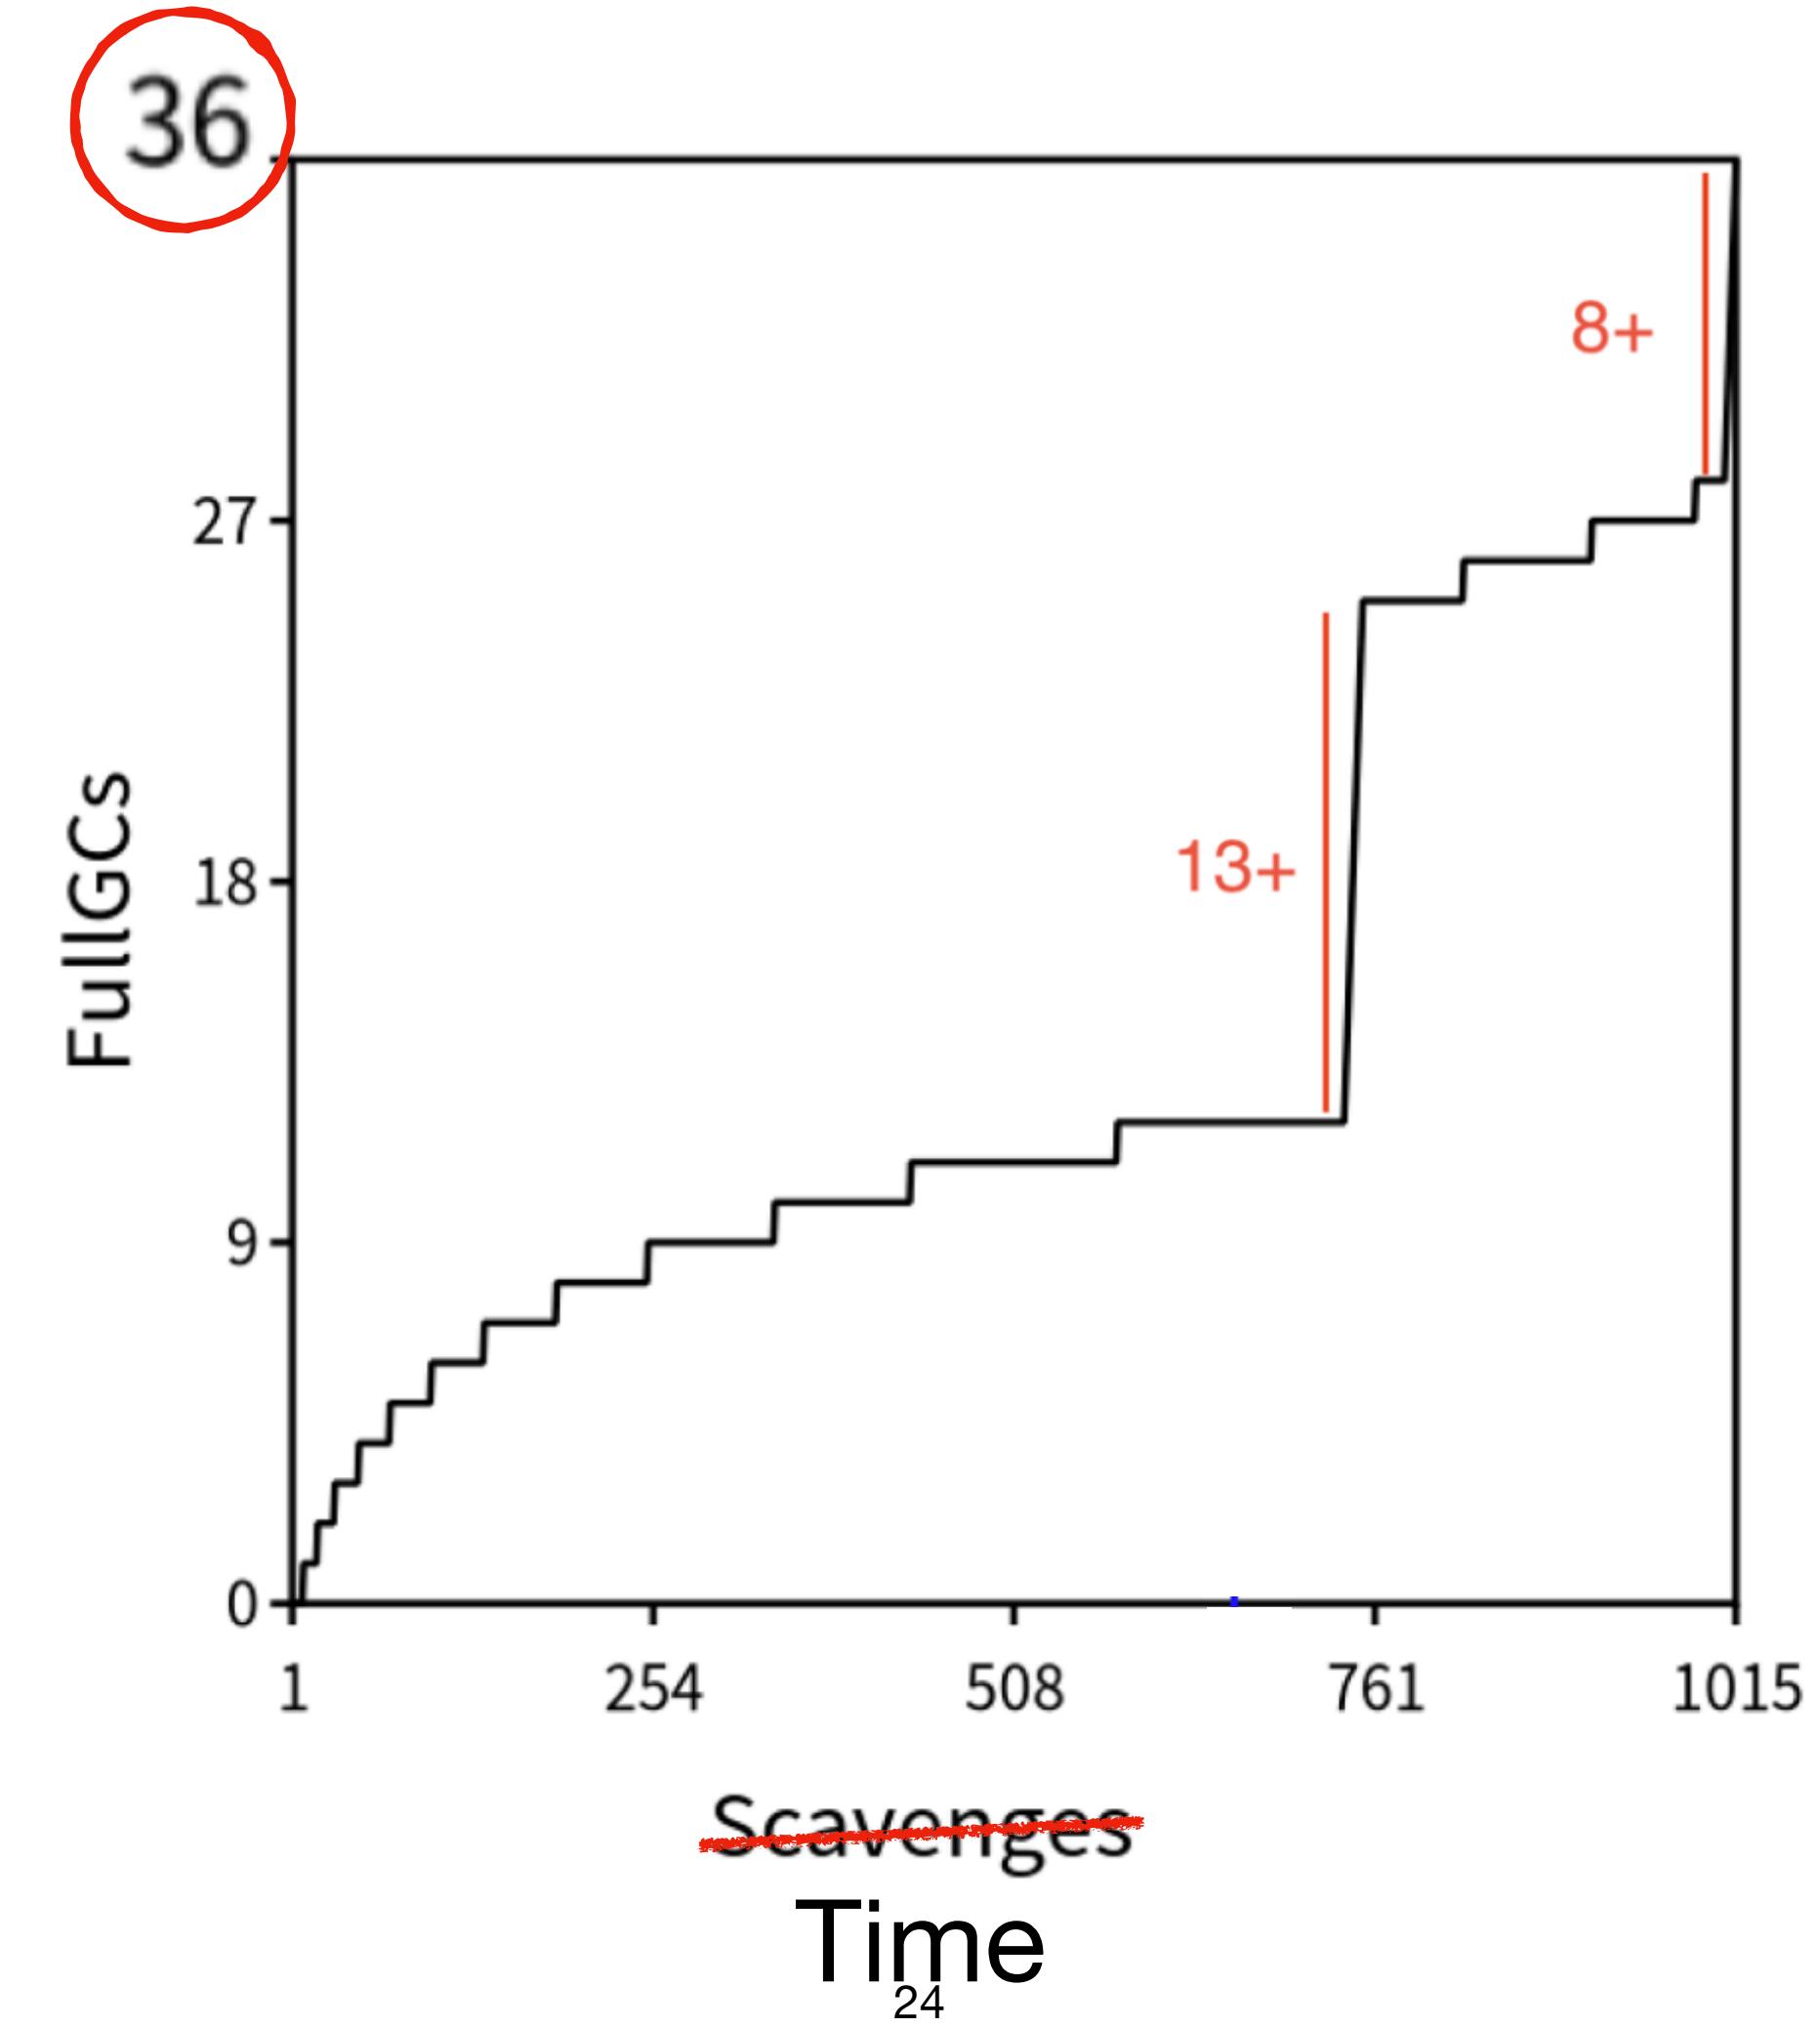
The overhead



How Memory Grows

The overhead

I run some FullGCs
when memory grows
so much

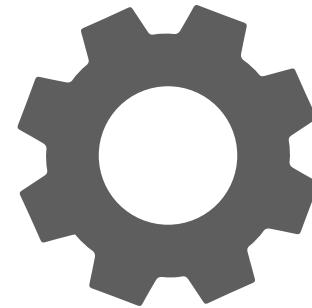


Don't do that

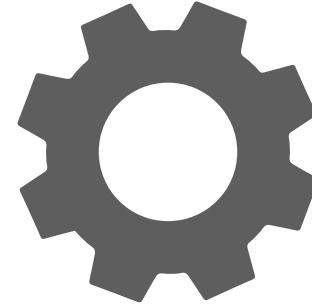


How Memory Grows

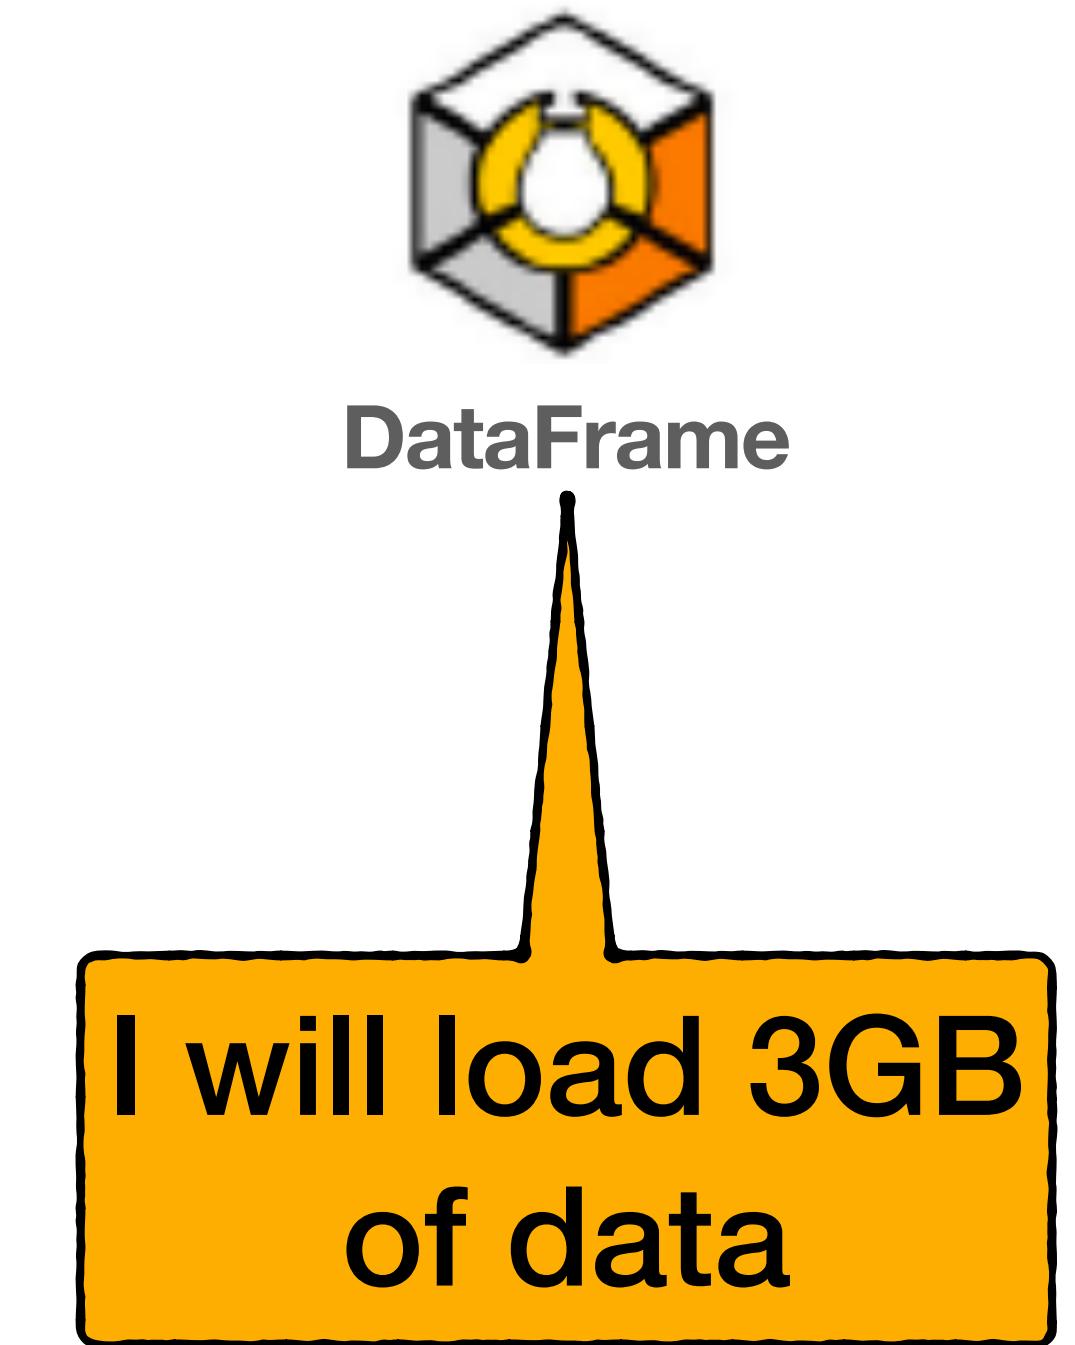
The tuning solution



FullGC Ratio - Threshold for triggering a FullGC when the old space grows more than expected

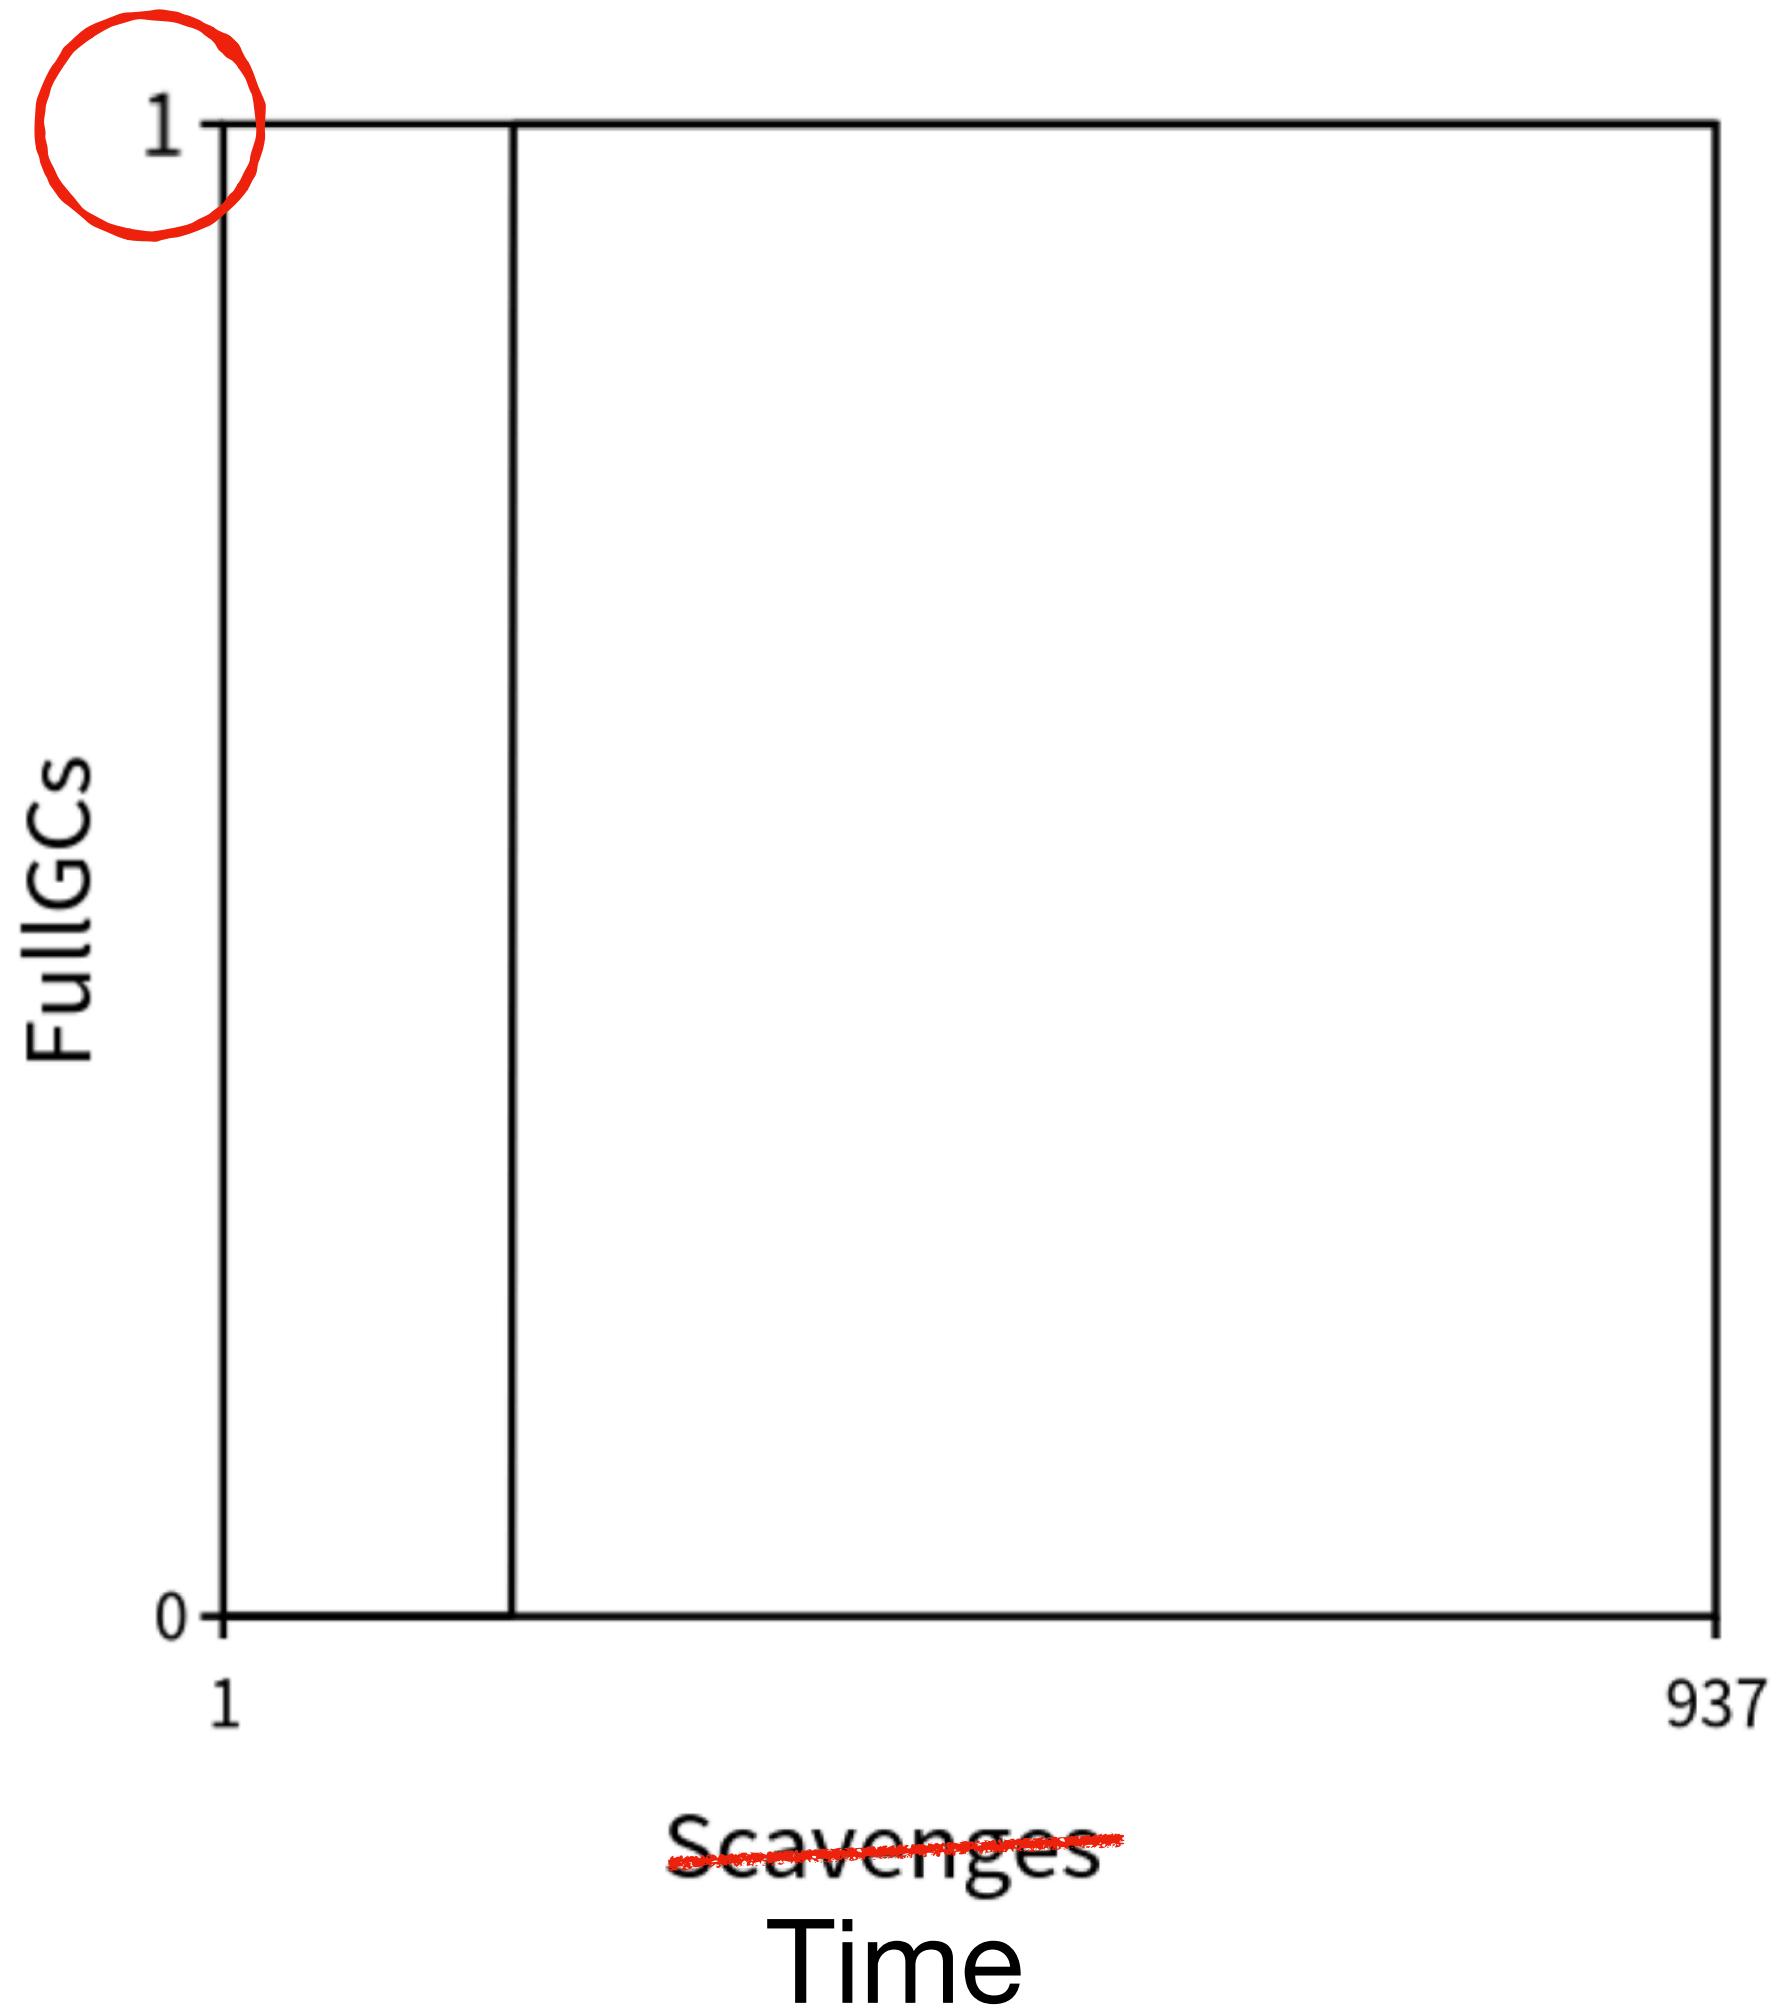


Grow Headroom - Minimum amount of memory that the GC will order from the OS



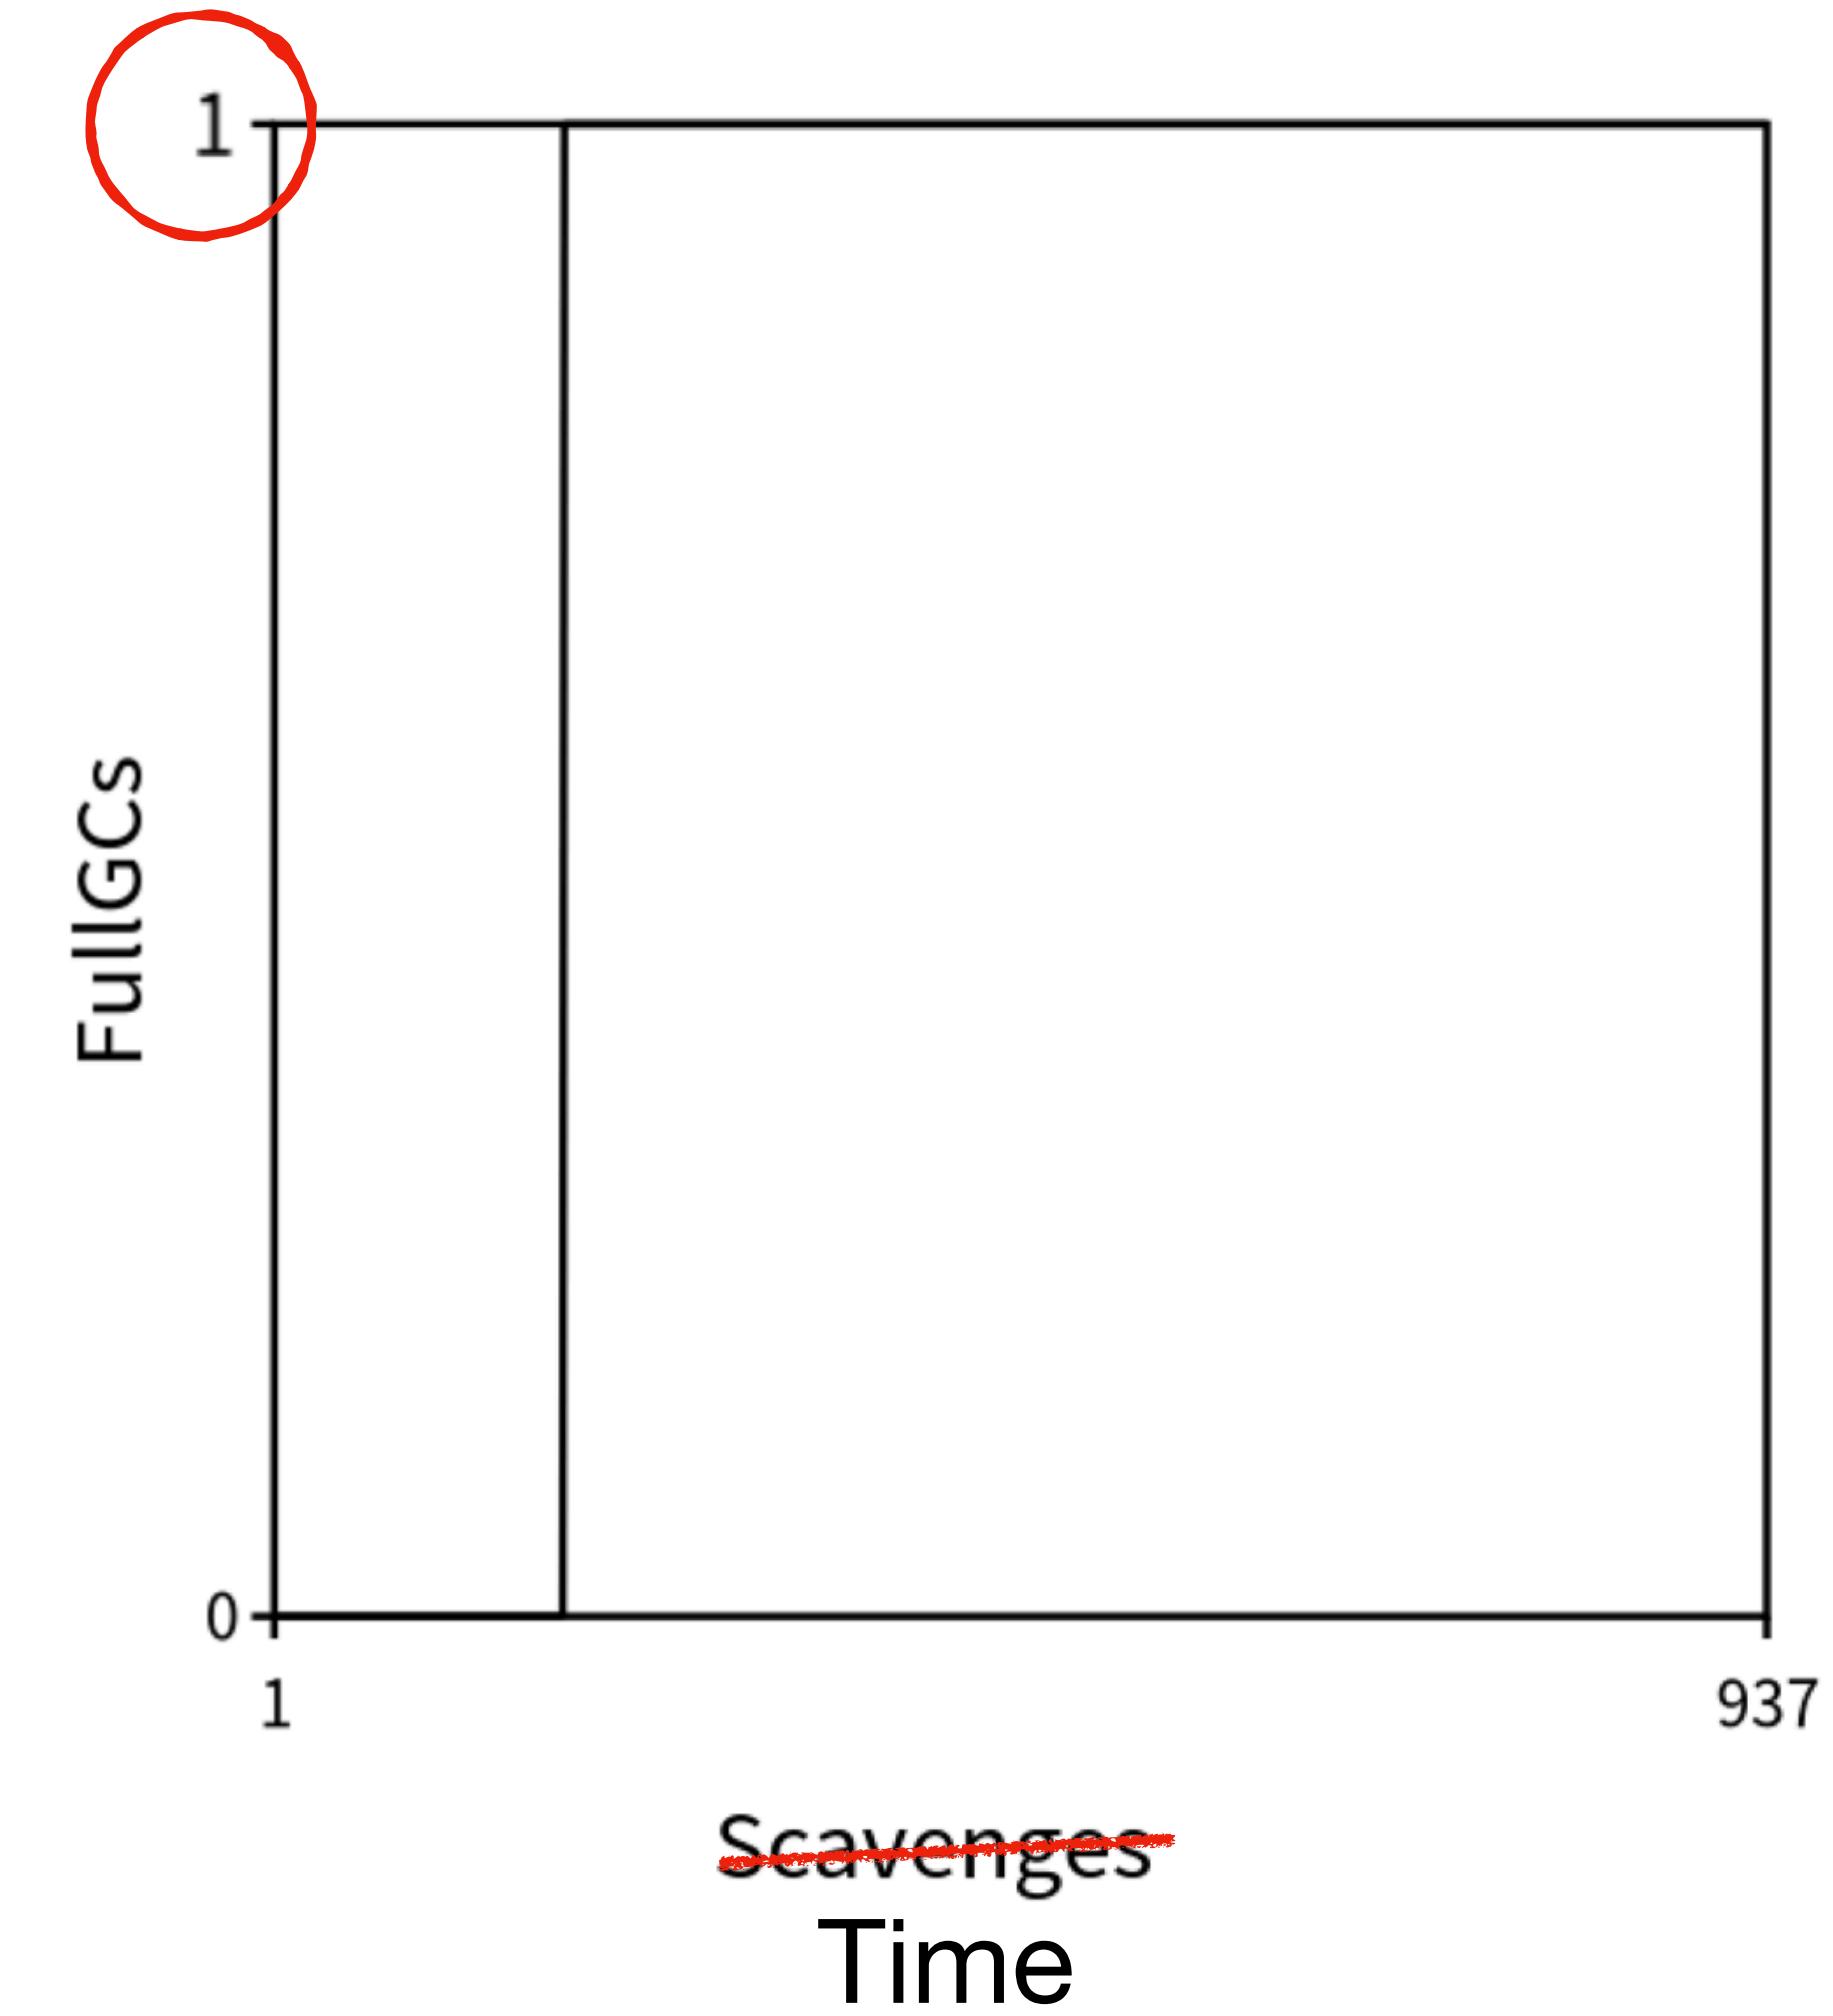
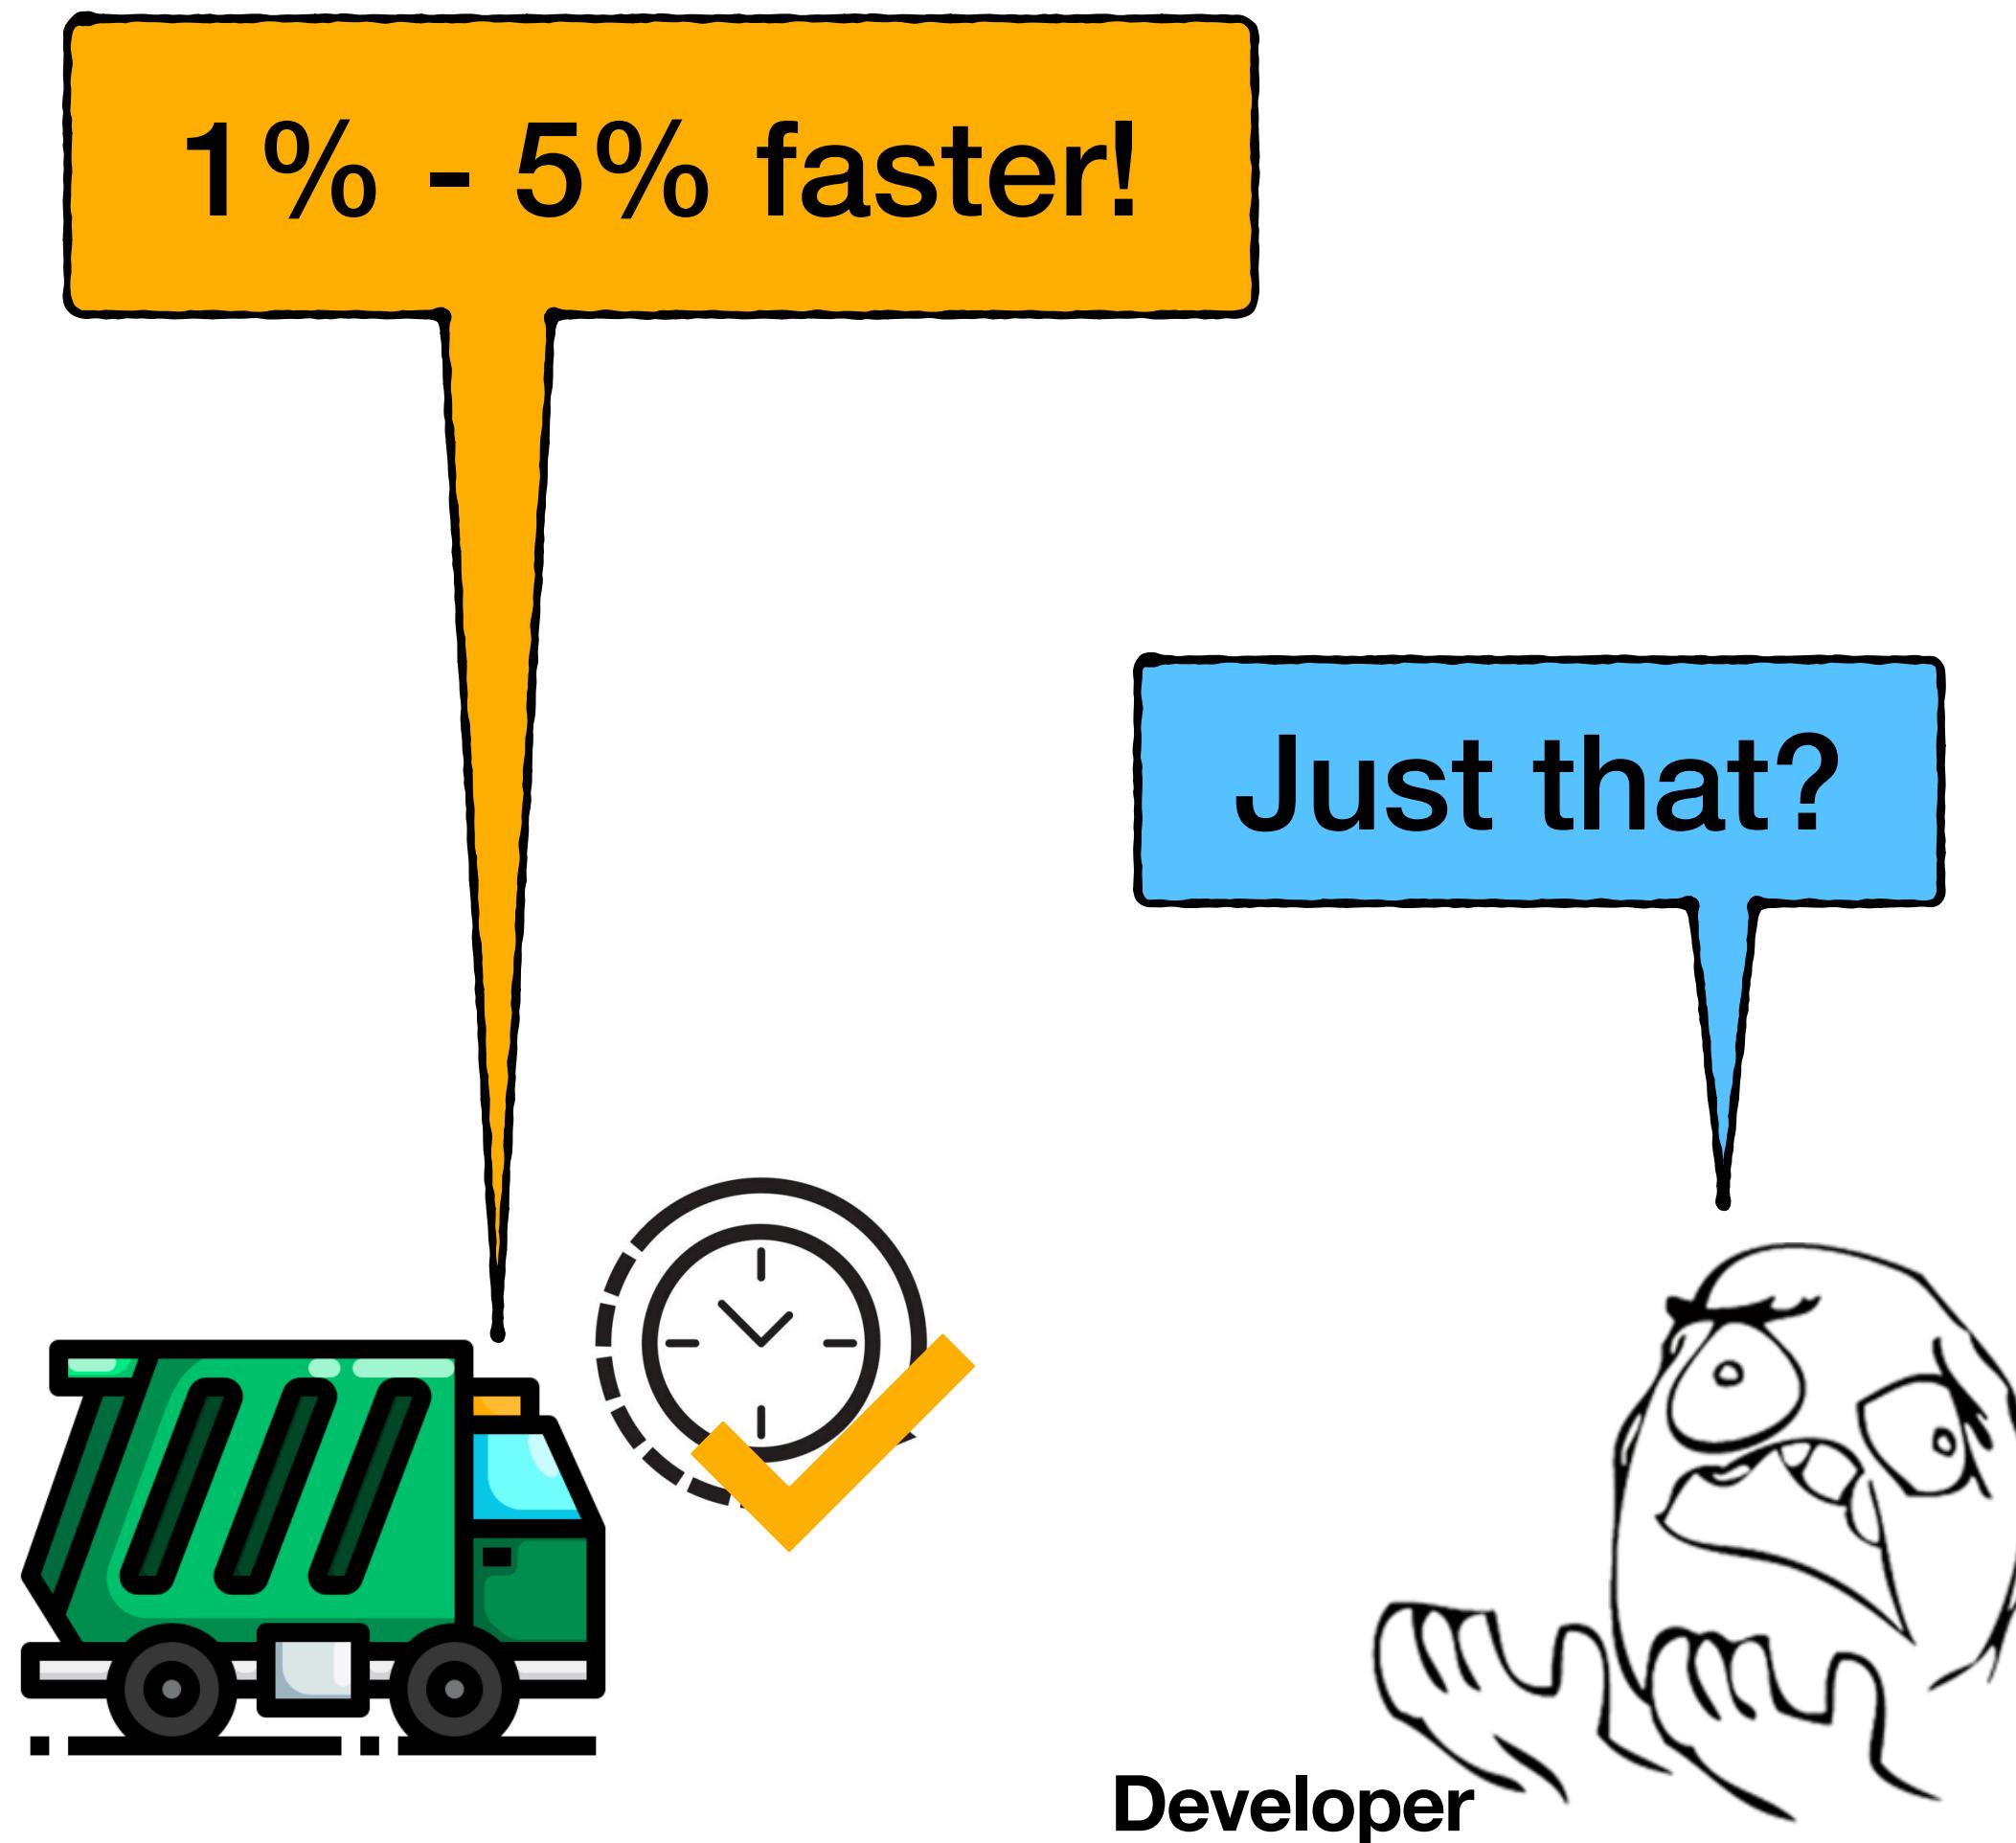
How Memory Grows

The tuning solution



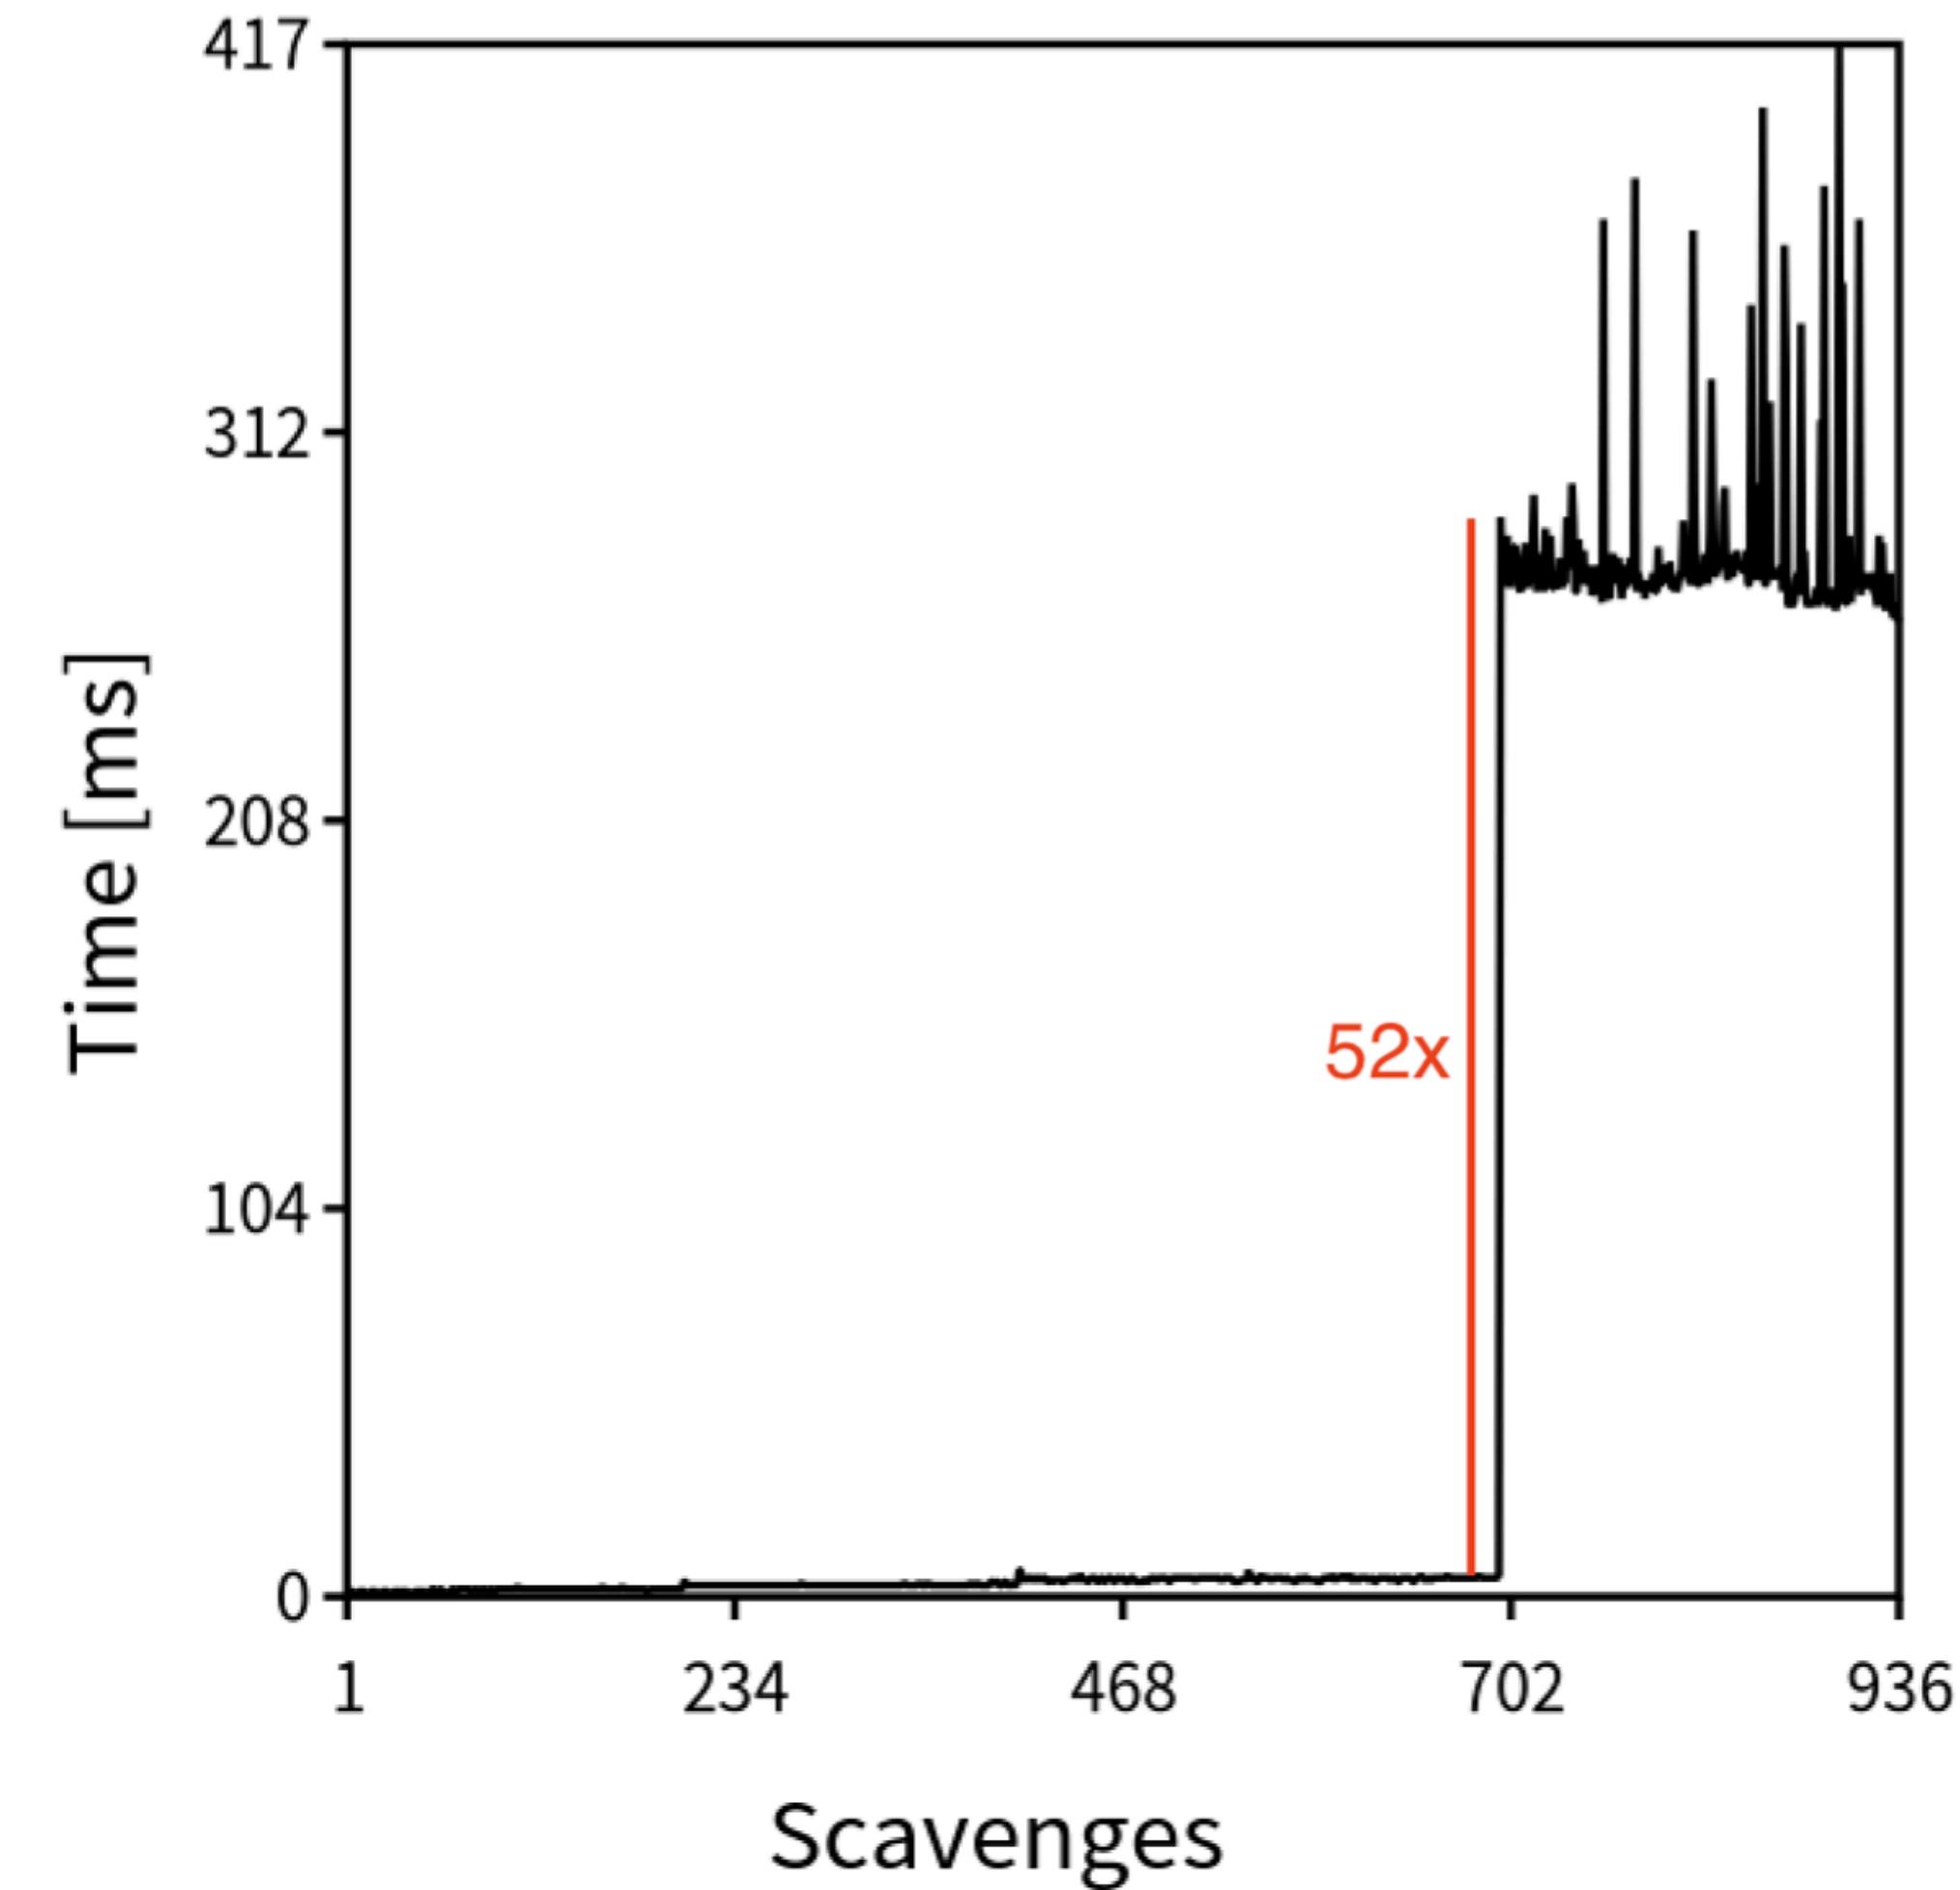
How Memory Grows

The tuning solution



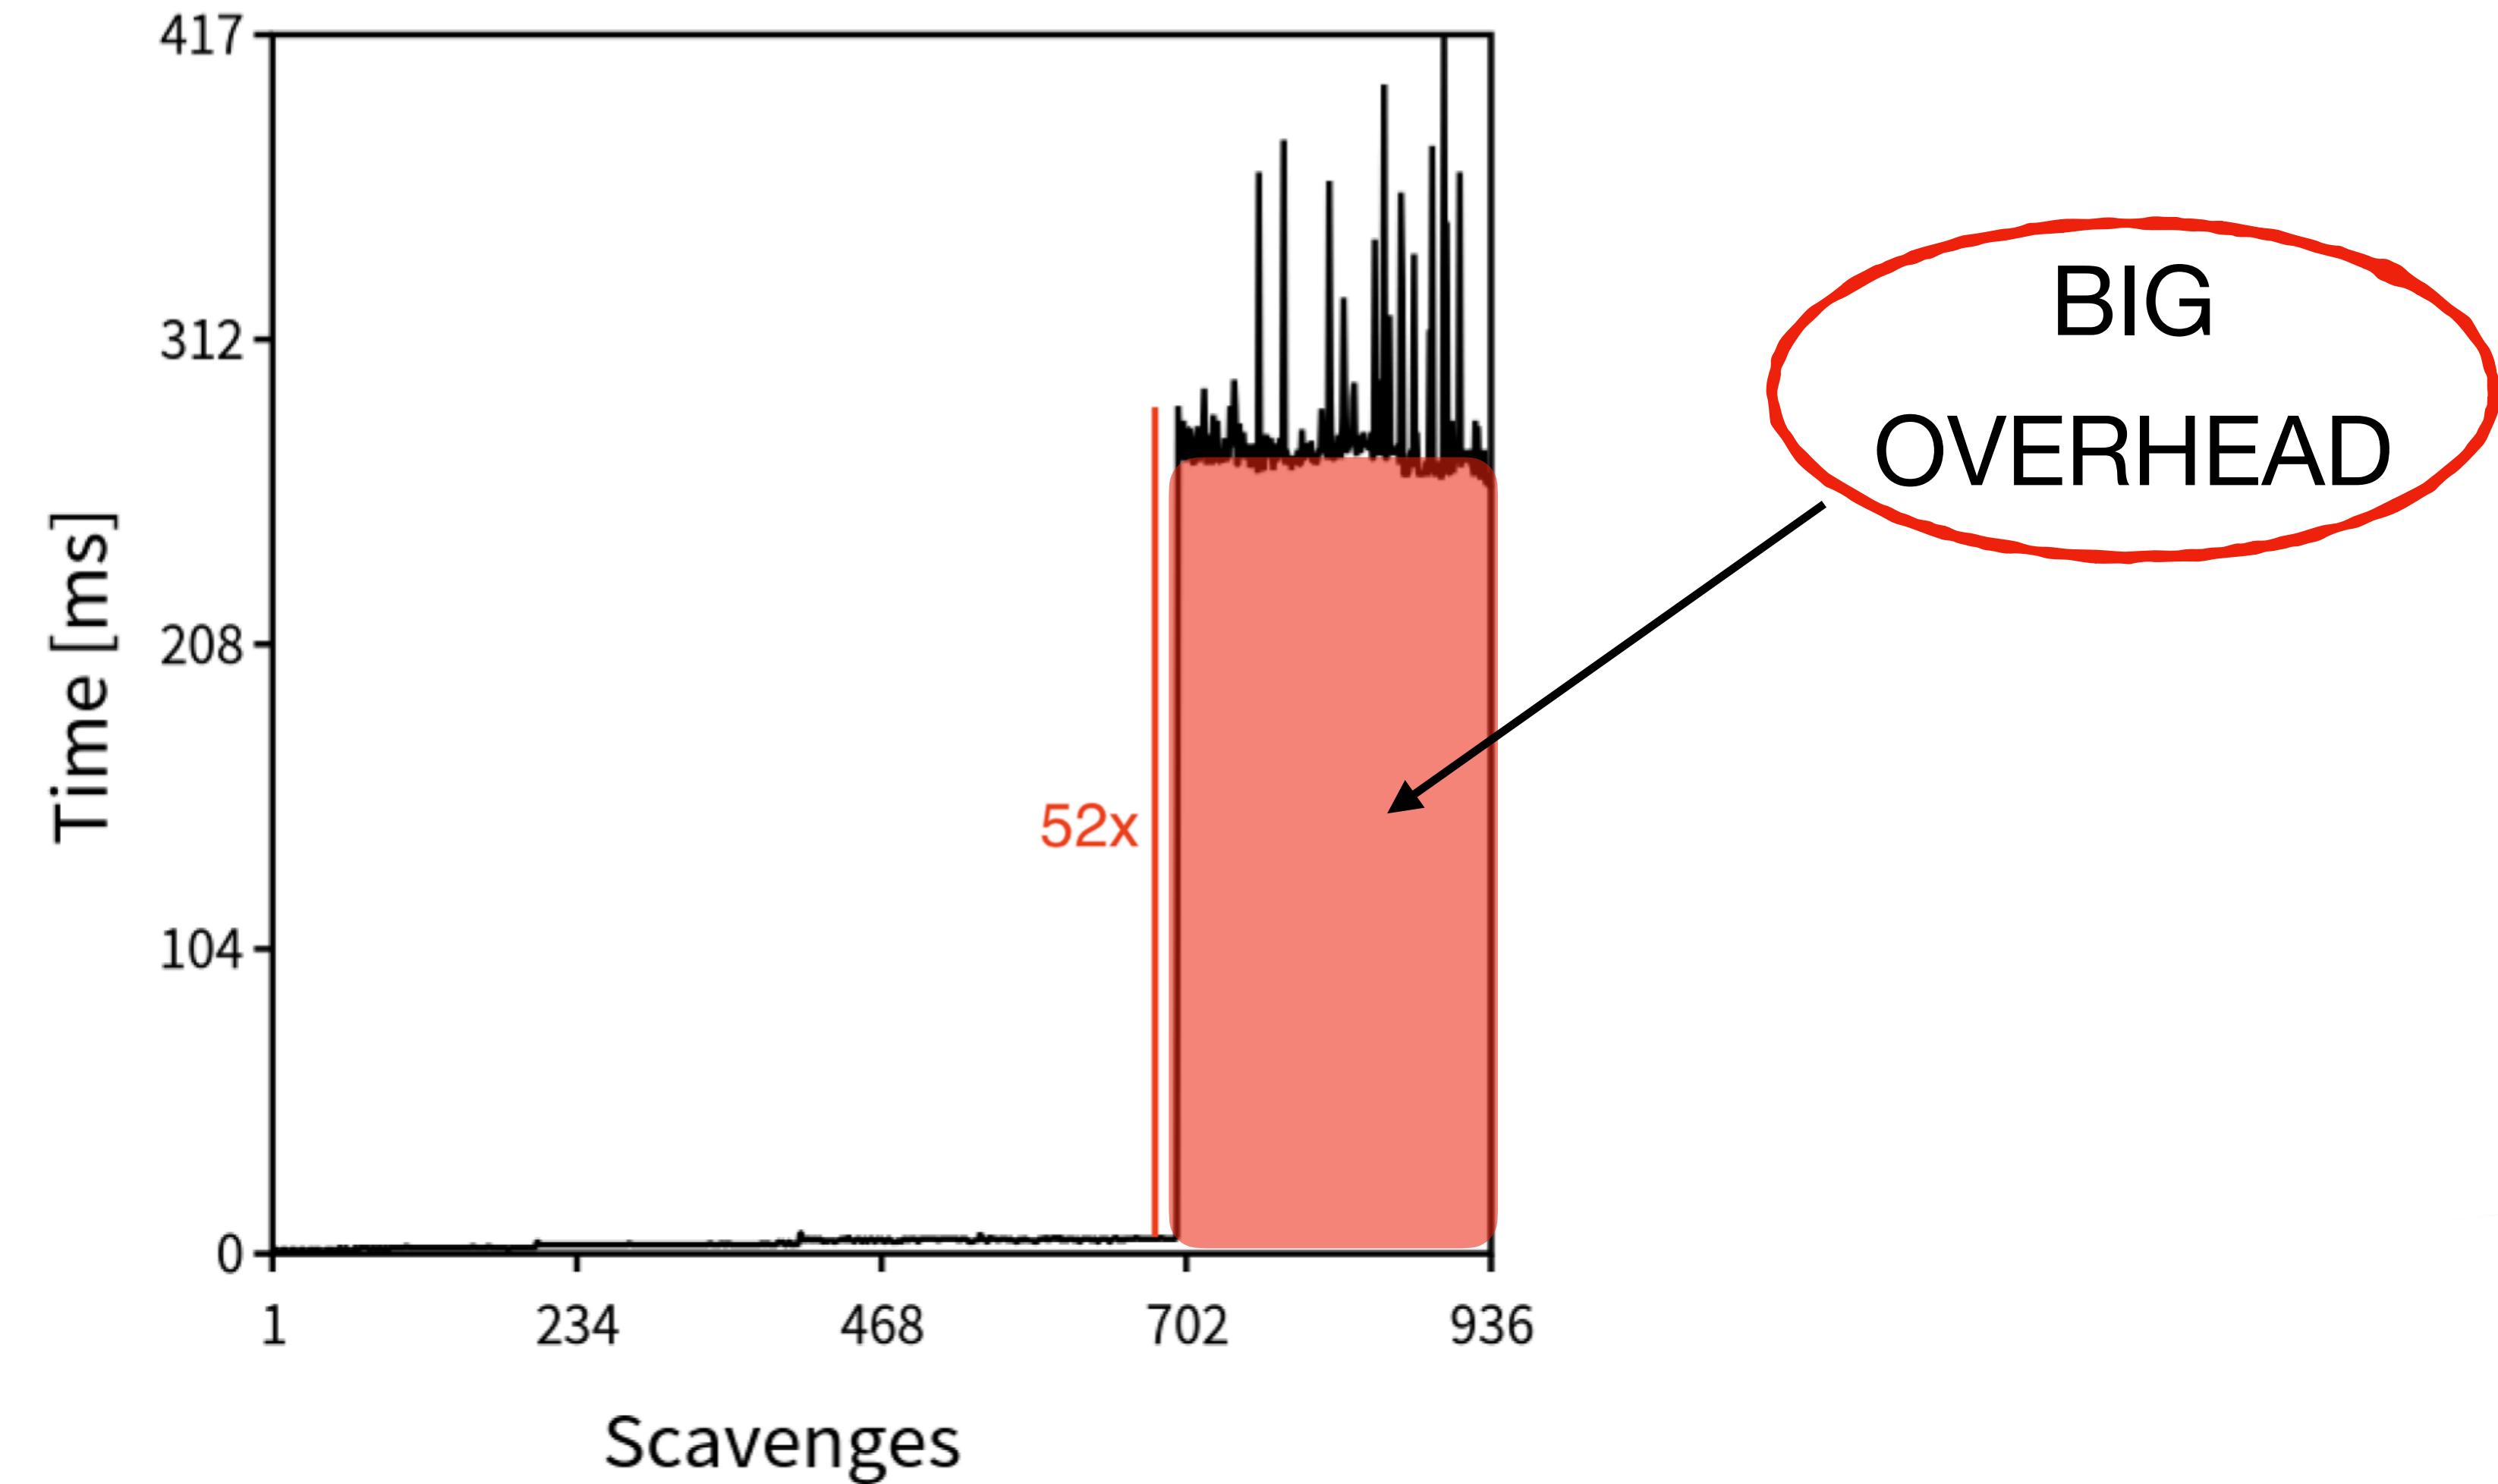
Deeper in the allocation pattern

Generational clash



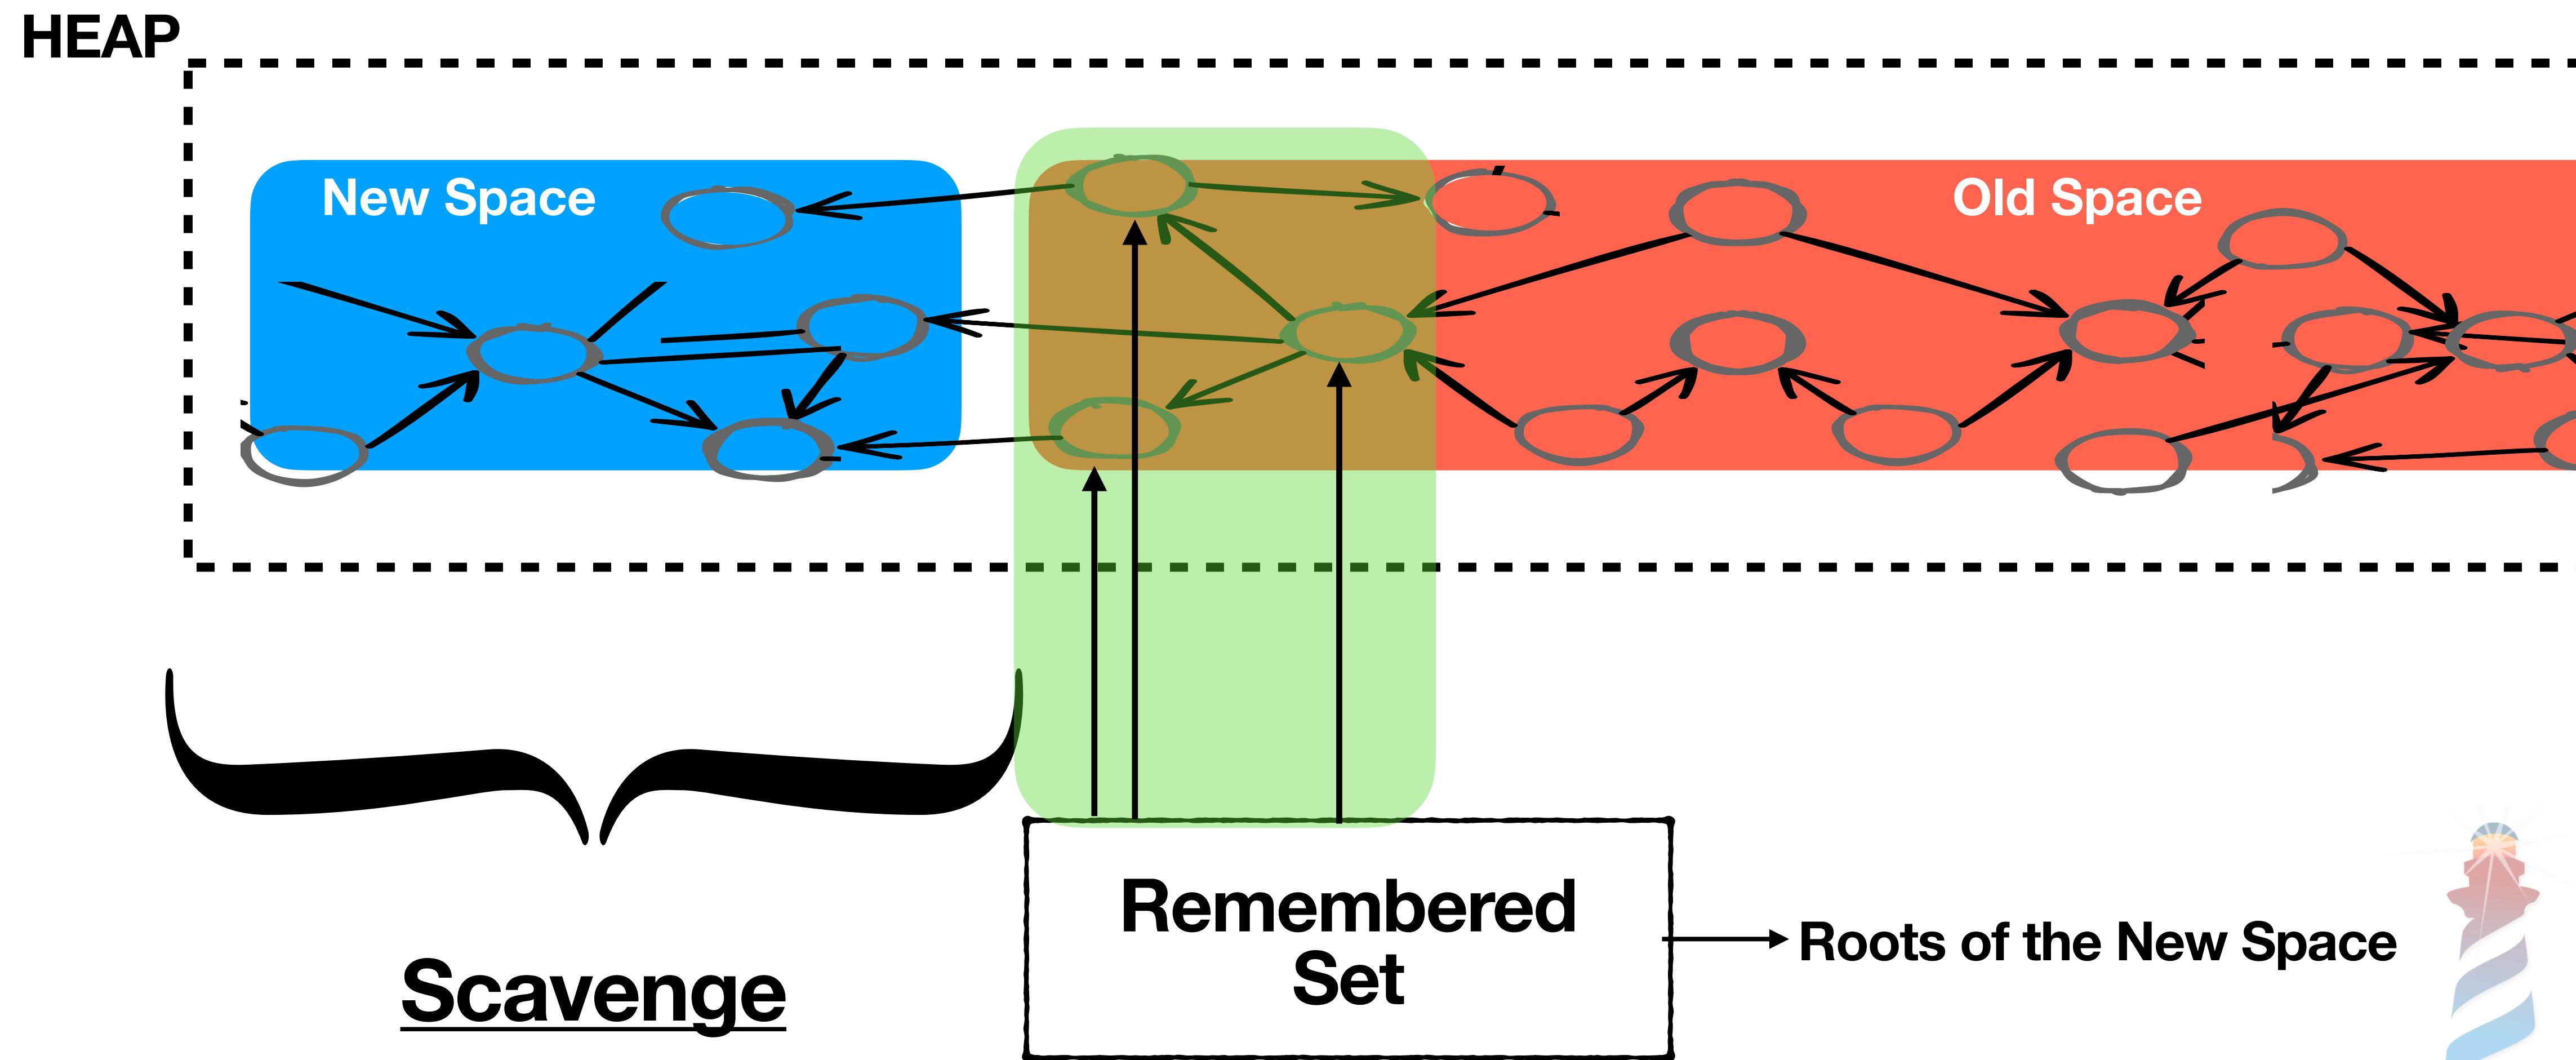
Deeper in the allocation pattern

Generational clash



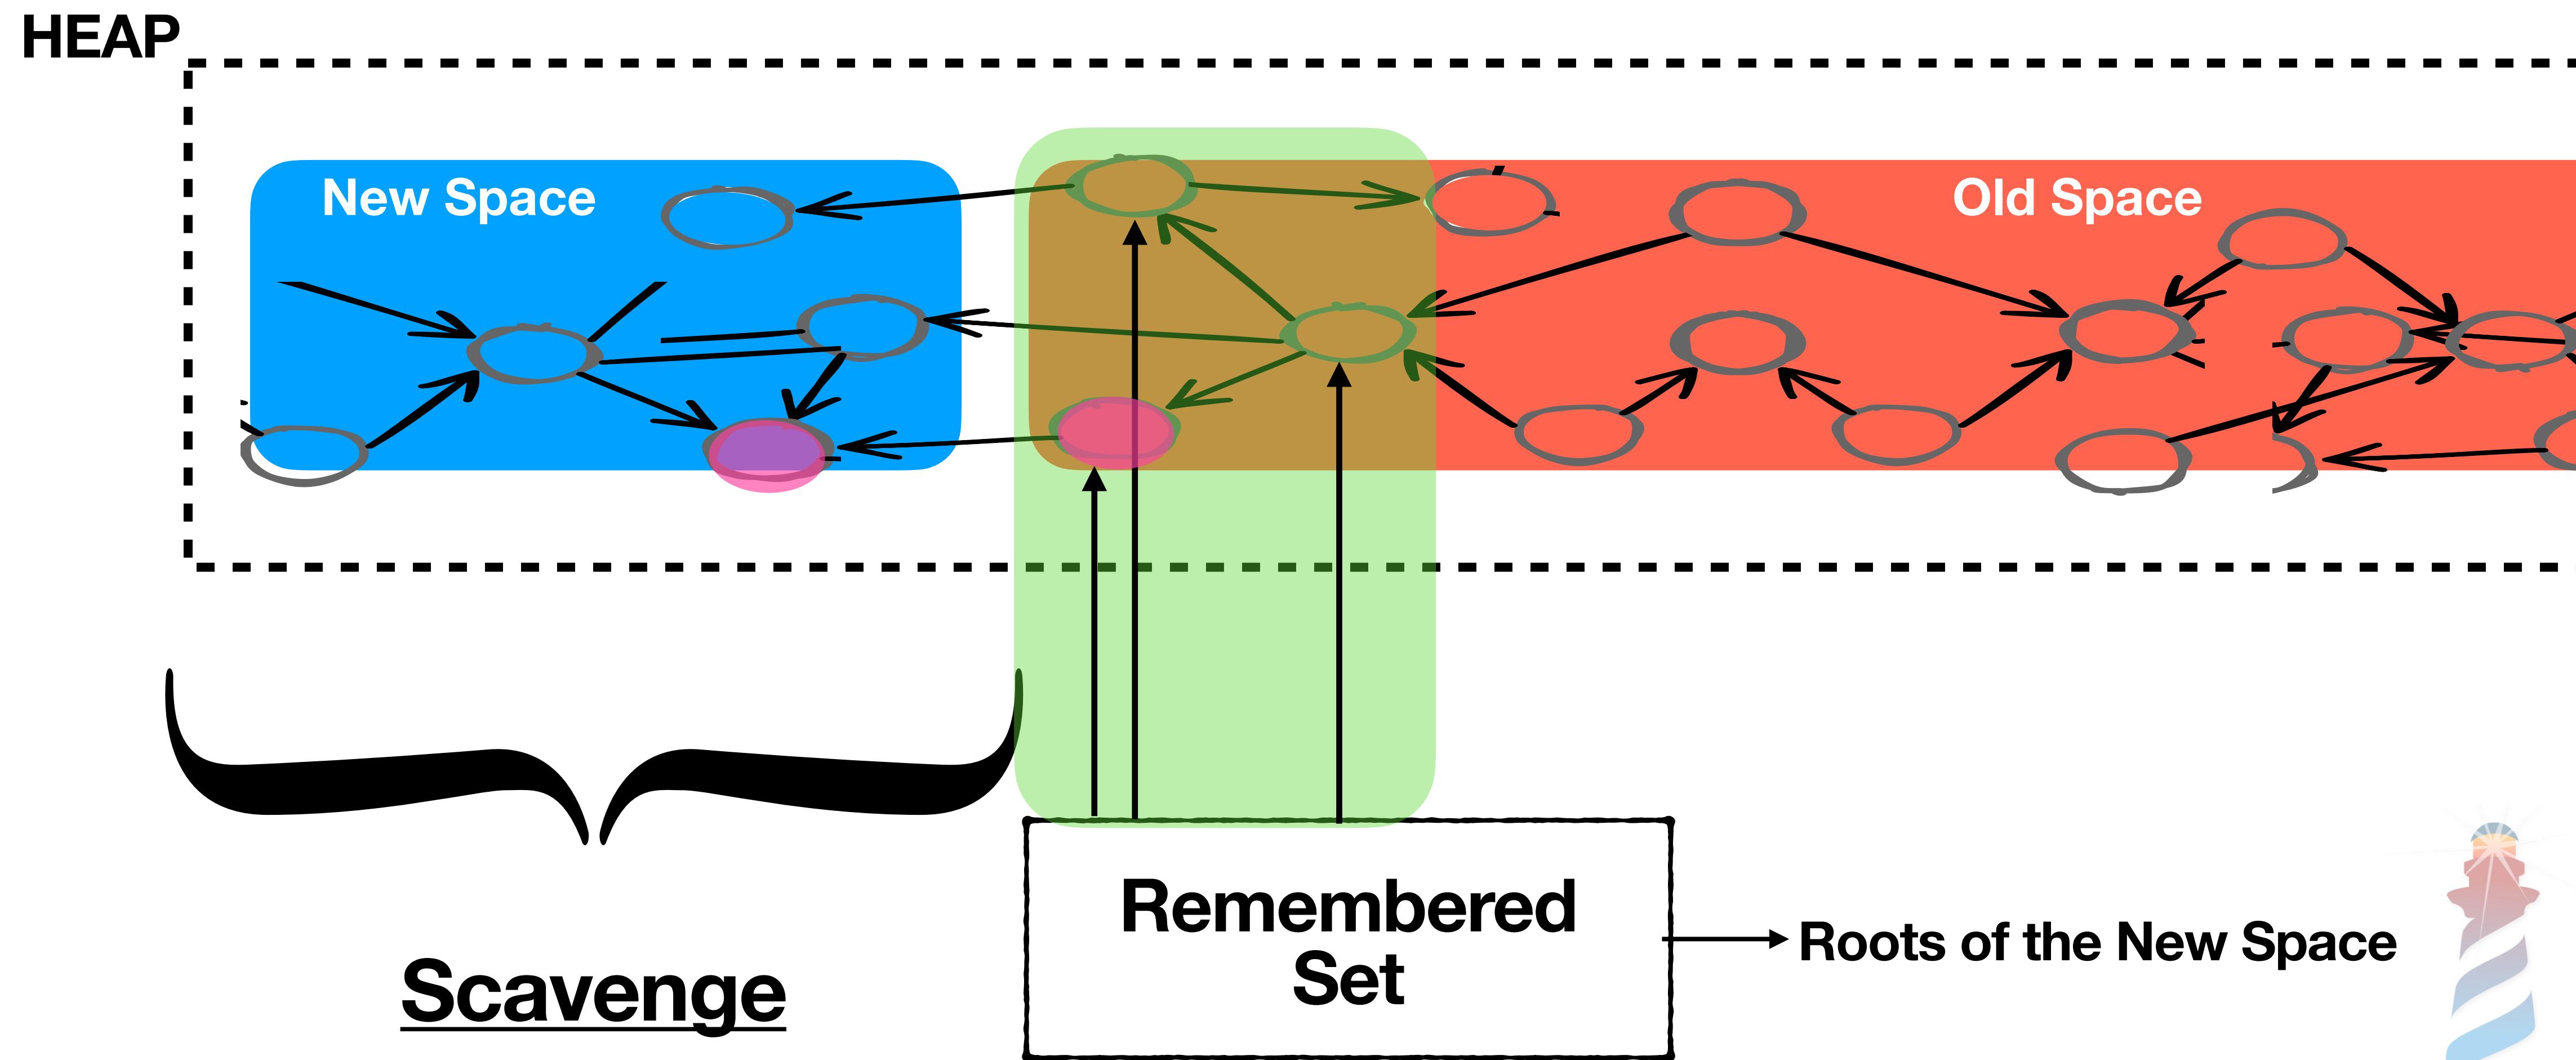
Deeper in the allocation pattern

Generational clash



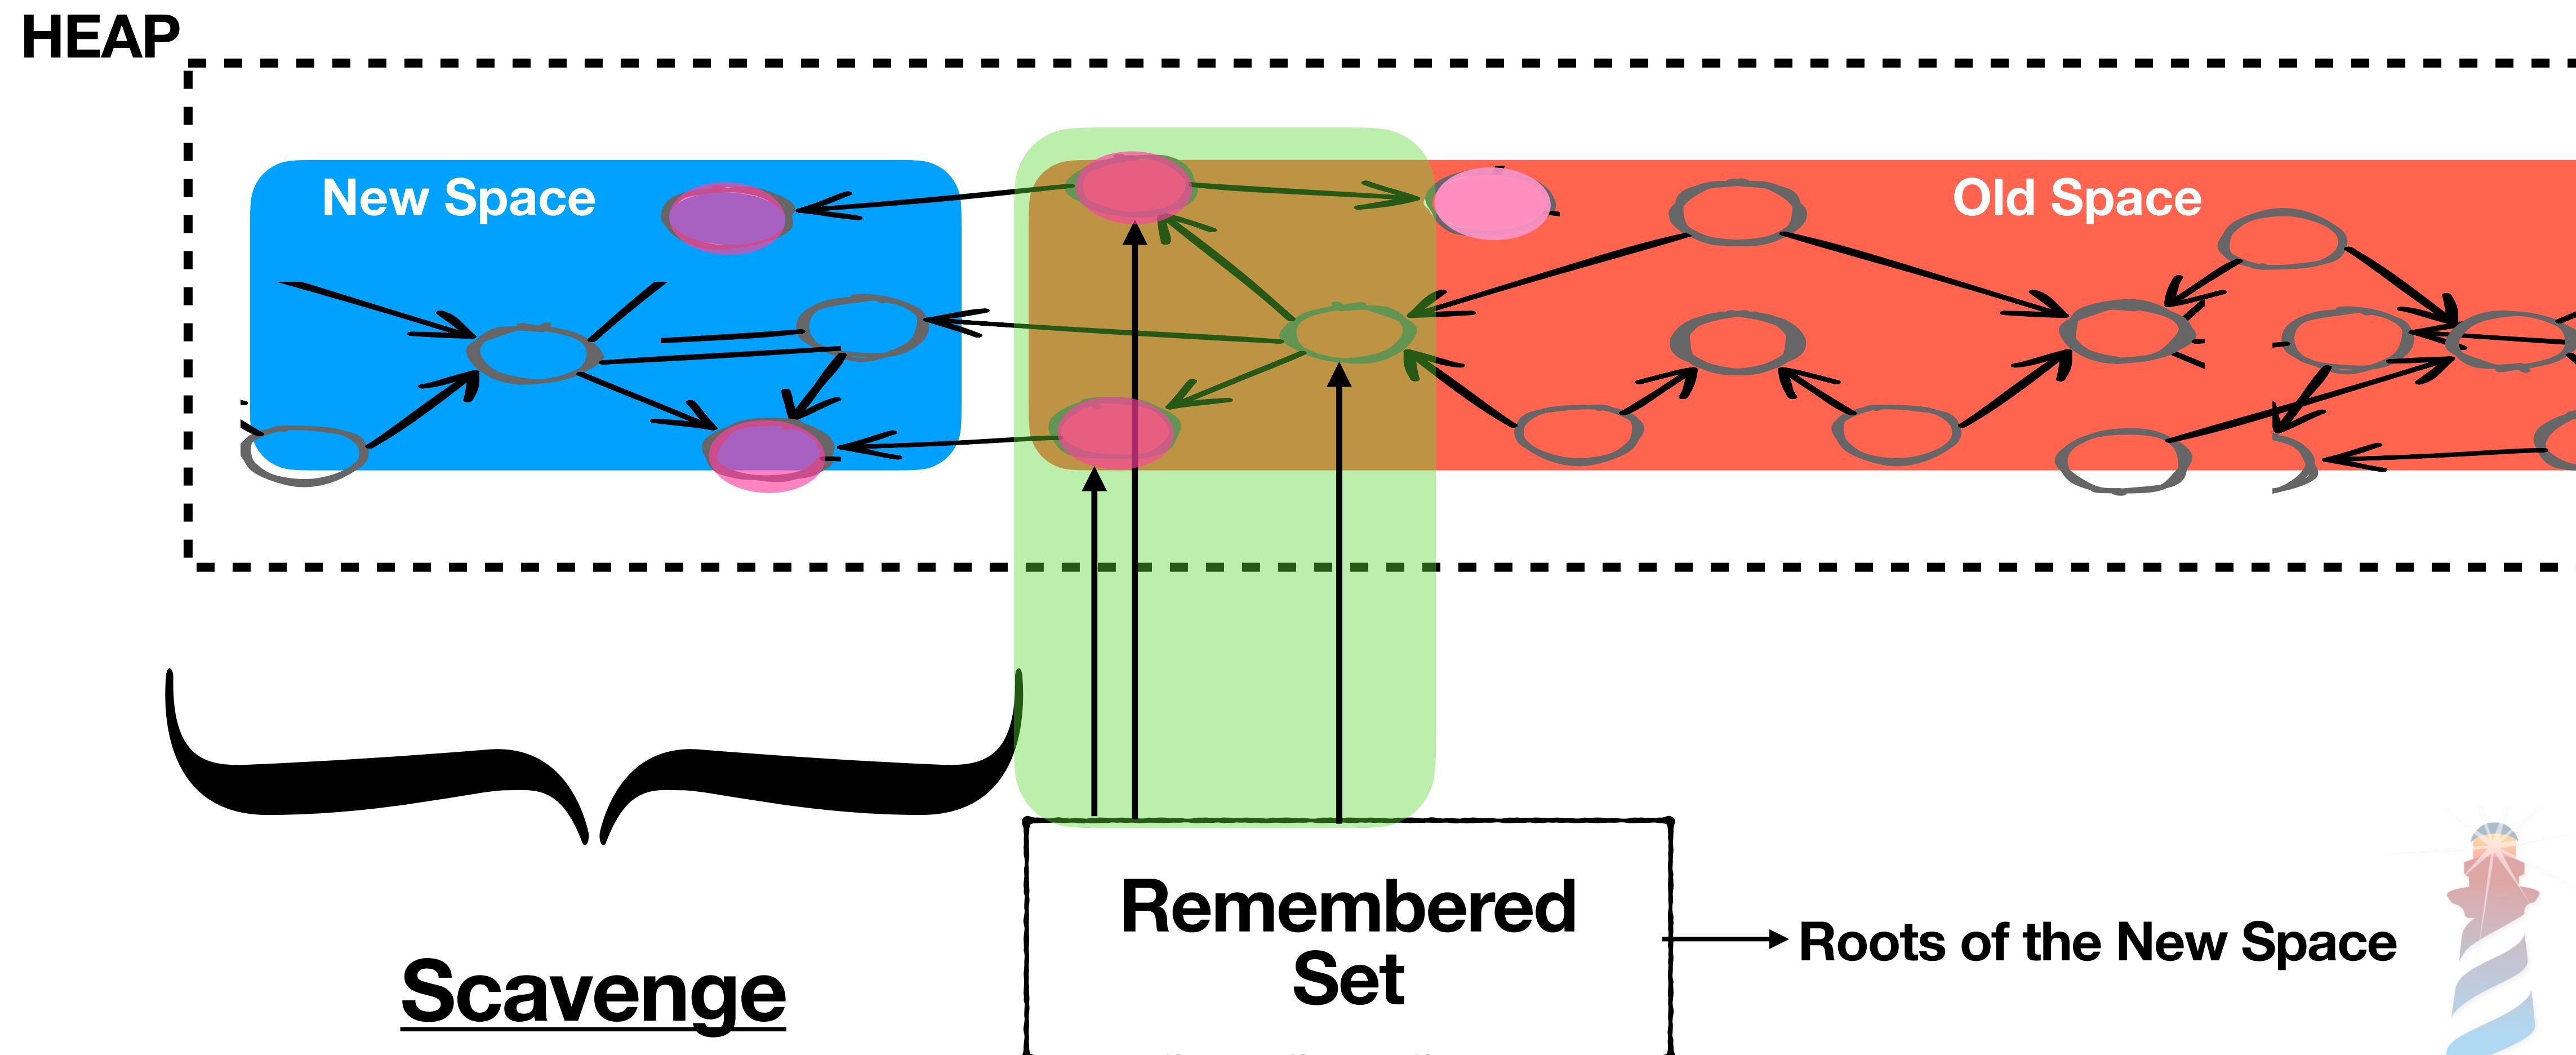
Deeper in the allocation pattern

Generational clash



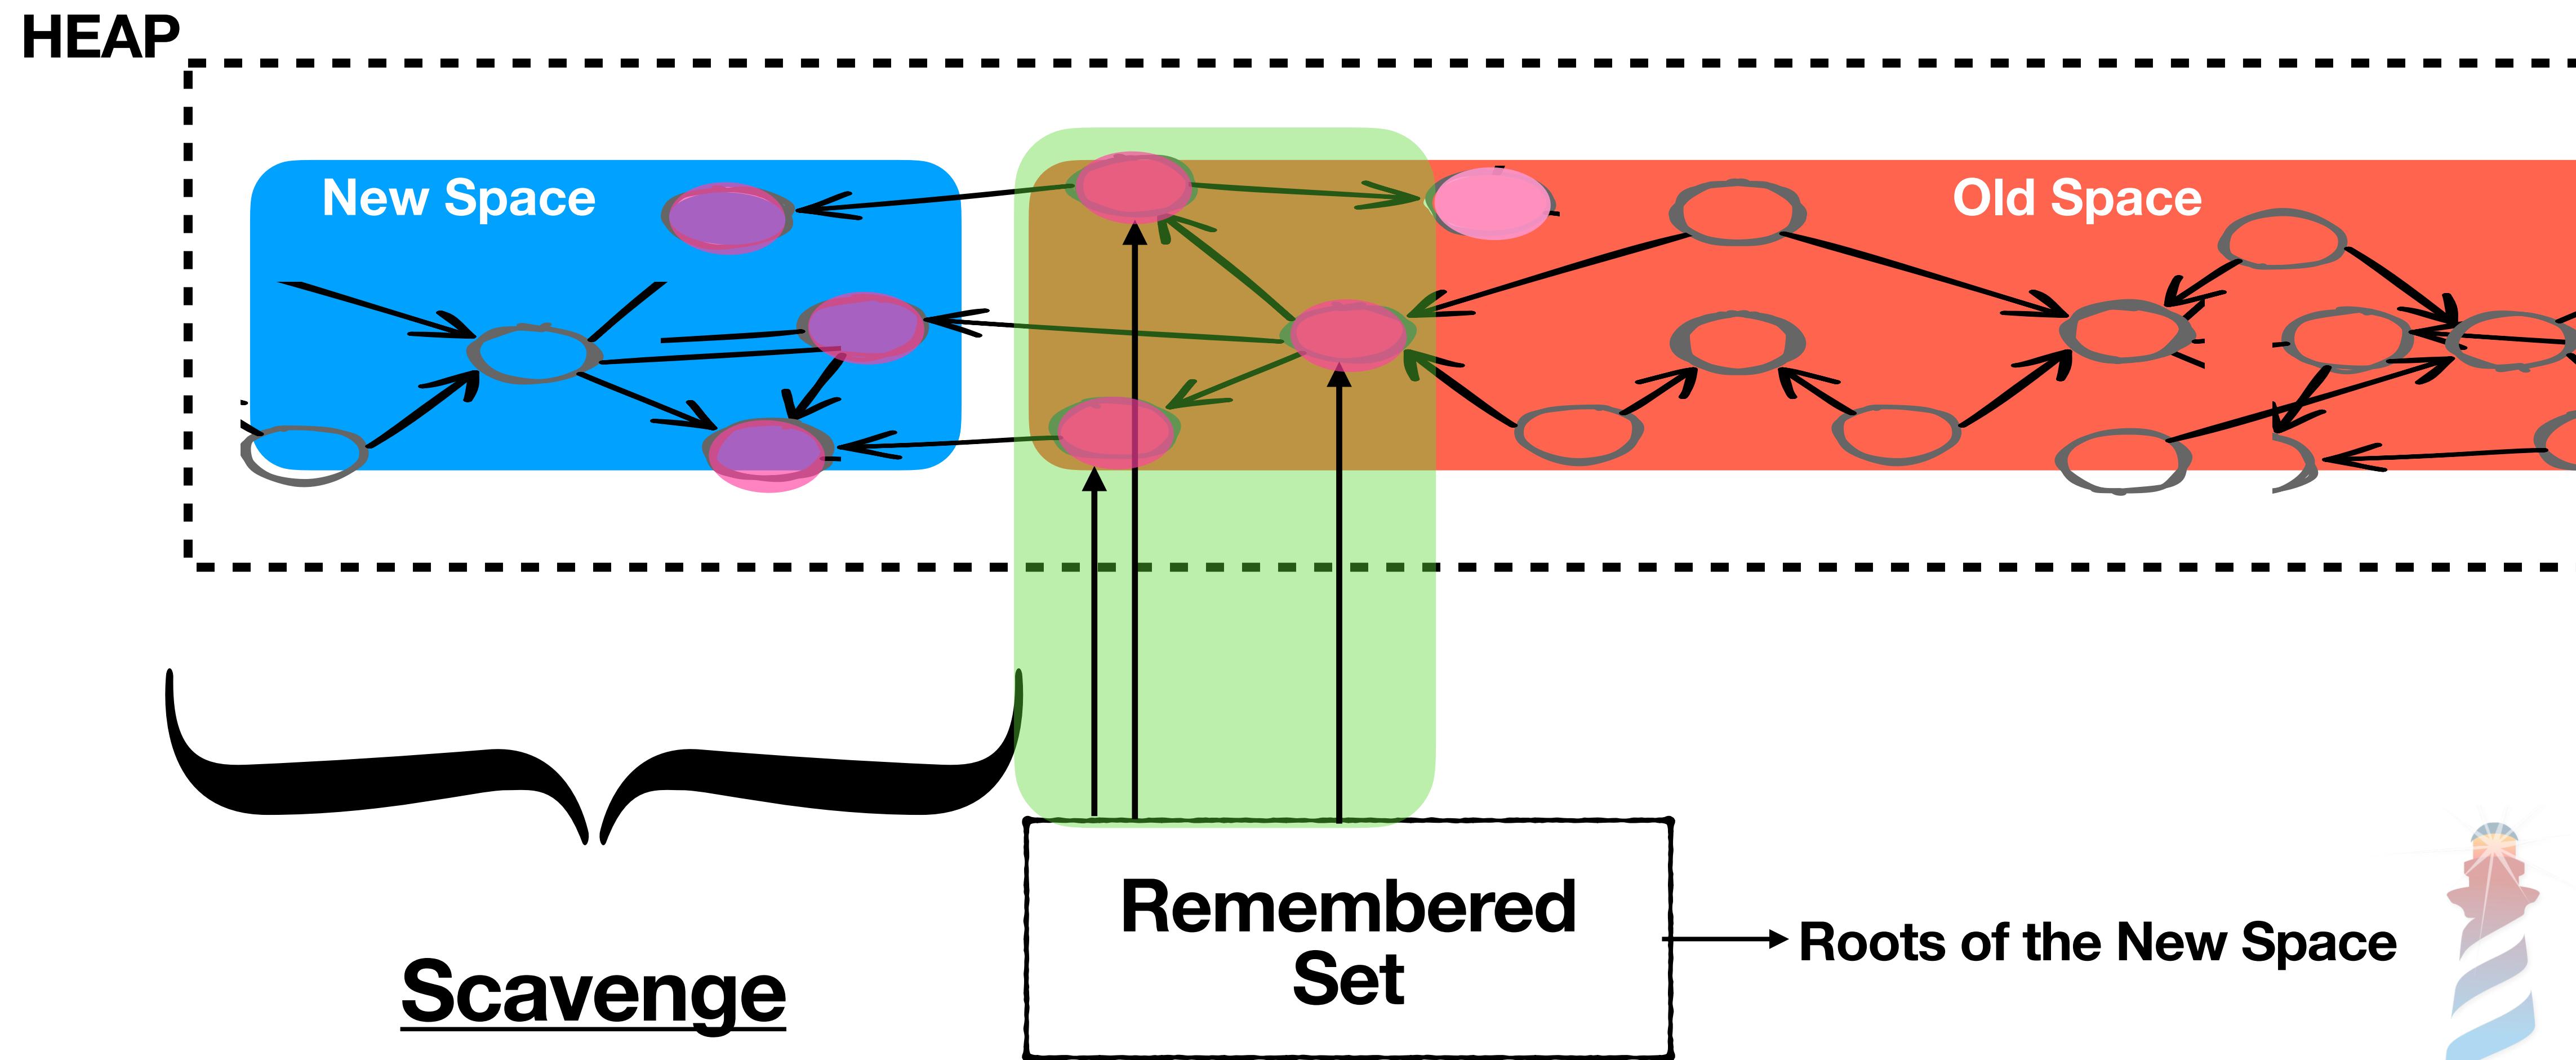
Deeper in the allocation pattern

Generational clash



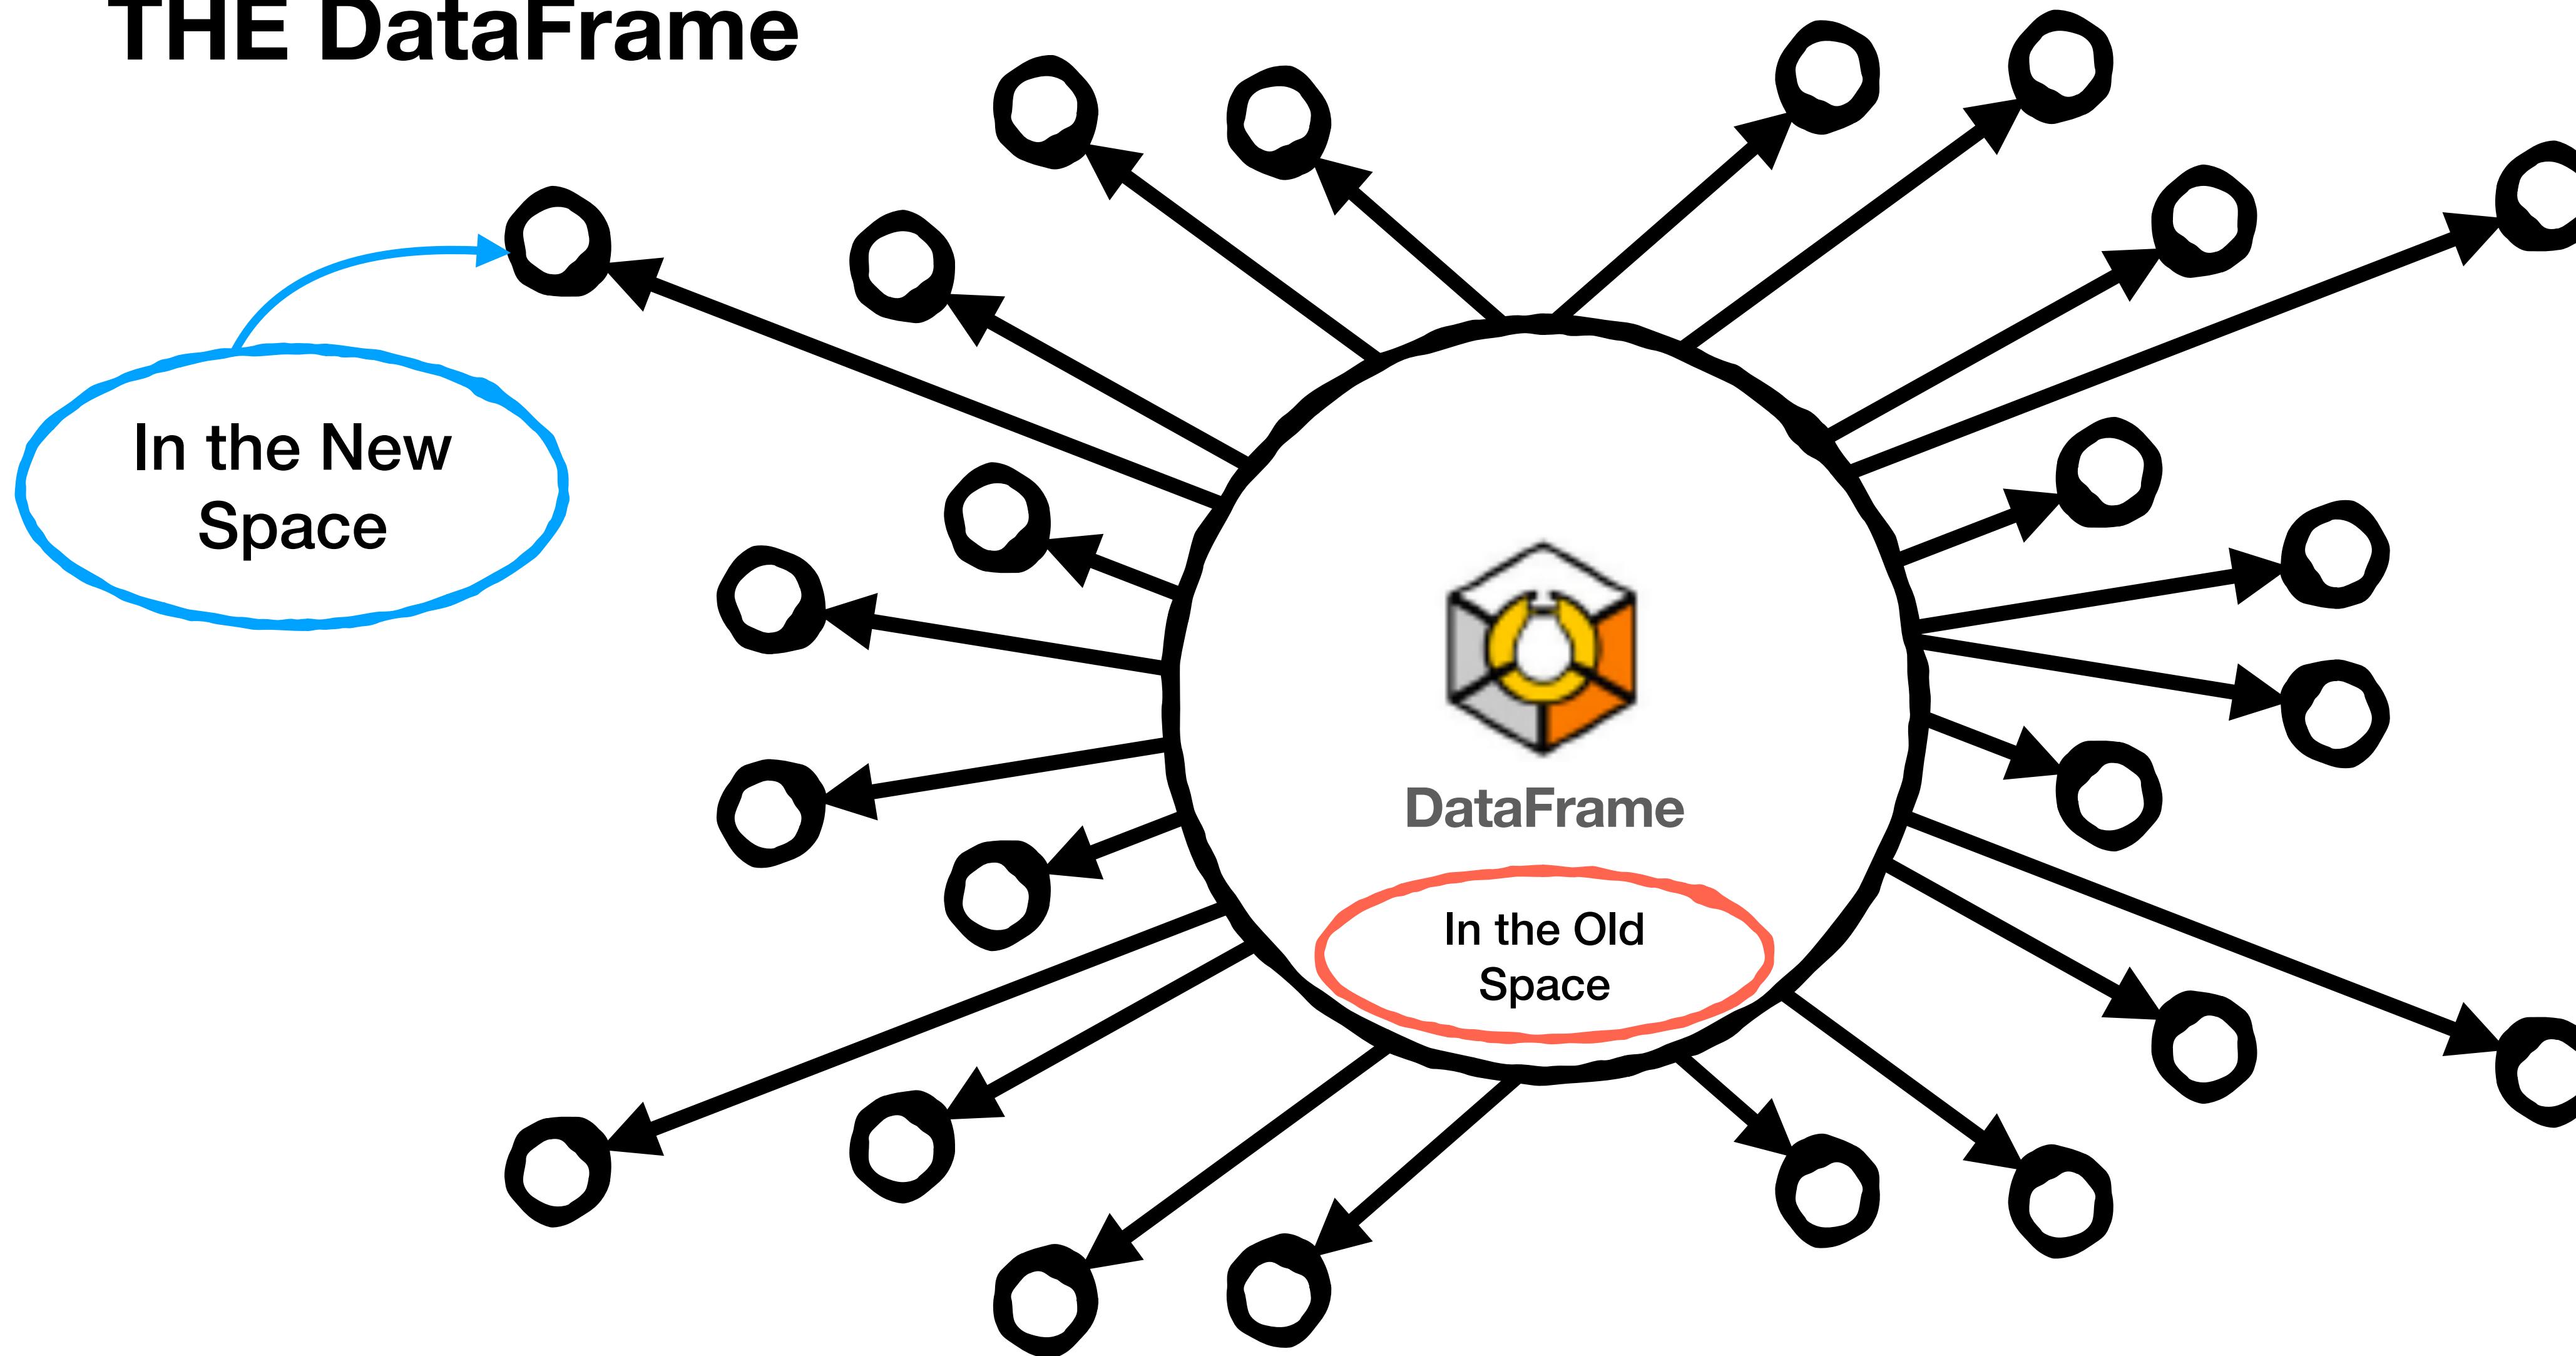
Deeper in the allocation pattern

Generational clash



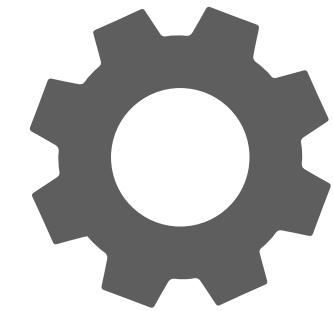
Deeper in the allocation pattern

THE DataFrame

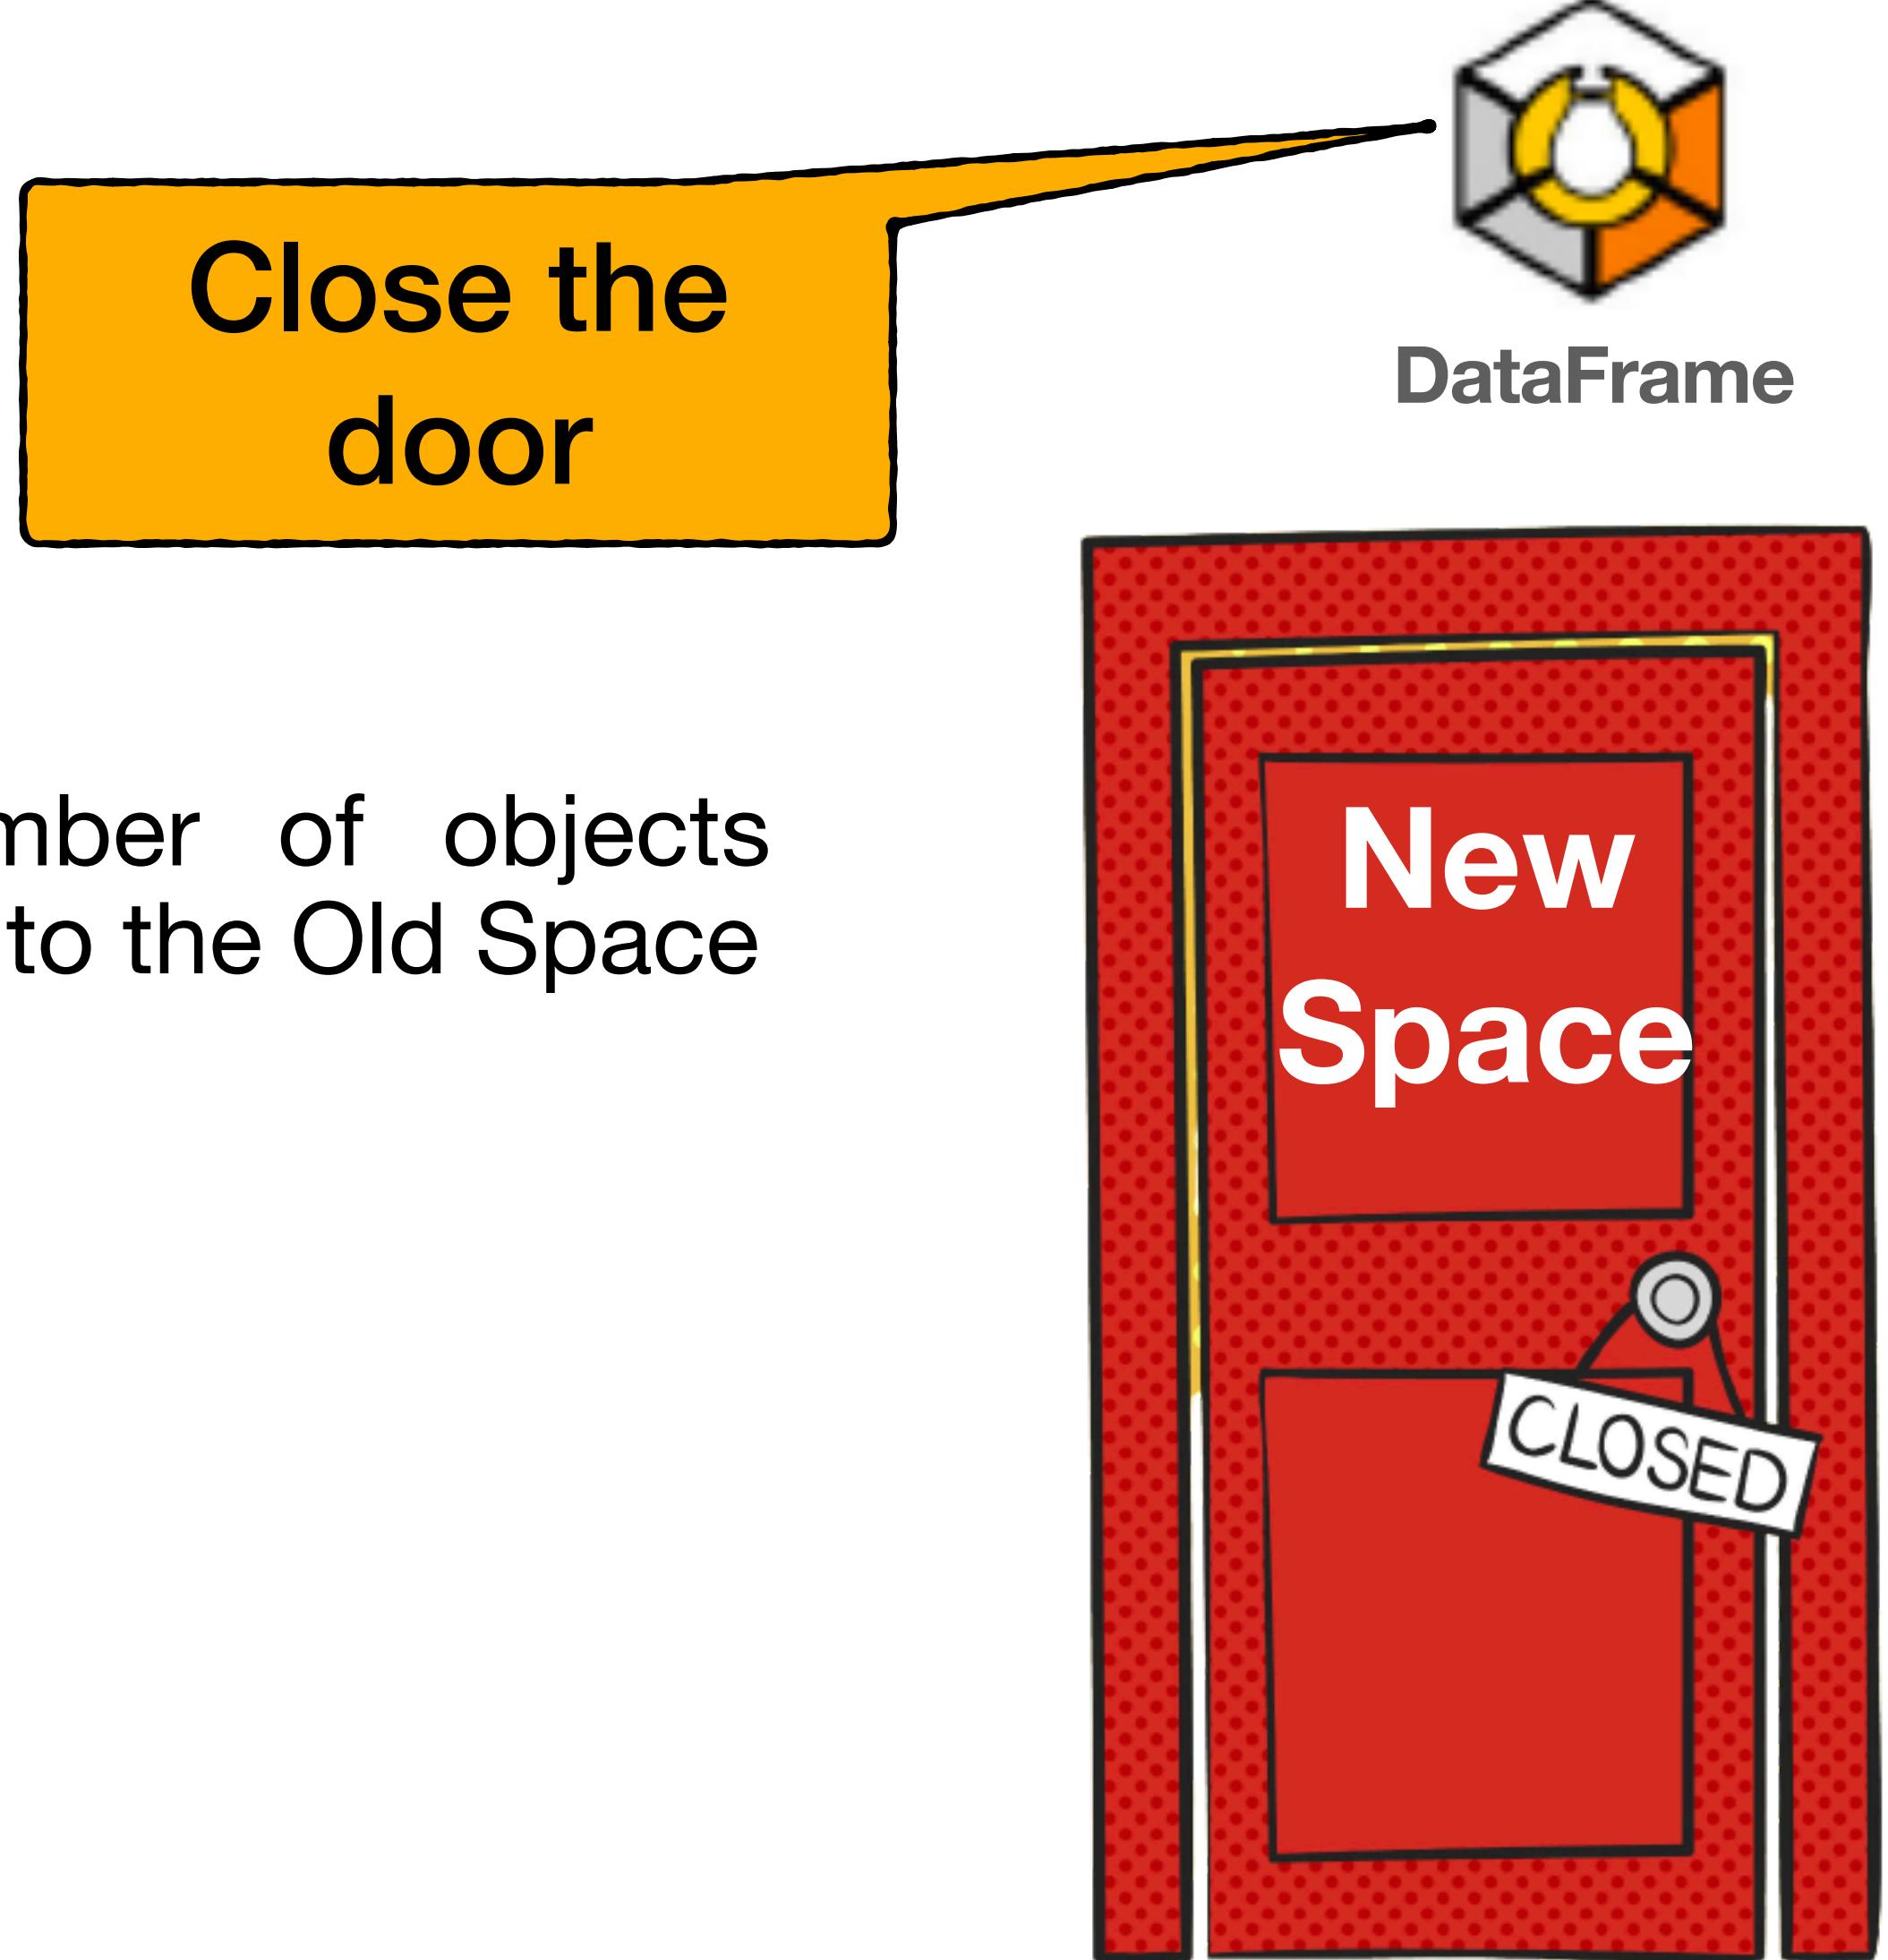


Long Scavenges

The tuning solution

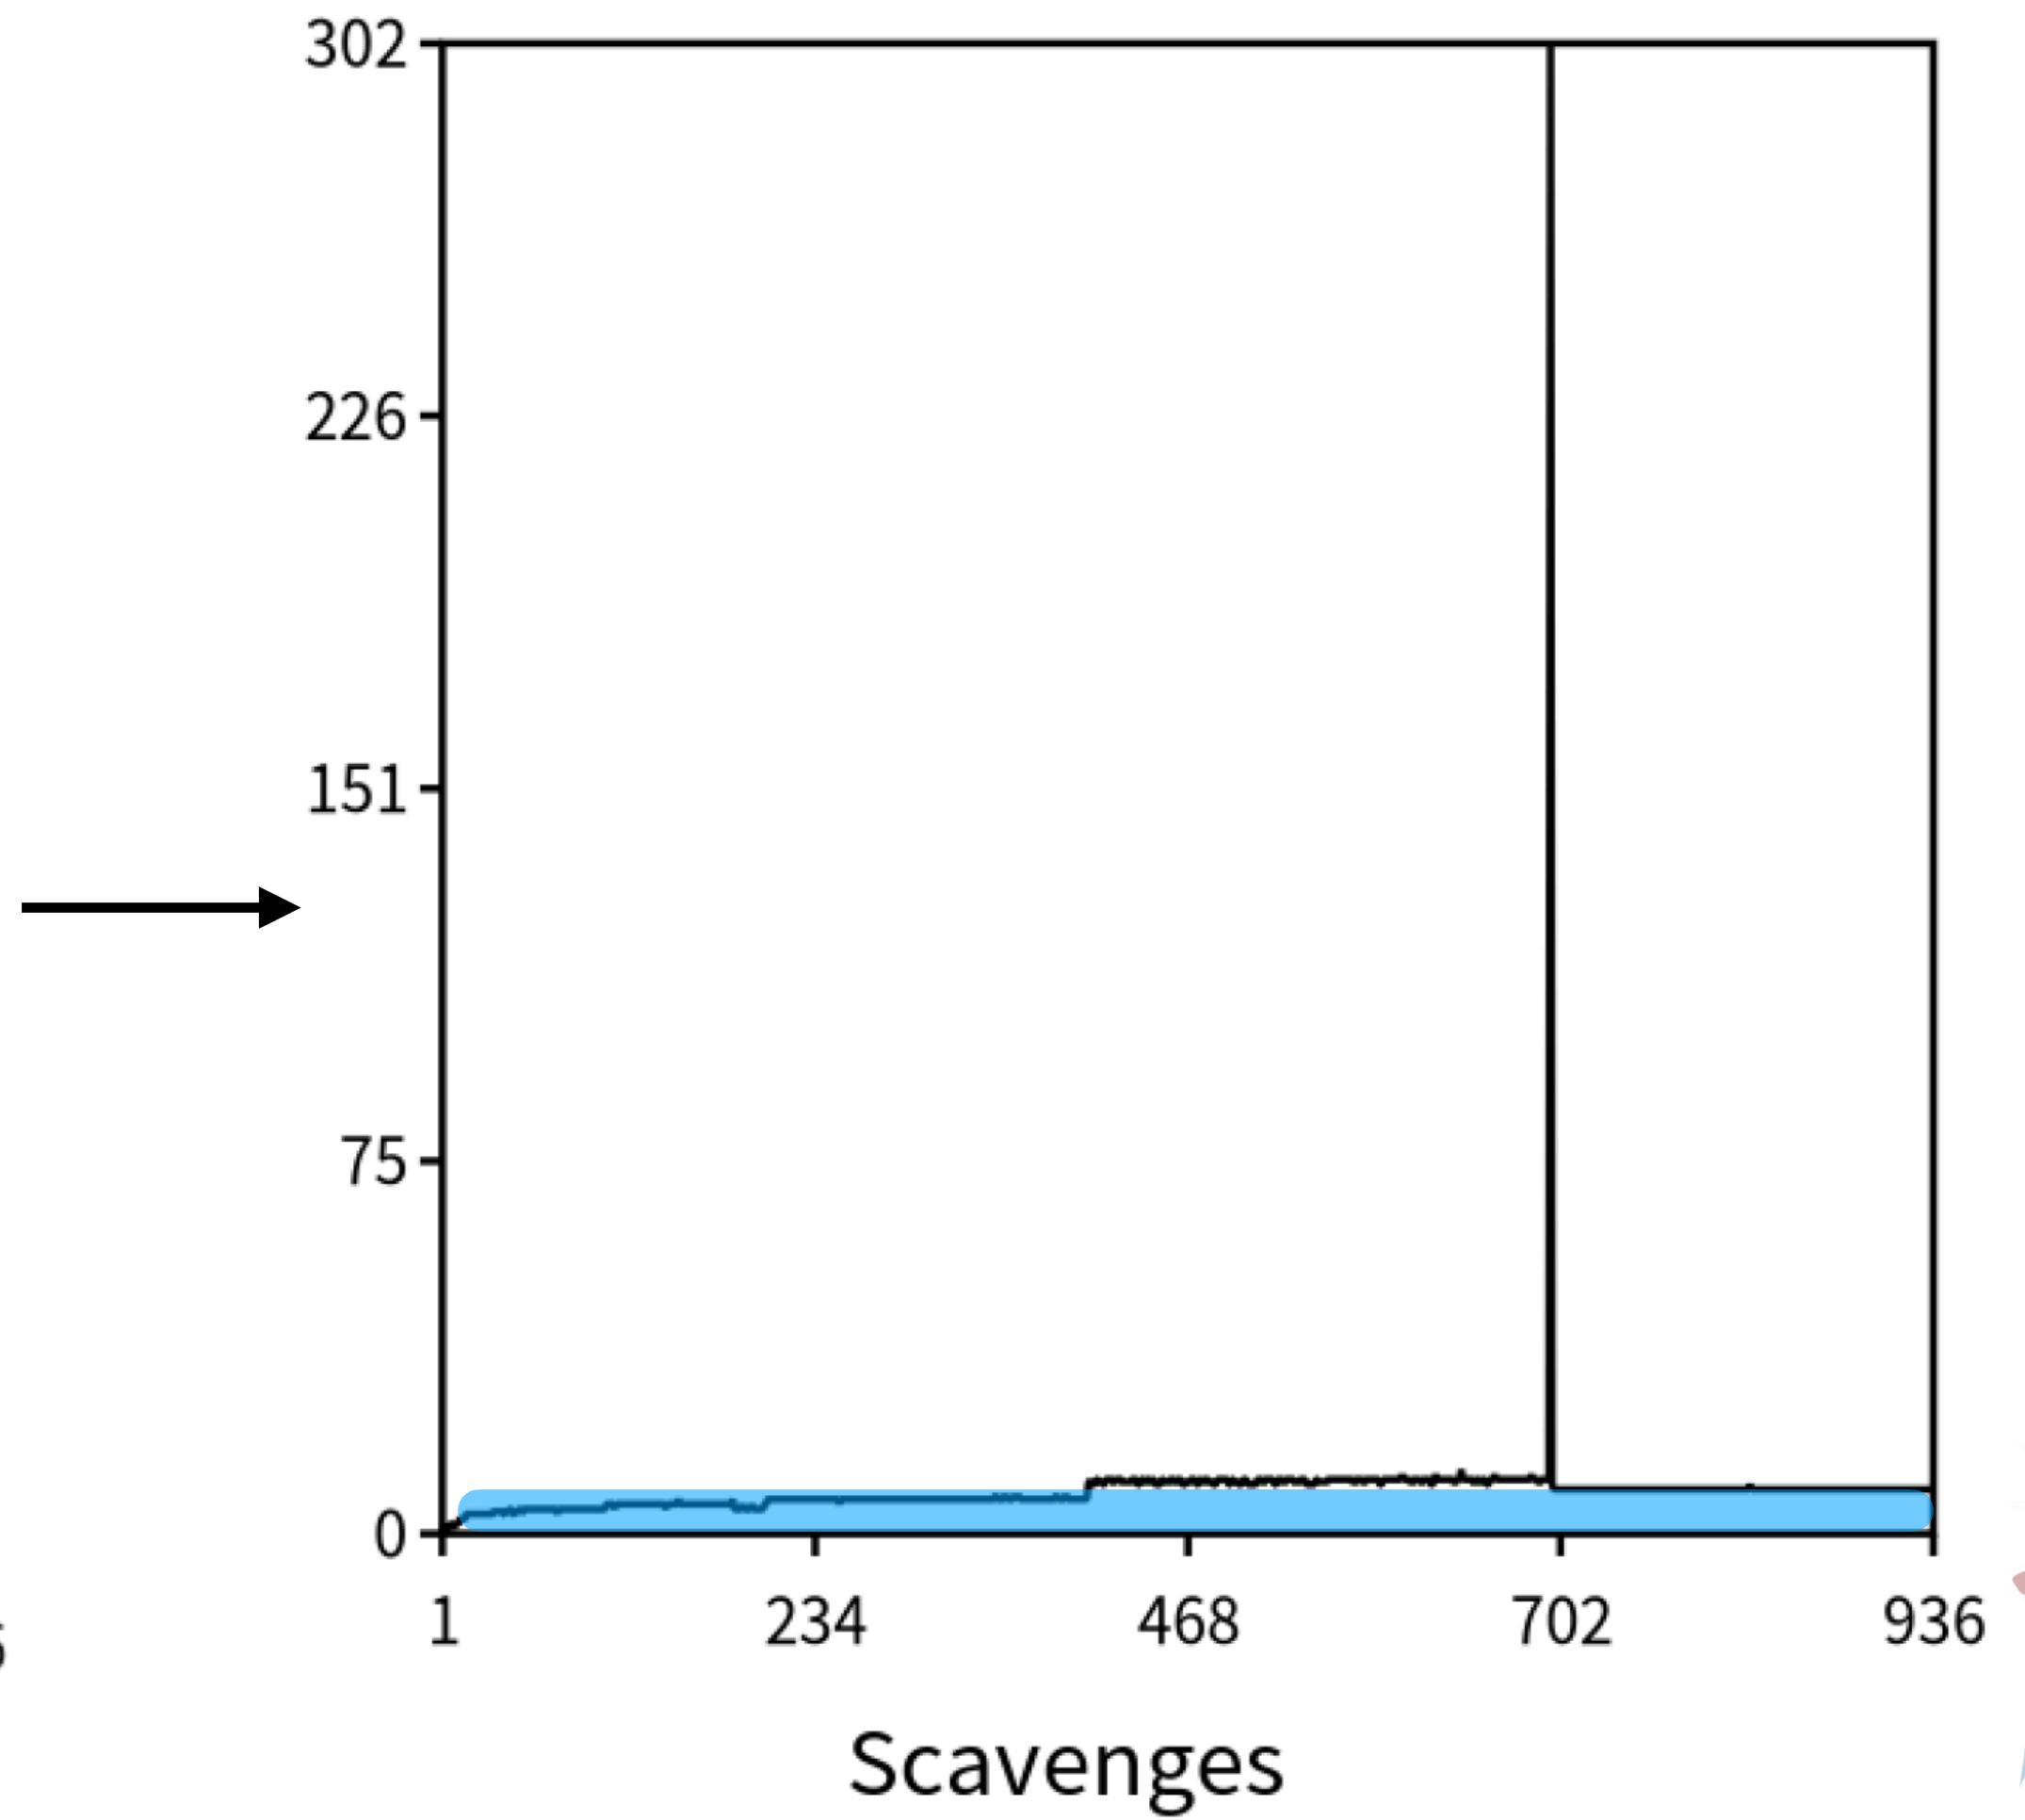
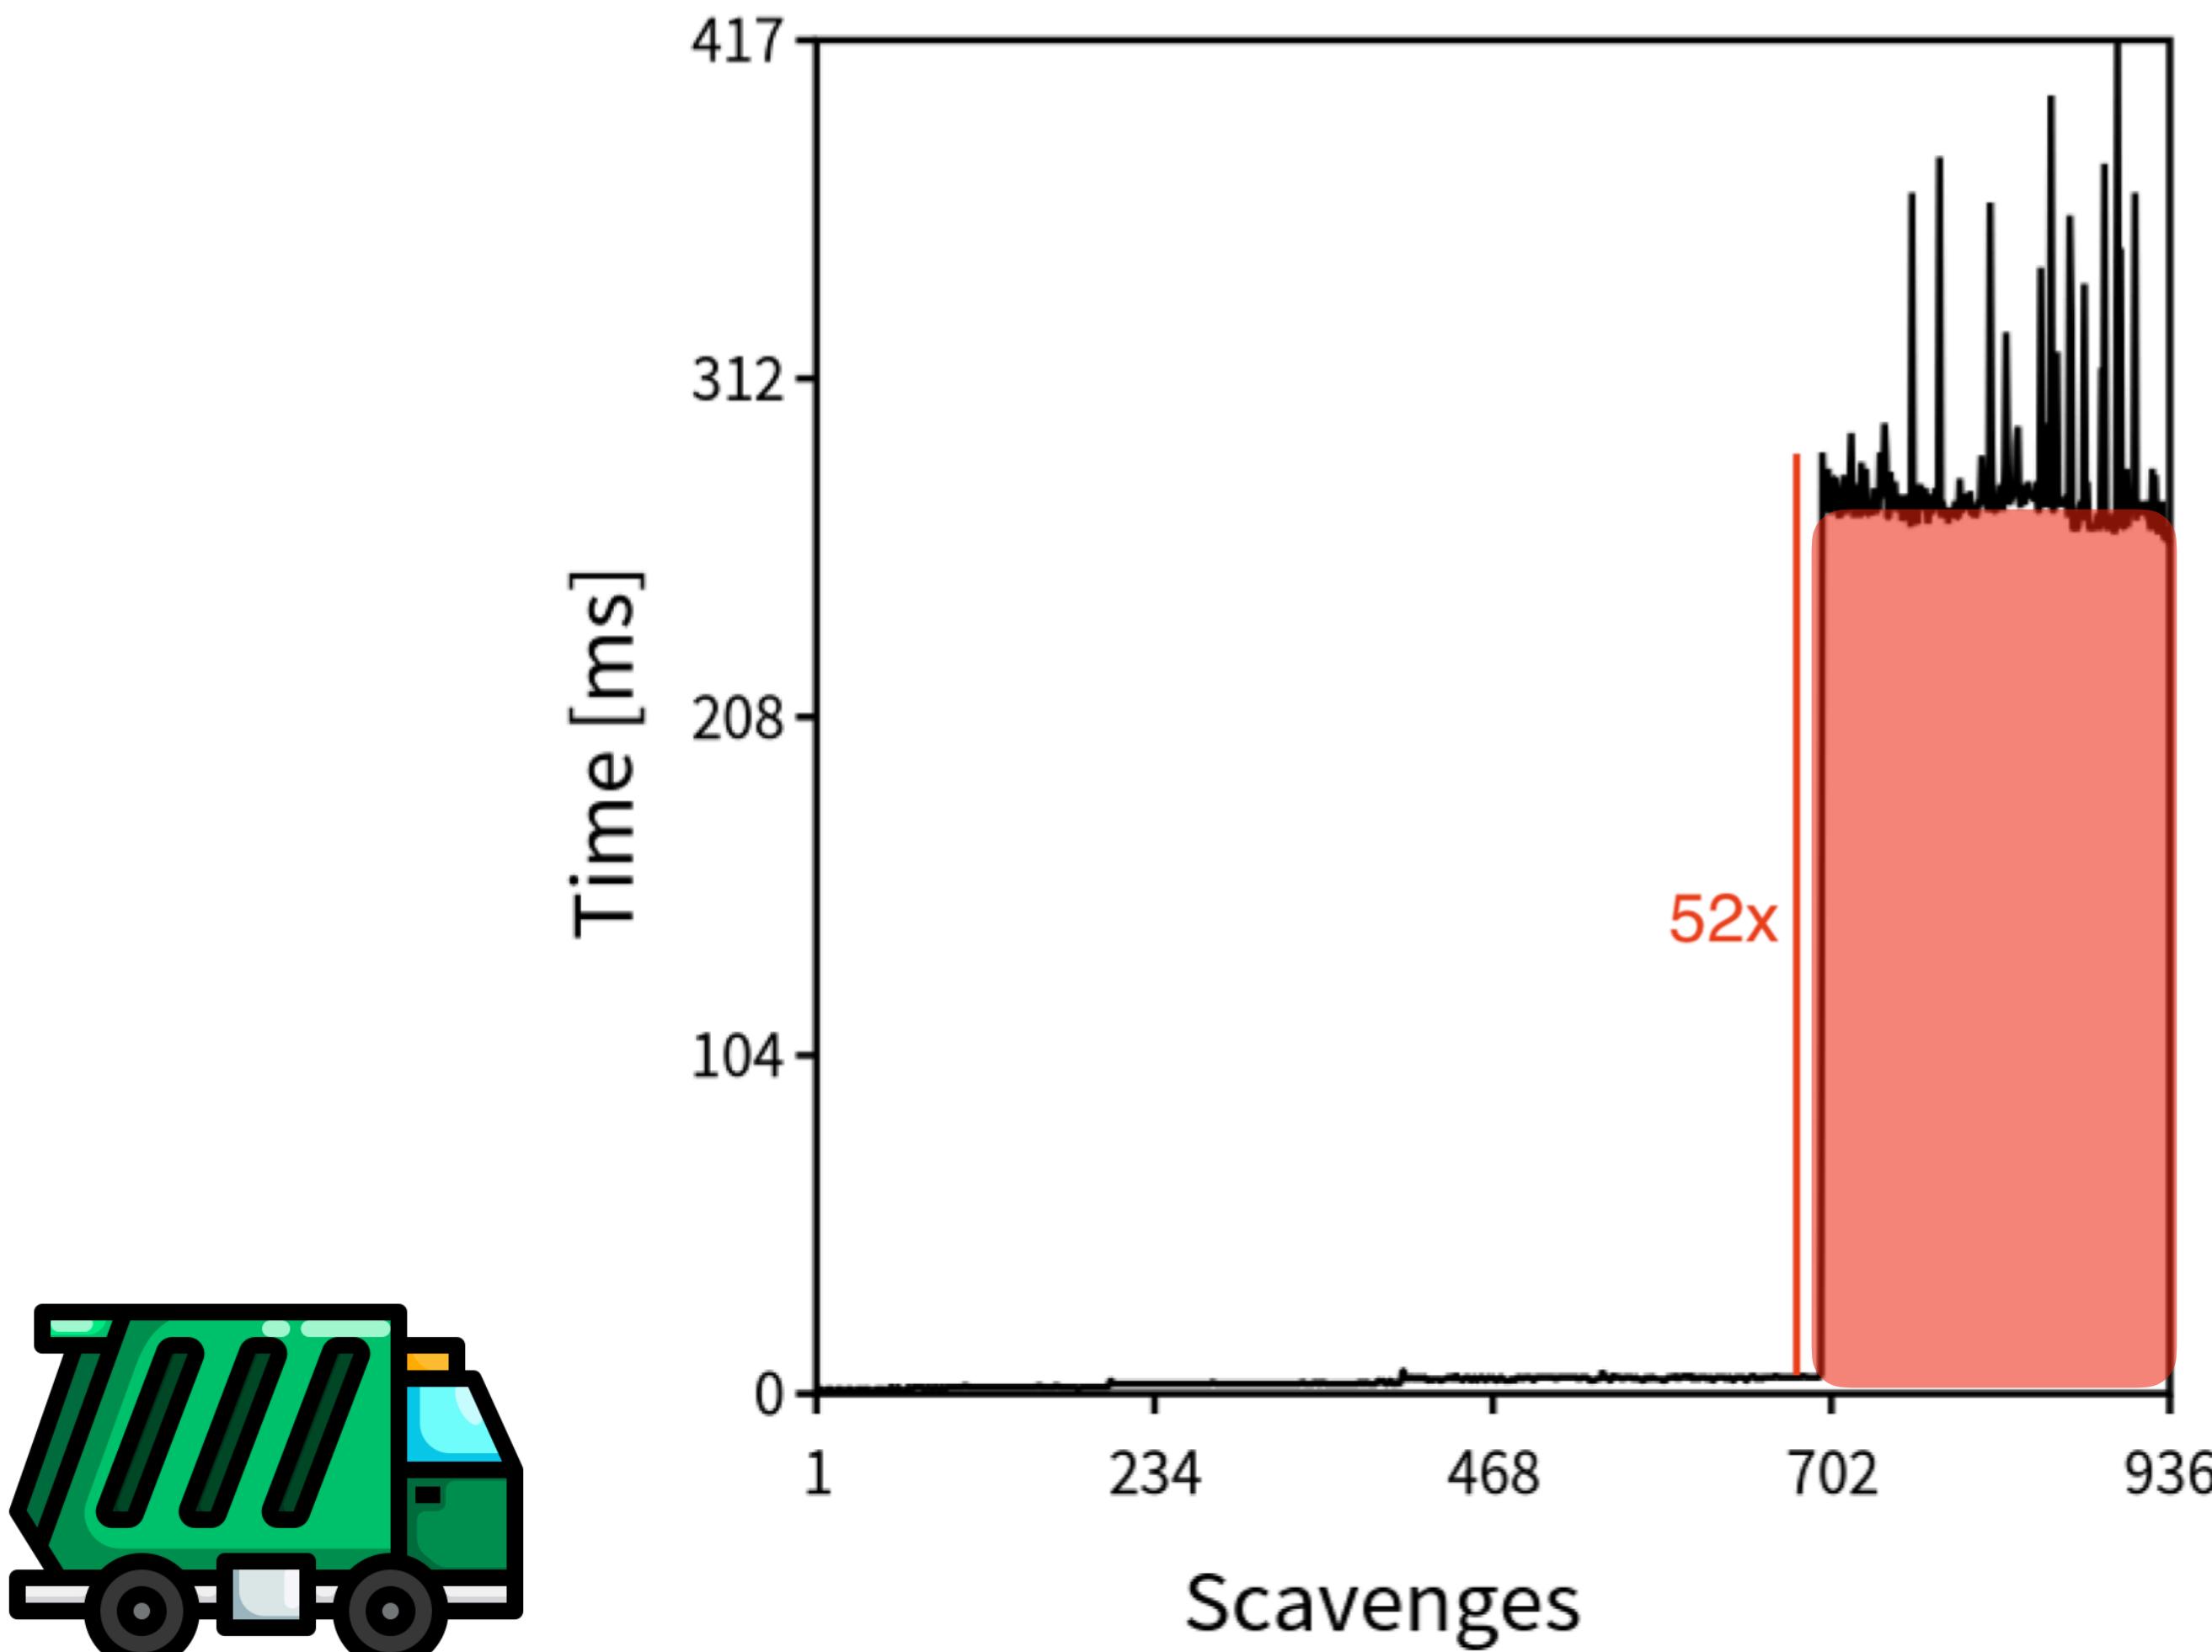


Tenuring threshold - Desired number of objects already in the New Space for tenuring to the Old Space



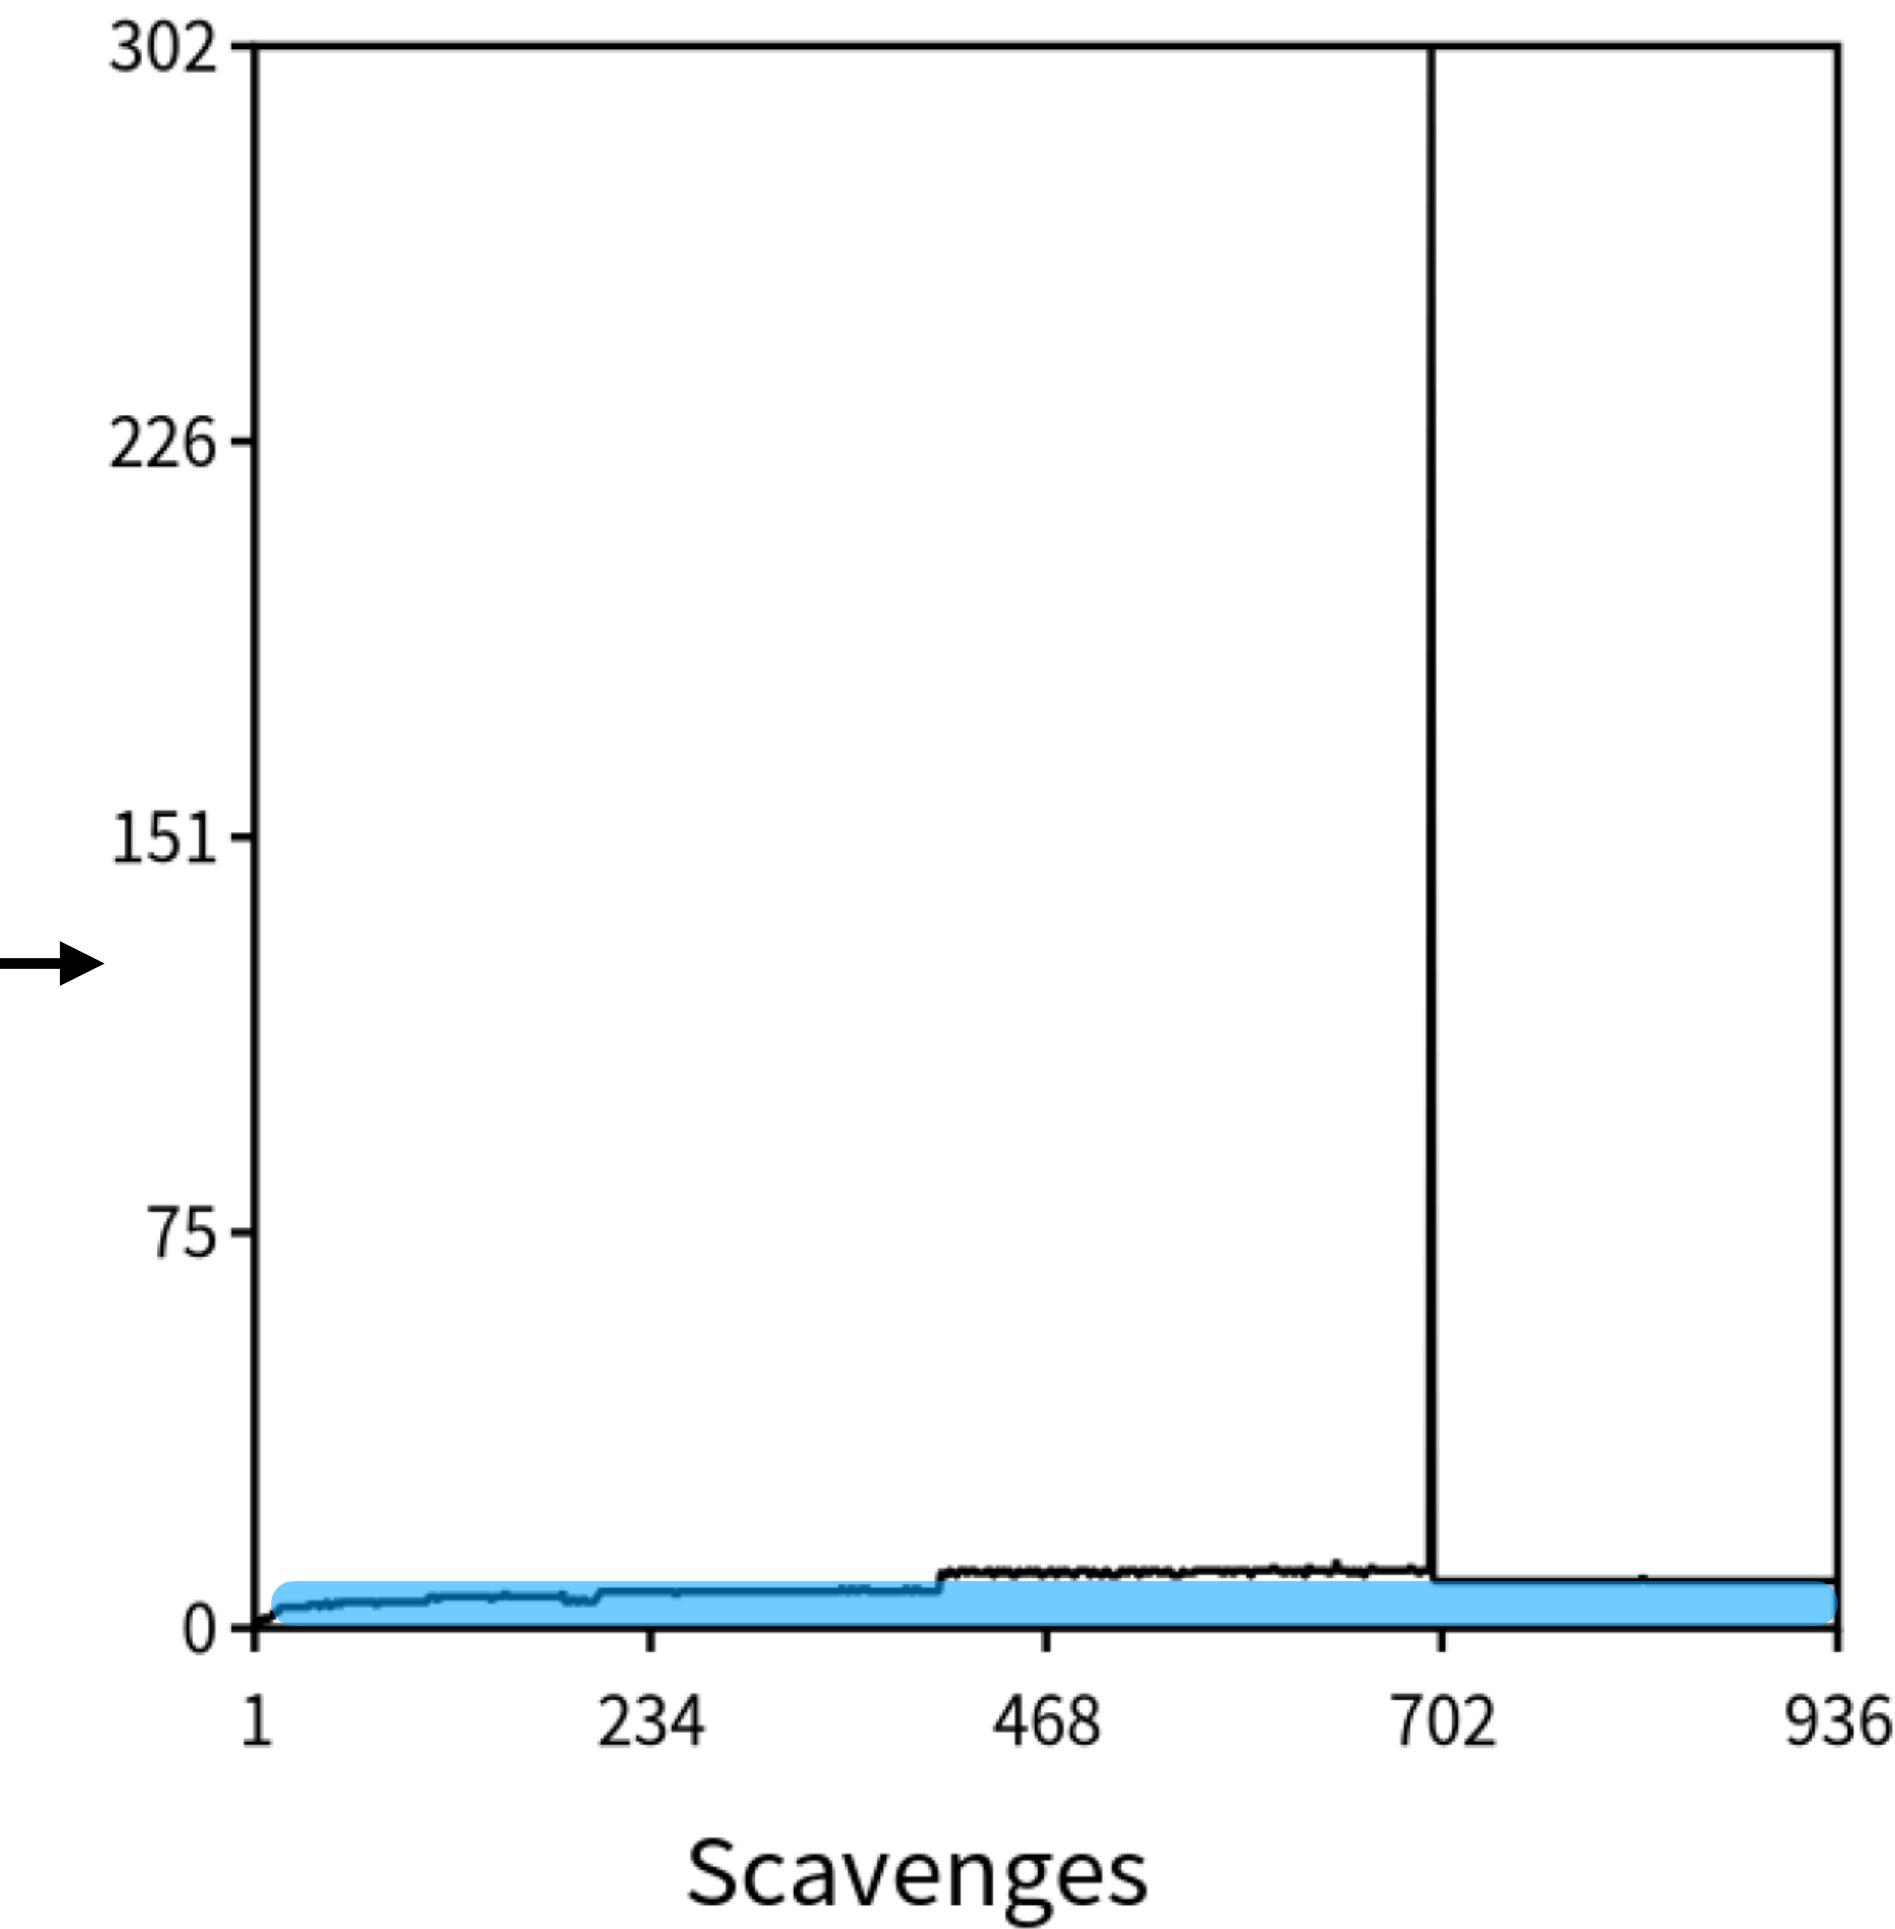
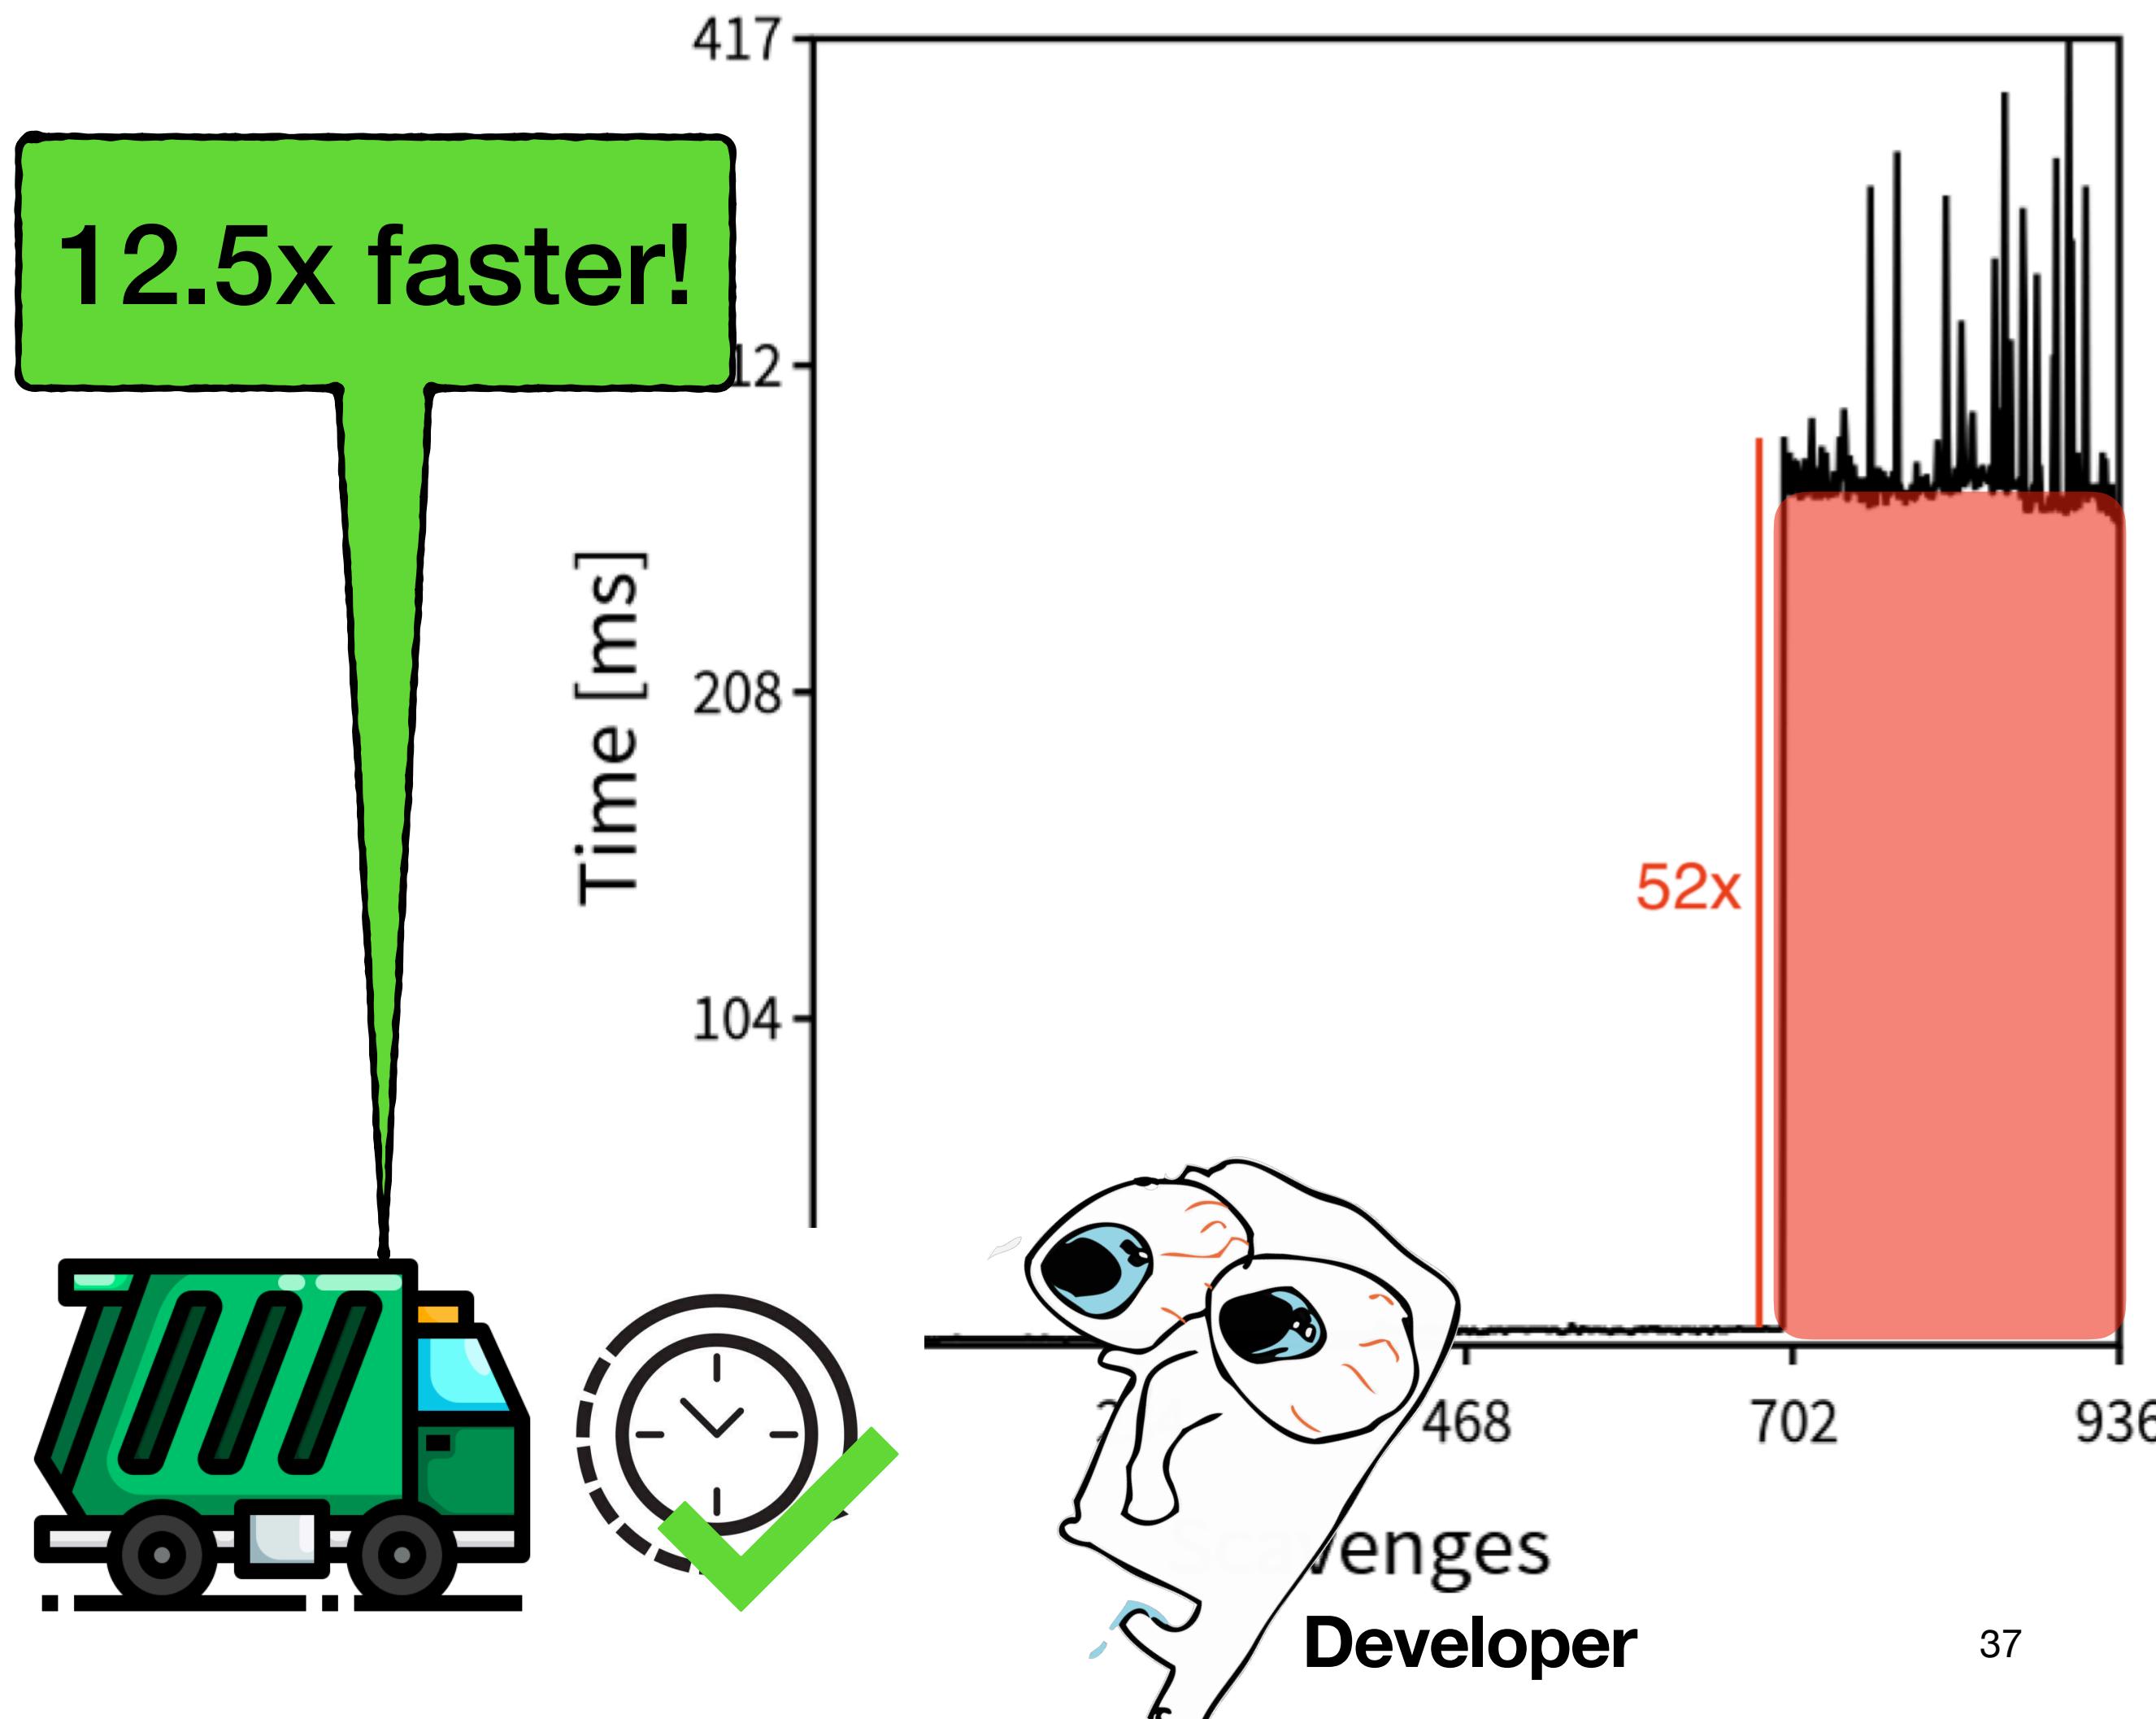
Long Scavenges

The tuning solution



Long Scavenges

The tuning solution



Final result

- 1) Have an infinite FullGC ratio** To reduce the number of FullGC when the Old Space grows.
- 2) Have a grow headroom equal to the loaded file** To avoid many FullGC together.
- 3) Keep all survivors in the semi-space** To tenure new objects to the Old Space quickly.

Data size	Total secs before	GC overhead before	Total secs after	GC overhead after
529 MB	43	16%	37 (1.1x)	5% (3.2x)
1.6 GB	150	25%	122 (1.2x)	7% (3.6x)
3.1 GB	5599 >1h30m	92%	440 (12.5x)	24% (3.8x)

~7mins





A large, stylized word cloud centered around the words "thank you" in various languages. The words are rendered in different colors and sizes, creating a dynamic and international feel. The background is white, making the colorful text stand out.

Understanding Garbage Collectors in object-oriented programming

Nahuel Palumbo



 PalumboN



 nahuel.palumbo@inria.fr

